

# PALETTE-Aug: A Simple Texture-Preserving Augmentation for Medical Imaging Domain Generalization

Anonymous CVPR submission

Paper ID \*\*\*\*\*

## Abstract

**001** ***Introduction.** Segmentation networks trained on one MRI*  
**002** *contrast or imaging modality often fail on another, and an-*  
**003** *notating data for every new domain is costly. Synthesis-*  
**004** *-based domain randomization such as SynthSeg generates*  
**005** *training images from label maps alone, but it needs dense*  
**006** *label maps and discards the texture of the real image. We*  
**007** *introduce PALETTE-Aug, a texture-preserving alternative*  
**008** *that needs only the task’s own labels, and ask when tex-*  
**009** *ture matters for generalization. **Methods.** PALETTE-Aug*  
**010** *remaps the real intensities of an unsupervised partition of*  
**011** *the image, so the texture inside every region survives. We*  
**012** *benchmark it against five methods on eight segmentation*  
**013** *tasks across MRI, CT and dental CBCT: five whose tar-*  
**014** *gets are bounded by an anatomical interface (healthy brain,*  
**015** *multi-organ, healthy spine, mandible, pelvis) and three de-*  
**016** *finied by tissue appearance (brain tumor, multiple sclero-*  
**017** *sis lesions, breast lesions). Each method trains on one*  
**018** *modality and is tested in-domain and out-of-domain. An*  
**019** *ablation over sixteen training settings holds the partition*  
**020** *fixed and replaces noise fill with the real-intensity remap.*  
**021** ***Results.** PALETTE-Aug wins the benchmark overall: it*  
**022** *has the highest Dice on seven of the eight tasks and, with*  
**023** *all tasks pooled, is significantly better than every competi-*  
**024** *tor. In the ablation, keeping texture raises Dice by +6.8*  
**025** *points on multiple sclerosis lesions and +4.6 on brain tu-*  
**026** *mor, both appearance-defined, against at most +3.5 on*  
**027** *interface-bounded tasks. Breast lesions, also appearance-*  
**028** *defined, gain only +0.5. **Conclusion.** Texture preservation*  
**029** *helps most on targets defined by tissue appearance rather*  
**030** *than an anatomical interface, though not on all of them.*  
**031** *PALETTE-Aug generalizes best overall across eight seg-*  
**032** *mentation tasks.*

## **033** 1. Introduction

**034** Medical image segmentation models are often trained on  
**035** one acquisition setting [18] and frequently applied to an-

other: a different scanner, sequence, or imaging modal-  
ity than the one seen during training [17, 30]. MRI con-  
trast is a protocol choice, not a physical constant, so the  
same anatomy has different intensities across sequences [5].  
A network trained only on T1-weighted brain images can  
fail outright on T2-weighted or FLAIR images [3]. Ac-  
quiring and annotating training data for every contrast a  
model might encounter is not practical: manual segmen-  
tation is expensive, and clinical sites rarely share raw data  
across institutions [39, 57]. The setting that matters for de-  
ployment is therefore single-source domain generalization:  
train on one labeled dataset in one contrast, then generalize  
to unseen contrasts and modalities without target-domain  
data [42, 51, 61, 63].

The dominant single-source strategy is label-generative  
domain randomization, popularized by SynthSeg [3].  
SynthSeg does not use the real image at all: it builds each  
training image from a label map, filling every labeled struc-  
ture with Gaussian noise of a randomly drawn mean and  
variance. The label map must be *dense*, a voxel-wise parcel-  
lation of the whole image rather than the sparse annotation a  
segmentation dataset normally provides. The resulting im-  
ages have arbitrary, unrealistic contrast, forcing the network  
to rely on shape and layout [3]. This is effective, but carries  
two structural costs. First, it is *label-driven*: the dense par-  
cellation it needs is richer than the task’s own annotation.  
SynthSeg trains on whole-brain parcellations with dozens  
of labeled structures [3], whereas an MS dataset such as  
Open-MS annotates only the lesions [35]; obtaining the  
parcellation reintroduces the expensive labeling the method  
aims to avoid. Second, and less discussed, it *does not pre-*  
*serve texture*: replacing every labeled region with piecewise  
Gaussian noise discards the real gradients, partial-volume  
boundaries, and unlabeled fine structure the acquired scan  
contains. Layout survives; texture does not.

We ask when keeping texture, rather than discarding it,  
helps contrast robustness. Our central finding is that the  
answer depends on the *task*: preserving real intensity struc-  
ture helps most where the target is defined by tissue appear-  
ance, and much less where an anatomical interface sepa-

rates it from its surroundings. Prior work shows that texture carries diagnostic signal in appearance-defined tasks: texture features separate MS lesions from normal-appearing white matter [62], and diffuse glioma infiltrates beyond every margin visible on conventional MRI [21, 53].

**PALETTE-Aug**, our single-source augmentation, keeps this texture and needs no dense label map. **PALETTE-Aug** groups the foreground intensities of a scan into classes with 1-D  $k$ -means, splits each class into spatial sub-regions with a Voronoi partition, and applies a signed random affine remap to the *real* intensities of each sub-region. Edges, gradients and partial-volume detail therefore survive inside each sub-region, because existing intensities are remapped rather than replaced by noise. Only the contrast between regions is randomized, and without a target: **PALETTE-Aug** can produce any contrast, including implausible ones such as inverted intensities, because it is not tied to a specific target contrast. An optional step also re-randomizes each labeled structure; it uses whatever labels exist, a dense label map or just the task’s own segmentation labels.

We then test when texture matters, with a controlled ablation in sixteen training settings. We keep **PALETTE-Aug**’s partition and region target means and change only what fills each region: the region’s own image content, affinely remapped so that its gradients and fine structure survive (texture kept), or Gaussian noise (texture removed). The four largest Dice gains are all on appearance-defined targets: brain tumor (+12.7 points from T1c training, +6.0 from T1n) and MS lesions (+7.7 from FLAIR, +5.9 from T1w), against at most +4.8 in any interface-bounded setting. Pooled per task, the swap adds +6.8 Dice on MS and +4.6 on brain tumor, against at most +3.5 on an interface-bounded task. The converse does not hold: breast lesions, also appearance-defined, gain only +0.5, and several appearance-defined settings show no gain, so we state the result in one direction only (§4.3). A segmentation-free gradient measurement confirms that the swap changes texture: noise fill keeps only the partition edges, and real fill recovers most of the source texture.

We train one nnU-Net [27] configuration, varying only the epoch budget per dataset, and evaluate cross-contrast and cross-modality transfer on eight tasks: healthy brain, MS lesion, brain tumor, abdominal organ, mandible, breast lesion, spine and pelvis. We compare against nnU-Net with its default augmentation; the released SynthSeg (SynthSeg-noEM) and its expectation-maximization variant for cardiac MRI [3, §5.4] (SynthSeg-EM), whose code was not released, so we re-implemented it; SRCSM [54]; and the Auglab augmentation pipeline [40]. **PALETTE-Aug** wins the benchmark overall: it has the best Dice on seven of the eight tasks and significantly beats every competitor on five of them. Pooling all tasks into one test, it is significantly better than each competitor on both Dice and boundary dis-

tance, though not on every task.

In summary, we make the following contributions:

- **When texture preservation matters, tested by intervention.** Replacing noise fill with the real-intensity remap, with the partition fixed, gives its largest gains on appearance-defined targets, though not on every one. A segmentation-free gradient measurement confirms that the swap changes texture.
- **PALETTE-Aug**, a single-source contrast augmentation that needs no dense label map, only the task’s own labels. In every comparison it runs as one added operator inside the Auglab pipeline; the ablation also tests it alone.
- **Breadth evidence.** Across cross-contrast and cross-modality transfer on eight tasks and seventeen datasets, with one configuration reused throughout, **PALETTE-Aug** wins the benchmark overall: it is best on seven tasks and significantly beats every competitor on five; where it does not win, we test whether the leader is truly ahead.

## 2. Related Work

### 2.1. Domain generalization for medical segmentation

Domain generalization (DG) targets unseen distributions without access to target data, unlike unsupervised domain adaptation, which assumes unlabeled target images and solves an easier problem [50]. Most medical DG methods are *multi-source*, interpolating across several labeled domains by meta-learning, feature alignment, frequency-space exchange, multi-frequency architectures, foundation-model scale, or diffusion-based stylization [2, 36, 39, 41, 48, 57] – all of which still rely on source diversity. We instead work in the *single-source* regime (one labeled dataset, no target data, no second source), which is routine in clinical deployment [42, 63]. This rules out cross-domain interpolation and makes architecture-coupled methods orthogonal to our contribution: our mechanism acts at the input level and leaves the segmentation network a fixed nnU-Net [27].

### 2.2. Synthesis-based contrast-agnostic training

The dominant single-source strategy is label-generative domain randomization, introduced by SynthSeg [3]: intensities are sampled from a per-label Gaussian mixture over an anatomical segmentation, yielding images of arbitrary contrast that force contrast-invariance. This family has the two costs noted in Sec. 1: it needs a dense anatomical label map per subject, and it replaces real texture with piecewise noise. SIAM [55] trades label *quantity* for *quality* with curated templates, but still needs a full dense parcellation. The lineage has been specialized to white-matter and MS lesions and to atlas-based expectation-maximization [7, 8, 34]. Among SynthSeg’s own variants




Figure 1. The PALETTE-Aug transform on one real CHAOS T2SPIR slice (top) and one real ON-Harmony T1w slice (bottom), in the order of its four steps (Sec. 3); every panel is a deterministic function of the same real image. (b) and (c) paint each region flat with its target mean  $\mu$  to show the partition alone:  $k$ -means intensity classes (b, step 1) split into Voronoi sub-regions (c, step 2). (d) applies the signed affine remap of Eq. (1) to the real intensities of each sub-region (step 3), which keeps their texture; (e) adds the label remap (step 4), an affine re-map inside each anatomical label (its straight edges are label boundaries, e.g. between the two hemispheres), and is the training input.

we take the original (*noEM*), which synthesizes from the label map alone and never looks at the acquired scan, and the *EM* extension, which first splits every label, background included, into intensity-coherent sub-regions by expectation-maximization clustering of the real image, so that sparse task labels yield a richer map to synthesize from. SynthSeg introduced the EM extension for cardiac MRI [3, §5.4]; we re-implement it.

### 2.3. Texture as a segmentation signal

ImageNet-trained CNNs classify more by texture than by shape, and training them on style-randomized images shifts them toward shape and improves robustness [15]; how strong this texture bias is depends mostly on the training data and augmentation [24]. SynthSeg-style randomization pulls the same lever on medical images: replacing texture with noise forces the network onto shape. We ask what that costs when the target is itself defined by its appearance. In medical images, texture carries such information: second-order gray-level statistics discriminate image content beyond intensity [22], and texture features from tumor regions carry prognostic signal across CT, MRI and PET [33]. Closer to our setting, Kang et al. disentangle structure from texture in cross-domain image translation with a texture co-occurrence loss, targeting the same failure mode we do — naive translation distorts fine texture and degrades downstream segmentation [29]; a texture-fidelity loss has like-

wise been added to GAN-based lesion augmentation [19]. For diffuse white-matter lesions, it has been argued directly that the lack of a sharp edge defeats boundary-based methods and favors texture-based segmentation [32]. None of this prior work, to our knowledge, connects texture preservation in an *augmentation* to *which* segmentation targets it helps; that connection, tested causally, is our contribution.

### 2.4. Single-source intensity and appearance augmentation

Where the first line of single-source methods synthesizes from a label map, a second manufactures appearance diversity directly from the real image, with no generative label model. General-purpose libraries stack intensity, spatial and artifact transforms on an nnU-Net trainer; the recent Auglab pipeline [27, 40] is one such GPU-accelerated framework and is the strongest competitor across our tasks, and BigAug stacks a comparable set of coarse intensity, quality and spatial perturbations under the same single-source DG protocol we use [63]. Our method inserts the PALETTE-Aug transform as one additional operator inside the Auglab pipeline, so the comparison against Auglab directly isolates PALETTE-Aug’s contribution. Closer to our operator are random-filter methods, which vary appearance with an untrained nonlinear filter: GIN applies randomly re-initialized convolutions for label-free appearance variants [42]; related variants establish the random-filter prin-

ciple for natural images, perturb intensities per class using the task mask, add learned diffusion generators, or perturb the Fourier amplitude spectrum instead of the image itself [9, 37, 45, 51, 61, 64]. Our operator instead partitions the foreground into intensity classes and Voronoi sub-regions and applies a signed *affine* remap of the real intensities within each, invertible and gradient-preserving by construction. The closest peer is SRCSCM [54]. Like PALETTE-Aug it perturbs region-wise in 3-D with one label-consuming step: it applies a fresh untrained convolution per anatomical class, smooths the resulting boundary artifacts, and histogram-matches test images to the source. They differ in the transform: SRCSCM applies a random nonlinear convolution, PALETTE-Aug an affine remap that keeps the relative intensity structure inside each region.

### 3. Method

We describe PALETTE-Aug (§3.1), how it is applied during training (§3.2), and the causal ablation used to isolate texture preservation as the source of its effect (§3.3).

#### 3.1. The PALETTE-Aug transform

Let  $x$  be a training volume min-max normalized to  $[0, 1]$ ,  $x_v$  its intensity at voxel  $v$ , and  $F = \{v : x_v > 0.01\}$  its foreground. PALETTE-Aug applies four steps at every training iteration, shown in order in Fig. 1(b)–(e).

**1. Intensity partition.** We run 1-D  $k$ -means on the foreground intensities  $\{x_v : v \in F\}$  with  $k$  sampled per iteration from  $\{2, \dots, 6\}$ , producing  $k$  intensity classes  $C_1, \dots, C_k$  (Fig. 1b). This partition groups voxels by intensity alone, with no label map, so it applies unchanged to any modality or contrast. Occasionally it is skipped and step 3 applies one global remap to the whole foreground; all step probabilities are in Sec. 7.

**2. Spatial sub-parcellation.** A Voronoi partition seeded with points sampled within each class  $C_i$  splits it into sub-regions  $R_{i,1}, \dots, R_{i,m_i}$  that are spatially localized and similar in intensity (Fig. 1c). Two voxels of equal intensity in different parts of the anatomy can thus fall in different sub-regions and receive independent contrast changes. Each class is usually, but not always, sub-split this way (Sec. 7).

**3. Signed affine remap.** For each sub-region  $R$ , let  $\bar{x}_R$  denote its mean intensity,  $\mu_R \sim U(0, 1)$  a resampled target mean, and draw  $\alpha_R \sim \pm U(0.5, 2.0)$  (a magnitude in  $[0.5, 2.0]$  with an independently sampled sign). Every voxel  $v \in R$  is remapped as

$$y_v = \mu_R + \alpha_R (x_v - \bar{x}_R). \quad (1)$$

Because Eq. (1) is a per-region affine function of the *real* intensity  $x_v$ , it preserves the relative structure of the signal within each region: edges, gradients, and partial-volume transitions between adjacent voxels are scaled, not replaced. Only the contrast between regions is randomized, and  $\alpha_R < 0$  also inverts local contrast (Fig. 1d). Background voxels are zeroed at this step (and remain zero through the rest of the pipeline). Half the time, the result is then Gaussian-blurred (Sec. 7).

**4. Anatomical-label affine remap (optional).** An intensity partition does not always separate the regions of interest: where a target’s boundary is already faint, steps 1–3 can place both sides in one region, or give them similar contrasts, and weaken it further. Step 4 therefore remaps each anatomical label  $\ell$  on its own: we draw a new target mean  $\mu_\ell \sim U(0, 1)$  and signed scale  $\alpha_\ell \sim \pm U(0.5, 2.0)$ , as in step 3, and apply Eq. (1) *again* to  $\ell$ ’s voxels of the step-3 output (Fig. 1e), so the label remap composes with the unsupervised one and keeps its internal structure. Since labels are independent of the unsupervised partition, the label remap restores a full-strength contrast step at every true label border, whether or not step 3 kept it; the anatomy itself is unchanged. It uses only the task’s own labels and is optional: without it, PALETTE-Aug reduces to steps 1–3, e.g. on unlabeled data. An optional second blur and a foreground z-score normalization then give the training input.

#### 3.2. Training protocol

PALETTE-Aug is applied on-the-fly as one contrast-augmentation operator inside an existing GPU-accelerated nnU-Net augmentation pipeline [27, 40], alongside that pipeline’s standard spatial and artifact augmentations; it adds no learned parameters and runs on the GPU inside that pipeline. One configuration ( $k$ -range, sub-region count, blur schedule,  $\alpha$  range; Sec. 7) is set once and reused unchanged for every dataset and modality: PALETTE-Aug’s hyperparameters are never tuned to a task.

#### 3.3. Causal ablation: isolating texture preservation

PALETTE-Aug differs from SynthSeg in two ways at once: the partition (unsupervised  $k$ -means and Voronoi vs. an anatomical label map) and the fill (remapped real intensities vs. Gaussian noise). A direct comparison cannot separate them.

We construct a controlled variant, *noise-fill*, that keeps PALETTE-Aug’s partition (steps 1 and 2) and draws the same target mean  $\mu_R$  for each sub-region, but replaces what fills it: instead of remapping the sub-region’s real intensities with Eq. (1), it fills the sub-region with SynthSeg-style Gaussian noise,  $\mu_R + \sigma_R \epsilon$  with  $\epsilon \sim \mathcal{N}(0, 1)$  i.i.d. per voxel and  $\sigma_R$  drawn uniformly from  $[0.05, 0.25]$  per sub-region; the label step (step 4) is noise-filled the same way. The par-

| Task                                                                    | Target                  | Trained on (one model per modality) | Tested on (every model; <sup>†</sup> from an external cohort)            |
|-------------------------------------------------------------------------|-------------------------|-------------------------------------|--------------------------------------------------------------------------|
| <b>Appearance-defined targets:</b> abnormal signal inside a tissue      |                         |                                     |                                                                          |
| Glioma                                                                  | glioma, 4 sub-regions   | T1n   T1c   T2w   FLAIR             | T1n, T1c, T2w, FLAIR                                                     |
| MS                                                                      | MS lesions              | FLAIR   T1w                         | FLAIR, T1w, T2w                                                          |
| Breast                                                                  | tumor                   | T1-CE   T2w                         | T1-CE <sup>(†)</sup> , T2w, pre-contrast <sup>†</sup> , DWI <sup>†</sup> |
| <b>Interface-bounded targets:</b> a structure bounded by another tissue |                         |                                     |                                                                          |
| Abdomen                                                                 | liver, kidneys, spleen  | T1in   T2spir                       | T1in, T1out, T2spir, T2w <sup>†</sup> , CT <sup>(†)</sup>                |
| Brain                                                                   | 31 brain regions        | T1w   T2w   DWI                     | T1w, T2w, DWI, BOLD, EPI, GRE                                            |
| Mandible                                                                | mandible                | CBCT                                | CBCT, CT <sup>†</sup> , MR-T1 <sup>†</sup>                               |
| Spine                                                                   | vertebrae, discs, canal | CT   Dixon in-phase                 | CT, Dixon in-phase, fat, water, T1w <sup>†</sup> , T2w <sup>†</sup>      |
| Pelvis                                                                  | 11 bones and muscles    | CT   MRI                            | CT, MRI                                                                  |

Table 1. The eight tasks and eighteen settings. Each training modality trains its own models, which are tested on every contrast listed for the task, the training contrast included (<sup>(†)</sup>: from the training cohort and an external cohort). All tasks but Spine have an ablation ladder. Datasets and case counts per contrast: Tab. 5.




Figure 2. The two arms of the fill swap on the brain slice of Fig. 1, with the same partition and region target means: Gaussian noise (left) or the remapped real intensities, PALETTE’s output (right). Insets: the boxed region magnified.

tion and the target means are PALETTE-Aug’s own, so a difference between the two isolates whether the real intensity structure inside each region is kept or replaced by noise (Fig. 2).

We report this comparison separately for *interface-bounded* targets (a normal structure bounded by another tissue class) and *appearance-defined* targets (a pathology identified by abnormal signal within a tissue), assigned by anatomy; the texture hypothesis predicts a larger gain on the latter.

## 4. Experiments

### 4.1. Setup

**Protocol.** Each model trains on one real source modality and is tested on every modality of the same task, the protocol Auglab uses for spine [40], here over eight tasks and with external cohorts wherever one exists. No target-domain data is seen in training. All methods share one nnU-Net [27] configuration and schedule and differ only in input augmentation; the epoch budget is fixed per dataset (Sec. 7). Our full method adds the PALETTE-Aug transform as one operator to the Auglab pipeline [40], so it and Auglab differ by exactly that transform. Each method is trained on

three train/validation folds, and all three models are scored on the same held-out test set. Significance tests are paired, use the patient (for Brain, the session) as the unit, and are Holm-corrected [25] (Sec. 9).

**Datasets.** Table 1 lists the eight tasks, the modalities each is trained on and the contrasts it is tested on, across MRI, CT and dental CBCT. Every (task, training modality) pair is one *setting*, eighteen in all; each setting trains its own models and is tested on every contrast of its task, the training contrast included. *Out-of-domain* (OOD) results exclude the training contrast, also in external cohorts. The training cohorts are, in table order, BraTS2024-GLI [11], Open-MS [35], I-SPY2 [38], CHAOS [31], ON-Harmony [58], ToothFairy2 [6], a CT/Dixon spine cohort and TotalSegmentator [10, 59]. The external test cohorts are Duke [14, 47], I-SPY1 [26] and ACRIN-6698 [43] (Breast), AMOS [28] and SLIVER07 [23] (Abdomen), HaN-Seg [44] and PDDCA [46] (Mandible), and SPIDER [56] (Spine); they are extra held-out test columns of their task, not separate tasks (sources per contrast: Tab. 5). ON-Harmony’s reference labels are silver-standard FreeSurfer [13] parcellations of the T1w, registered to the other contrasts. External volumes that image more of the body than the training cohort are cropped to its field of view before prediction, identically for every method (Sec. 8).

**Baselines.** Table 2 compares the five competitors: nnU-Net with its default augmentation (*Baseline*), SynthSeg [3] as released, developed for the brain (*SynthSeg-noEM*), and its EM label-subdivision variant, developed for cardiac MRI [3, §5.4] and re-implemented here since its code was not released (*SynthSeg-EM*), both given each task’s training labels, which are a dense parcellation only on Brain; SRCSM [54]; and Auglab [40], a competitor in its own right. Each competitor is run as its authors propose: the

| Method             | Intensities        | Varied per   | No dense labels | Real texture |
|--------------------|--------------------|--------------|-----------------|--------------|
| Baseline           | real               | —            | ✓               | ✓            |
| SynthSeg-noEM      | Gaussian           | label        | ×               | ×            |
| SynthSeg-EM        | Gaussian           | EM sub-label | ✓               | ×            |
| SRCSM              | random convolution | label        | ✓               | ~            |
| Auglab             | real, generic aug. | image        | ✓               | ✓            |
| <b>PALETTE-Aug</b> | real, affine remap | sub-region   | ✓               | ✓            |

Table 2. What each method trains on. Sub-region: unsupervised  $k$ -means and Voronoi partition (Sec. 3). No dense labels: works from the task’s own labels (SynthSeg-noEM runs on them but fills everything unlabeled with one Gaussian). ~: the real image passes through a random filter. Augmentation probabilities: Tab. 4.

SynthSeg variants and SRCSM synthesize on every training iteration with their own augmentation recipe, and Auglab uses its default configuration (Tab. 4). SRCSM’s published pipeline also histogram-matches each test image to the average source histogram, part of the method as published, so we apply it. On sparse task labels, SynthSeg-noEM (designed for dense labels) is a reference point for synthesizing from labels alone, not a tuned competitor. Our method adds one operator to Auglab, so comparing the two tests that operator; the other competitors run without the Auglab pipeline (Sec. 5).

## 4.2. Main comparison

**PALETTE-Aug is the best method overall.** PALETTE-Aug has the best Dice on seven of the eight tasks. Across tasks it is significantly better than every competitor (Tab. 3):  $p = 1.1 \times 10^{-13}$  against Auglab, the closest, and  $p \leq 10^{-35}$  against the others. Within single tasks it significantly outperforms every competitor on five tasks (Glioma, MS, Breast, Spine, Pelvis), all but SRCSM on Abdomen (85.6 vs. 85.9), and three of five on Brain and Mandible. The cross-task test generalizes to new patients on these eight tasks, not to new tasks. Counting each task as one unit instead, PALETTE-Aug beats Auglab on seven of eight tasks (one-sided sign test  $p = 0.035$ ) and SRCSM on six of the seven tasks where SRCSM was run ( $p = 0.063$ ).

SynthSeg-noEM falls below the baseline on the six sparse-label tasks; on Brain, the one dense-label task, it is within 2.6 Dice of PALETTE-Aug.

**No competitor is reliably second.** The best competitor alternates between Auglab (Glioma, MS, Breast, Spine, Pelvis) and SRCSM (Abdomen, Brain, Mandible), and SRCSM fails on Breast (14.7 Dice). Picking the right one would require knowing the answer in advance. PALETTE-Aug is first on every task but Mandible.

**Where PALETTE-Aug does not win.** The one-sided test above cannot show that a competitor is ahead, so on Mandible we also ran the reverse test on the same patient-level design (Holm-corrected across competitors within the task). There Auglab is significantly ahead (+0.6 Dice,  $p = 3 \times 10^{-3}$ ), SRCSM leads by 2.1 but not significantly ( $p = 0.15$ ), and PALETTE-Aug is third; PALETTE alone matches Auglab (78.3).

**PALETTE without Auglab.** The transform alone (the ladder’s real-fill rung) is as strong as the full method: over the seven tasks with both, it is significantly better than every competitor, Auglab included (Dice  $p = 4 \times 10^{-5}$ ), and does not differ from PALETTE-Aug (two-sided  $p = 0.60$ ). PALETTE-Aug’s margin over the competitors therefore does not come from the Auglab pipeline. Adding Auglab raises task Dice on Glioma and Breast, lowers it on MS (significantly), Abdomen and Mandible and leaves Pelvis unchanged (Tab. 3); we have not identified which Auglab operator is responsible.

**Boundary distance.** On HD95 (Tab. 9), PALETTE-Aug is lowest on six of the eight tasks and significantly better than every competitor across tasks; it ties Auglab on Breast and trails Auglab and the baseline on Mandible.

## 4.3. Causal ablation: when does texture matter?

Does keeping the real texture help beyond PALETTE-Aug’s partition, and on which targets? We test this with a ladder: baseline  $\rightarrow +k$ -means  $\rightarrow +$ label remap  $\rightarrow +$ Voronoi  $\rightarrow$  real-intensity fill  $\rightarrow +$ Auglab. The baseline is nnU-Net with its default augmentation. The next four rungs share one pipeline (Sec. 3) whose only other augmentation is spatial (with flips), so the first step also drops nnU-Net’s intensity augmentation; from there on, each rung adds one ingredient. The three partition rungs fill every region with SynthSeg-style Gaussian noise. The step from +Voronoi to real fill, the *fill swap*, changes only the fill: partition, sub-region count, target-mean distribution and both blur passes are unchanged. Any change at this step therefore comes from keeping the real within-region intensity structure instead of noise.

In our data every large gain from real fill is on an appearance-defined target, but not every appearance-defined target shows one (Fig. 3 and Tab. 10). The swap is tested once per task, pooling its training modalities (Sec. 9); the per-setting values are a breakdown of that test. Pooled per task, the swap adds +6.8 Dice on MS (all eight test patients improve) and +4.6 on Glioma, against +3.5 on Pelvis, +2.6 on Abdomen, +0.1 on Mandible and -0.7 on Brain; each of these tasks passes the correction, Mandible with a negligible gain and Brain with a loss. The four largest per-setting gains are Glioma T1c (+12.7 Dice), MS

|                                   | Glioma      | MS          | Breast      | Abdomen     | Brain       | Mandible    | Spine       | Pelvis      | $p$ vs ours           |
|-----------------------------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-----------------------|
| Baseline                          | 23.4        | 21.0        | 35.0        | 29.5        | 34.5        | 59.2        | 48.5        | 29.8        | $3 \times 10^{-75}$   |
| SynthSeg-noEM                     | 8.8         | 4.3         | 16.5        | 60.9        | 64.1        | 53.6        | 47.5        | 15.0        | $2 \times 10^{-86}$   |
| SynthSeg-EM                       | 43.1        | 33.9        | 39.6        | 84.0        | 66.1        | 76.3        | 80.6        | 62.5        | $1 \times 10^{-35}$   |
| SRCSM                             | 36.9        | 27.5        | 14.7        | 85.6        | 66.5        | <b>79.8</b> | —           | 62.9        | $8 \times 10^{-66}$   |
| Auglab                            | 44.8        | 34.3        | 40.1        | 84.6        | 66.0        | 78.3        | 80.9        | 64.3        | $1.1 \times 10^{-13}$ |
| PALETTE alone (ours)              | 43.7        | 41.1        | 40.3        | 86.4        | 67.0        | 78.3        | —           | 65.7        | 0.70 <sup>§</sup>     |
| <b>PALETTE-Aug (ours)</b>         | <b>45.9</b> | <b>38.5</b> | <b>41.2</b> | <b>85.9</b> | <b>66.7</b> | 77.7        | <b>82.3</b> | <b>65.6</b> | —                     |
| competitors beaten ( $p < 0.05$ ) | 5/5         | 5/5         | 5/5         | 4/5         | 3/5         | 3/5         | 4/4         | 5/5         |                       |

Table 3. Task-level Dice (%). Each column pools the task’s test cases of every contrast (the training contrast included), all source modalities and any external cohorts, weighting contrast groups equally so a modality with several external sources does not outvote one with fewer (Sec. 9).  $p$ : Holm-corrected one-sided sign-flip test on the macro-averaged per-patient difference, stratified by task with tasks weighted equally. We print no cross-task mean Dice, since whole-liver and breast-lesion Dice are not on a common scale. The bottom row counts the competitors PALETTE-Aug beats within each task alone (same test). SRCSM includes its test-time source matching (Sec. 4.1). Spine has no SRCSM run (Tab. 5). PALETTE alone: the PALETTE transform without the Auglab pipeline (the ladder’s real-fill rung, Sec. 4.3), left out of the bolding; Spine has no ladder. <sup>§</sup>Same test, PALETTE-Aug against PALETTE alone, over the seven tasks with both. HD95 is in Tab. 9.

FLAIR (+7.7), Glioma T1n (+6.0) and MS T1w (+5.9); no interface-bounded setting gains more than +4.8 (Pelvis CT). Breast gains only +0.5, all from T2w training (+2.1; T1-CE training loses 1.2), and Glioma T2w and FLAIR show no gain (+1.1, −1.3). Relative to the noise-fill Dice (Fig. 3) the ranking at the top is unchanged.

ON-Harmony labels 31 brain regions in the FreeSurfer label set [13]. Nearly all are bounded by another tissue class (CSF against tissue, grey against white matter), so we group it with Abdomen as interface-bounded. At the task level it is the only task where the swap significantly hurts (−0.7 Dice), with sessions as units. It is also the only task labeling nearly the whole brain; with one such task, we cannot tell whether that is why.

**Does the fill swap change texture?** The ladder shows that the fill swap changes Dice, not that it changes texture. We check this without any segmentation, with the normalized-gradient-field (NGF) similarity [20] between each augmented volume and its source scan. NGF measures how well image gradients align: it is 1 when all edges and fine structure are kept and about 1/3 for isotropic noise. Figure 4 reports it for every method on the seven ablation tasks. The noise fill keeps only the partition edges (mean NGF 0.44); the real fill, on the same partition, raises it to 0.83, the highest of all methods. PALETTE alone keeps more texture than every competitor on six tasks and ties with Auglab on Brain (Tab. 18). SynthSeg keeps little (0.50 with EM, 0.37 without, near the noise floor). PALETTE-Aug keeps less (0.70): it applies PALETTE in half of its draws, on top of Auglab.

**Supporting analyses.** The supplement gives per-setting results (Sec. 10), every ladder rung, every training → test pair and the HD95 ablation (Sec. 11), and a component ablation on all test contrasts (Sec. 13).

## 5. Discussion and Limitations

**A possible reason for the largest effects.** Liver, spleen and kidney are encapsulated organs with a clear physical boundary, and edge-following methods have long segmented them well [52]. An augmentation that scrambles their internal texture but keeps the boundary loses little. Diffuse glioma has no such boundary: tumor cells infiltrate beyond both the enhancing margin and the FLAIR abnormality [21, 53], so the target must be inferred from appearance. MS lesions, though typically well demarcated, are identified by abnormal signal within white matter, and texture features separate them from normal-appearing white matter [62]. This account interprets the intervention of Sec. 4.3; it is not independent evidence.

**Not every appearance-defined target gains much.** Glioma T2w and FLAIR show no Dice gain, and Breast gains only slightly: +2.1 Dice from T2w training but −1.2 from T1-CE. We have two hypotheses. First, the effect depends sharply on the training contrast (Glioma: +12.7 from T1c, +6.0 from T1n, +1.1 and −1.3 from T2w and FLAIR), so whether its texture transfers may matter as much as the target; with T2w training, real fill helps some labels and hurts others (Tab. 15). Second, Breast pools four cohorts from different trials, which may dilute a cohort-specific effect. On interface-bounded tasks, where texture adds little relative to the starting Dice, most of the gain arrives with the texture-free partition rungs (Tab. 12),




Figure 3. The fill swap (shaded band), one panel per task with a ladder, grouped by target type. Thin lines: one model, i.e. one training modality tested on one held-out contrast (external cohorts included), colored at the swap by that pair’s significance; values per pair: Tab. 13. Bold line: the task average. Annotation: the task’s Dice change at the swap in points, and relative to its noise-fill Dice in brackets; \*: significant (paired Wilcoxon on patients pooling all training modalities, Holm-corrected across the seven tasks: the one correction family of the fill swap; Brain’s eight sessions come from four subjects). Per setting: Tab. 10; HD95: Fig. 5.

that is, from the wide range of region-to-region contrasts PALETTE-Aug samples.

**Limitations.** We assign each target to a group by its anatomy, not by a measured boundary statistic, and the grouping is ours. Target type is confounded with difficulty: all interface-bounded rows but one start above 63% noise-




Figure 4. Texture kept by each augmentation: NGF similarity of the augmented volume to its source scan on the seven ablation tasks (violin: all scans, each averaged over 5 draws; dots: task means; bar: mean of the task means). The noise fill keeps only the partition edges; the real fill keeps the most.

fill Dice and every appearance-defined row below 42%. We therefore lead with the absolute Dice change; dividing by the starting Dice favors the appearance-defined rows, and as a share of the remaining error (100% minus noise-fill Dice) Abdomen T2spir would rank first. The exception, Pelvis CT, starts at 35% and gains +4.8, below the four large appearance-defined gains. The between-group comparison is descriptive: a permutation test on the seven per-task gains gives  $p = 0.2$ . MS has eight test patients and Brain eight sessions of four subjects, the floor of a paired test, and Brain’s session-level tests are optimistic. Spine was run on a separate pipeline and has no SRCSM arm. The cross-task test covers new patients, not new tasks; with the task as the unit, the advantage over SRCSM is not significant. We applied each competitor as its paper proposes and did not place SRCSM or SynthSeg inside the Auglab pipeline; PALETTE alone, without Auglab, already beats each of them (Sec. 4.2).

## 6. Conclusion

Preserving real image texture helps most where the target is defined by tissue appearance (brain tumor, MS lesions: +4.6 and +6.8 Dice per task) rather than by an anatomical interface (at most +3.5), though not on every appearance-defined target. PALETTE-Aug, which preserves texture, needs no dense label map and was never tuned to a task, wins the eight-task benchmark overall. All code is public (link withheld for review), as are all datasets but the Spine training cohort.

## References

- [1] Michela Antonelli, Annika Reinke, Spyridon Bakas, Keyvan Farahani, Annette Kopp-Schneider, Bennett A. Landman, Geert Litjens, Bjoern Menze, Olaf Ronneberger, Ronald M. Summers, et al. The Medical Segmentation Decathlon. *Nature Communications*, 13(1):4128, 2022. 4
- [2] Jie Bao, Zhixin Zhou, Wen Jung Li, and Rui Luo. Structure-aware stylized image synthesis for robust medical image segmentation. *arXiv preprint arXiv:2412.04296*, 2024. 2
- [3] Benjamin Billot, Douglas N. Greve, Oula Puonti, Axel Thielscher, Koen Van Leemput, Bruce Fischl, Adrian V. Dalca, and Juan Eugenio Iglesias. SynthSeg: Segmentation of brain MRI scans of any contrast and resolution without retraining. *Medical Image Analysis*, 86:102789, 2023. 1, 2, 3, 5
- [4] Benjamin Billot, Colin Magdamo, You Cheng, Steven E. Arnold, Sudeshna Das, and Juan Eugenio Iglesias. Robust machine learning segmentation for large-scale analysis of heterogeneous clinical brain MRI datasets. *Proceedings of the National Academy of Sciences*, 120(9):e2216399120, 2023. 4
- [5] Richard Bitar, General Leung, Richard Perng, Sameh Tadros, Alan R. Moody, Josee Sarrazin, Caitlin McGregor, Monique Christakis, Sean Symons, Andrew Nelson, and Timothy P. Roberts. MR pulse sequences: What every radiologist wants to know but is afraid to ask. *RadioGraphics*, 26(2):513–537, 2006. 1
- [6] Federico Bolelli, Luca Lumetti, et al. Multi-structure segmentation in CBCT volumes: The ToothFairy2 challenge. *Medical Image Analysis*, 112:104095, 2026. 5, 3
- [7] Vicent Caselles-Ballester, Eloy Martínez-Heras, Giuseppe Pontillo, Zoe Mendelsohn, Elena M. Marrón, Juan Luis García Fernández, Laia Subirats, Jon Stutters, Jeremy Chataway, Frederik Barkhof, Sara Llufrú, and Ferran Prados. TimeLesSeg: Unified contrast-agnostic cross-sectional and longitudinal MS lesion segmentation via a stochastic generative model. *arXiv preprint arXiv:2605.07955*, 2026. 2
- [8] Stefano Cerri, Oula Puonti, Dominik S. Meier, Jens Wuerfel, Mark Muehlau, Hartwig R. Siebner, and Koen Van Leemput. A contrast-adaptive method for simultaneous whole-brain and lesion segmentation in multiple sclerosis. *NeuroImage*, 225:117471, 2021. 2
- [9] Boqi Chen, Yuanzhi Zhu, Yunke Ao, Sebastiano Caprara, Reto Sutter, Gunnar Räscher, Ender Konukoglu, and Anna Susmelj. Generalizable single-source cross-modality medical image segmentation via invariant causal mechanisms. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 3592–3602, 2025. 4
- [10] Tugba Akinci D’Antonoli, Lucas K. Berger, Ashraya K. Indrakanti, et al. TotalSegmentator MRI: Robust sequence-independent segmentation of multiple anatomic structures in MRI. *Radiology*, 314(2):e241613, 2025. 5, 3
- [11] Maria Correia de Verdier, Rachit Saluja, Louis Gagnon, Dominic LaBella, Ujjwal Baid, Nourel Hoda Tahon, Martha Foltyn-Dumitru, Jikai Zhang, Ahmed W. Moawad, et al. The 2024 brain tumor segmentation (BraTS) challenge: Glioma segmentation on post-treatment MRI. *arXiv preprint arXiv:2405.18368*, 2024. 5, 2, 3
- [12] Michael D. Ernst. Permutation methods: A basis for exact inference. *Statistical Science*, 19(4):676–685, 2004. 5
- [13] Bruce Fischl. FreeSurfer. *NeuroImage*, 62(2):774–781, 2012. 5, 7, 2
- [14] Lidia Garrucho, Kaisar Kushibar, Claire-Anne Reidel, Smriti Joshi, Richard Osuala, et al. A large-scale multicenter breast cancer DCE-MRI benchmark dataset with expert segmentations. *Scientific Data*, 12(1):453, 2025. 5, 2, 3
- [15] Robert Geirhos, Patricia Rubisch, Claudio Michaelis, Matthias Bethge, Felix A. Wichmann, and Wieland Brendel. ImageNet-trained CNNs are biased towards texture; increasing shape bias improves accuracy and robustness. In *International Conference on Learning Representations (ICLR)*, 2019. 3
- [16] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020. 9
- [17] Mohsen Ghafoorian, Alireza Mehrtash, Tina Kapur, Nico Karssemeijer, Elena Marchiori, Mehran Pesteie, Charles R. G. Guttmann, Frank-Erik de Leeuw, Clare M. Tempny, Bram van Ginneken, Andriy Fedorov, Purang Abolmaesumi, Bram Platel, and William M. Wells. Transfer learning for domain adaptation in MRI: Application in brain lesion segmentation. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, pages 516–524, 2017. 1
- [18] Hao Guan and Mingxia Liu. Domain adaptation for medical image analysis: A survey. *IEEE Transactions on Biomedical Engineering*, 69(3):1173–1185, 2022. 1
- [19] Qiu Guan, Yizhou Chen, Zihan Wei, Ali Asghar Heidari, Haigen Hu, Xu-Hua Yang, Jianwei Zheng, Qianwei Zhou, Huiling Chen, and Feng Chen. Medical image augmentation for lesion detection using a texture-constrained multichannel progressive GAN. *Computers in Biology and Medicine*, 145:105444, 2022. 3
- [20] Eldad Haber and Jan Modersitzki. Intensity gradient based registration and fusion of multi-modal images. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 726–733, 2006. 7
- [21] Alexander F. Haddad, Jacob S. Young, Ramin A. Morshed, and Mitchel S. Berger. FLAIRctomy: Resecting beyond the contrast margin for glioblastoma. *Brain Sciences*, 12(5):544, 2022. 2, 7
- [22] Robert M. Haralick, K. Shanmugam, and Its’hak Dinstein. Textural features for image classification. *IEEE Transactions on Systems, Man, and Cybernetics*, SMC-3(6):610–621, 1973. 3
- [23] Tobias Heimann, Bram van Ginneken, Martin A. Styner, Yulia Arzhaeva, Volker Aurich, Christian Bauer, Andreas Beck, Christoph Becker, Reinhard Beichel, György Bekes, et al. Comparison and evaluation of methods for liver segmentation from CT datasets. *IEEE Transactions on Medical Imaging*, 28(8):1251–1265, 2009. 5, 2, 3
- [24] Katherine L. Hermann, Ting Chen, and Simon Kornblith. The origins and prevalence of texture bias in convolutional

- neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. 3
- [25] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979. 5, 4
- [26] Nola M. Hylton et al. Locally advanced breast cancer: MR imaging for prediction of response to neoadjuvant chemotherapy—results from ACRIN 6657/I-SPY TRIAL. *Radiology*, 263(3):663–672, 2012. 5, 2, 3
- [27] Fabian Isensee, Paul F. Jaeger, Simon A. A. Kohl, Jens Petersen, and Klaus H. Maier-Hein. nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18(2):203–211, 2021. 2, 3, 4, 5, 1
- [28] Yuanfeng Ji, Haotian Bai, Chongjian Ge, Jie Yang, Ye Zhu, Ruimao Zhang, Zhen Li, Lingyan Zhang, Wanling Ma, Xiang Wan, and Ping Luo. AMOS: A large-scale abdominal multi-organ benchmark for versatile medical image segmentation. In *Advances in Neural Information Processing Systems*, pages 36722–36732, 2022. 5, 2, 3
- [29] Myeongkyun Kang, Philip Chikontwe, Dongkyu Won, Miguel Luna, and Sang Hyun Park. Structure-preserving image translation for multi-source medical image domain adaptation. *Pattern Recognition*, 144:109840, 2023. 3
- [30] Neerav Karani, Krishna Chaitanya, Christian Baumgartner, and Ender Konukoglu. A lifelong learning approach to brain MR segmentation across scanners and protocols. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, pages 476–484, 2018. 1
- [31] A. Emre Kavur, N. Sinem Gezer, Mustafa Barış, Sinem Aslan, Pierre-Henri Conze, Vladimir Groza, Duc Duy Pham, Soumick Chatterjee, Philipp Ernst, Savaş Özkan, et al. CHAOS challenge – combined (CT-MR) healthy abdominal organ segmentation. *Medical Image Analysis*, 69:101950, 2021. 5, 2, 3
- [32] Frithjof Kruggel, Joseph Suresh Paul, and Hermann-Josef Gertz. Texture-based segmentation of diffuse lesions of the brain’s white matter. *NeuroImage*, 39(3):987–996, 2008. 3
- [33] Philippe Lambin, Emmanuel Rios-Velazquez, Ralph Leijenaa, Sara Carvalho, Ruben G. P. J. van Stiphout, Patrick Granton, Catharina M. L. Zegers, Robert Gillies, Ronald Boellard, Andre Dekker, and Hugo J. W. L. Aerts. Radiomics: Extracting more information from medical images using advanced feature analysis. *European Journal of Cancer*, 48(4):441–446, 2012. 3
- [34] Pablo Laso, Stefano Cerri, Annabel Sorby-Adams, Jennifer Guo, Farrah Mateen, Philipp Goebel, Jiaming Wu, Peirong Liu, Hongwei Li, Sean I. Young, Benjamin Billot, Oula Puonti, et al. Quantifying white matter hyperintensity and brain volumes in heterogeneous clinical and low-field portable MRI. In *IEEE International Symposium on Biomedical Imaging (ISBI)*, pages 1–5, 2024. 2
- [35] Žiga Lesjak, Alfiia Galimzianova, Aleš Koren, Matjaž Lukač, Koen Van Leemput, and Žiga Špiclin. A novel public MR image dataset of multiple sclerosis patients with lesion segmentations based on multi-rater consensus. *Neuroinformatics*, 16(1):51–63, 2018. 1, 5, 2, 3
- [36] Da Li, Yongxin Yang, Yi-Zhe Song, and Timothy M. Hospedales. Learning to generalize: Meta-learning for domain generalization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018. 2
- [37] Heng Li, Haojin Li, Jianyu Chen, Mingyang Ou, Hai Shu, and Heng Miao. RaffesDG: Random frequency filtering enabled single-source domain generalization for medical image segmentation. *arXiv preprint arXiv:2405.01228*, 2024. 4
- [38] Wen Li, David C. Newitt, Jessica Gibbs, Lisa J. Wilmes, Ella F. Jones, Vignesh A. Arasu, et al. I-SPY 2 breast dynamic contrast enhanced MRI trial (ISPY2). The Cancer Imaging Archive, 2022. Version 1, data set. 5, 2, 3
- [39] Quande Liu, Cheng Chen, Jing Qin, Qi Dou, and Pheng-Ann Heng. FedDG: Federated domain generalization on medical image segmentation via episodic learning in continuous frequency space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1013–1023, 2021. 1, 2
- [40] Nathan Molinier, Hendrik Möller, Thomas Dagonneau, Anna Curto-Vilalta, Robert Graf, Matan Atad, Daniel Rueckert, Jan S. Kirschke, and Julien Cohen-Adad. One sequence to segment them all: Efficient data augmentation for CT and MRI cross-domain 3D spine segmentation. *arXiv preprint arXiv:2605.03098*, 2026. 2, 3, 4, 5
- [41] Ju-Hyeon Nam, Nur Suriza Syazwany, Su Jung Kim, and Sang-Chul Lee. Modality-agnostic domain generalizable medical image segmentation by multi-frequency in multi-scale attention. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11480–11491, 2024. 2
- [42] Cheng Ouyang, Chen Chen, Surui Li, Zeju Li, Chen Qin, Wenjia Bai, and Daniel Rueckert. Causality-inspired single-source domain generalization for medical image segmentation. *IEEE Transactions on Medical Imaging*, 42(4):1095–1106, 2023. 1, 2, 3, 9
- [43] Savannah C. Partridge et al. Diffusion-weighted MRI findings predict pathologic response in neoadjuvant treatment of breast cancer: The ACRIN 6698 multicenter trial. *Radiology*, 289(3):618–627, 2018. 5, 2, 3
- [44] Gašper Podobnik, Primož Stojan, Primož Peterlin, Bulat Ibragimov, and Tomaž Vrtovec. HaN-Seg: The head and neck organ-at-risk CT and MR segmentation dataset. *Medical Physics*, 50(3):1917–1927, 2023. 5, 3
- [45] Qiang Qiao, Wenyu Wang, Meixia Qu, Kun Su, Bin Jiang, and Qiang Guo. Medical image segmentation via single-source domain generalization with random amplitude spectrum synthesis. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 435–445, 2024. 4
- [46] Patrik F. Raudaschl et al. Evaluation of segmentation methods on head and neck CT: Auto-segmentation challenge 2015. *Medical Physics*, 44(5):2020–2036, 2017. 5, 3
- [47] Ashirbani Saha, Michael R. Harowicz, Lars J. Grimm, Connie E. Kim, Sujata V. Ghatge, Ruth Walsh, and Maciej A. Mazurowski. A machine learning approach to radiogenomics of breast cancer: a study of 922 subjects and 529 DCE-MRI features. *British Journal of Cancer*, 119(4):508–516, 2018. 5, 2, 3

- [48] Sourya Sengupta, Satrajit Chakrabarty, Keerthi Sravan Ravi, Gopal Avinash, and Ravi Soni. SynthFM: Training modality-agnostic foundation models for medical image segmentation without real medical data. In *IEEE International Symposium on Biomedical Imaging (ISBI)*, pages 1–4, 2025. 2
- [49] Harshay Shah, Kaustav Tamuly, Aditi Raghunathan, Prateek Jain, and Praneeth Netrapalli. The pitfalls of simplicity bias in neural networks. In *Advances in Neural Information Processing Systems*, 2020. 9
- [50] Hyungseob Shin, Hyeongyu Kim, Sewon Kim, Yohan Jun, Taejoon Eo, and Dosik Hwang. SDC-UDA: Volumetric unsupervised domain adaptation framework for slice-direction continuous cross-modality medical image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7412–7421, 2023. 2
- [51] Zixian Su, Kai Yao, Xi Yang, Qiufeng Wang, Jie Sun, and Kaizhu Huang. Rethinking data augmentation for single-source domain generalization in medical image segmentation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 2366–2374, 2023. 1, 4
- [52] Kenji Suzuki, Hieu Trung Huynh, Yipeng Liu, Dominic Calabrese, Karen Zhou, Aytekin Oto, and Masatoshi Hori. Computerized segmentation of liver in hepatic CT and MRI by means of level-set geodesic active contouring. In *Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 2984–2987, 2013. 7
- [53] Jun Tang, Liangpeng Chen, Ziyang Li, Cong Li, and Deling Li. Imaging strategies for delineating glioma margins: Aiming at precise therapy. *Fundamental Research*, 6(4):2644–2653, 2026. 2, 7
- [54] Franz Thaler, Martin Urschler, Mateusz Koziński, Matthias A. F. Gsell, Gernot Plank, and Darko Štern. Semantic-aware random convolution and source matching for domain generalization in medical image segmentation. *IEEE Access*, 14: 64338–64356, 2026. 2, 4, 5
- [55] Romain Valabregue, Ines Khemir, Eric Badinet, François Rousseau, Guillaume Auzias, and Reuben Dorent. SIAM: Head and brain MRI segmentation from few high-quality templates via synthetic training. *arXiv preprint arXiv:2605.02737*, 2026. 2
- [56] Jasper W. van der Graaf, Miranda L. van Hooff, Constantinus F. M. Buckens, Matthieu Rutten, Job L. C. van Susante, Robert Jan Kroeze, Marinus de Kleuver, Bram van Ginneken, and Nikolas Lessmann. Lumbar spine segmentation in MR images: a dataset and a public benchmark. *Scientific Data*, 11:264, 2024. 5, 3
- [57] Shujun Wang, Lequan Yu, Kang Li, Xin Yang, Chi-Wing Fu, and Pheng-Ann Heng. DoFE: Domain-oriented feature embedding for generalizable fundus image segmentation on unseen datasets. *IEEE Transactions on Medical Imaging*, 39(12):4237–4248, 2020. 1, 2
- [58] Shaun Warrington, Andrea Torchi, Olivier Mougin, Jon Campbell, Asante Ntata, Martin Craig, Stephania Assimpopoulos, Fidel Alfaro-Almagro, Karla L. Miller, Mark Jenkinson, Paul S. Morgan, and Stamatios N. Sotiropoulos. A multi-site, multi-modal travelling-heads resource for brain MRI harmonisation. *Scientific Data*, 12:609, 2025. 5, 2, 3
- [59] Jakob Wasserthal et al. TotalSegmentator: Robust segmentation of 104 anatomic structures in CT images. *Radiology: Artificial Intelligence*, 5(5):e230024, 2023. 5, 3
- [60] Frank Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945. 4
- [61] Zhenlin Xu, Deyi Liu, Junlin Yang, Colin Raffel, and Marc Niethammer. Robust and generalizable visual representation learning via random convolutions. In *International Conference on Learning Representations (ICLR)*, 2021. 1, 4
- [62] Jing Zhang, Longzheng Tong, Lei Wang, and Ning Li. Texture analysis of multiple sclerosis: a comparative study. *Magnetic Resonance Imaging*, 26(8):1160–1166, 2008. 2, 7
- [63] Ling Zhang, Xiaosong Wang, Dong Yang, Thomas Sanford, Stephanie Harmon, Baris Turkbey, Bradford J. Wood, Holger Roth, Andriy Myronenko, Daguang Xu, and Ziyue Xu. Generalizing deep learning for medical image segmentation to unseen domains via deep stacked transformation. *IEEE Transactions on Medical Imaging*, 39(7):2531–2540, 2020. 1, 2, 3
- [64] Haoyu Zhao, Wenhui Dong, Rui Yu, Zhou Zhao, Bo Du, and Yongchao Xu. MoreStyle: Relax low-frequency constraint of Fourier-based image reconstruction in generalizable medical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 434–444, 2024. 4

# PALETTE-Aug: A Simple Texture-Preserving Augmentation for Medical Imaging Domain Generalization

## Supplementary Material

### 874 7. Training configuration

875 All methods share one nnU-Net [27] configuration; only the  
876 input augmentation differs. Nothing below is tuned per task.

877 **Training schedule.** 3D full-resolution nnU-Net, SGD  
878 with Nesterov momentum 0.99, initial learning rate  $10^{-2}$   
879 with the nnU-Net `polyLR` decay, 250 mini-batches per  
880 epoch, combined Dice and cross-entropy loss, deep supervision on. Patch size, spacing and batch size are whatever the  
881 nnU-Net planner selects for that dataset, identically for every  
882 method. The epoch budget is fixed per task and shared  
883 by all methods on it: 2500 (Glioma), 2000 (Brain, MS),  
884 1000 (Mandible, Breast) and 200 (Abdomen, Pelvis); Spine  
885 was trained on a separate pipeline. Each method is trained  
886 as three models on three train/validation folds, all scored on  
887 the same held-out test set; folds are the same across methods.  
888  
889

890 **PALETTE-Aug hyper-parameters.** Set once and never  
891 tuned per dataset:  $k$ -means class count drawn per sample  
892 from  $\{2, \dots, 6\}$ , skipped in favor of a single global  
893 remap with probability 0.10; Voronoi sub-parcellation of  
894 each intensity class with probability 0.60 (else left as one  
895 region); signed affine coefficient  $\alpha \in \pm U(0.5, 2.0)$  applied  
896 per sub-region as in Eq. (1) (always applied, to whatever  
897 partition results from the  $k$ -means/Voronoi steps); per-label  
898 affine decoupling applied independently per label with  
899 probability 0.5 (skipped for any label with under 4 foreground  
900 voxels in that sample); Gaussian blur  $\sigma$  drawn from  
901  $\{0, 0, 0, 0.3, 0.5, 0.8\}$ , applied twice (after the region-level  
902 and after the label-level remap), so applied half the time at  
903 each point. The transform is applied on the GPU, on-the-fly,  
904 to a fixed 50% of training iterations; the remaining half  
905 of iterations use the un-transformed real image.

906 **Validation-time configuration.** Every reported model,  
907 every competitor and every ladder rung included, selects its  
908 best checkpoint on real-contrast validation images (*val000*).  
909 *val100* applies PALETTE-Aug to the entire validation split  
910 instead; it is reported only in Sec. 12.

911 **Competitor configurations.** Both SynthSeg variants run  
912 through Auglab's GPU implementation of the SynthSeg  
913 generator, with uniform intensity priors and with the generator's  
914 own spatial and flip steps disabled (the pipeline's spatial  
915 augmentation is shared by all methods). They receive

each task's training label map; SynthSeg-EM additionally  
subdivides it on the fly by expectation-maximization clustering  
of the image (2 sub-regions per foreground label, 3–10 for the  
background, 20 iterations), where the original work precomputes  
these maps [3, §5.4]. Each sub-region then receives a randomly  
sampled intensity, as in SynthSeg. Both SynthSeg variants and  
SRCSM synthesize on every training iteration because their original  
papers train that way, so we keep that setting when running them;  
PALETTE-Aug synthesizes on 50% of iterations (Tab. 4 lists every  
augmentation's probability). SRCSM is run as published, including  
its test-time source matching: each test image is histogram-matched  
to the average intensity histogram of the training images (re-implemented  
from the authors' code), with two adjustments for our data: MR  
intensities are clipped at 0 first (resampling leaves small negative  
values that would otherwise map the background to a tissue intensity)  
and the CT window is  $[-1024, 3071]$  instead of  $[-1023, 1024]$  (CBCT  
and bone exceed 1024). Held-out cases are never part of the source  
histogram. Every method is predicted with the same nnU-Net predictor,  
without test-time augmentation. SRCSM and Auglab otherwise use  
their published default configurations. No competing method was  
re-tuned in our favor, and all use the same folds, schedule and epoch  
budget as above.

**Compute.** Training was performed on NVIDIA L40S and H100 GPUs.  
A single fold costs from under an hour (Abdomen, 200 epochs) to  
roughly three days (Glioma, 2500 epochs); the full study is dominated  
by the six methods  $\times$  sixteen settings  $\times$  three folds of the main  
comparison (Spine's two settings were trained on a separate pipeline,  
without SRCSM), plus the sixteen ablation ladders of Sec. 4.3.

|           | Transform                           | Baseline         | SynthSeg-noEM | SynthSeg-EM | SRCSM | Auglab      | PALETTE-Aug |
|-----------|-------------------------------------|------------------|---------------|-------------|-------|-------------|-------------|
| Synthesis | SynthSeg generator                  | –                | 1.0           | 1.0         | –     | –           | –           |
|           | semantic random convolution         | –                | –             | –           | 1.0   | –           | –           |
|           | PALETTE transform                   | –                | –             | –           | –     | –           | 0.5         |
| Intensity | Gaussian noise                      | 0.1              | 0.1           | 0.1         | –     | 0.1         | 0.1         |
|           | Gaussian blur                       | 0.2              | 0.2           | 0.2         | –     | 0.2         | 0.2         |
|           | brightness                          | 0.15             | 0.15          | 0.15        | –     | 0.15        | 0.15        |
|           | contrast                            | 0.15             | 0.15          | 0.15        | –     | 0.15        | 0.15        |
|           | gamma / inverted gamma              | 0.3 / 0.1        | 0.3 / 0.1     | 0.3 / 0.1   | –     | 0.3 / 0.1   | 0.3 / 0.1   |
|           | simulated low resolution            | 0.25             | 0.25          | 0.25        | –     | 0.25        | 0.25        |
|           | bias field                          | –                | 0.15          | 0.15        | –     | 0.15        | 0.15        |
|           | random intensity function           | –                | 0.1           | 0.1         | –     | 0.1         | 0.1         |
|           | Scharr / unsharp mask               | –                | –             | –           | –     | 0.25 / 0.25 | 0.25 / 0.25 |
|           | random convolution                  | –                | –             | –           | –     | 0.2         | 0.2         |
|           | label-wise intensity redistribution | –                | –             | –           | –     | 0.4         | 0.4         |
|           | intensity inversion                 | –                | –             | –           | –     | 0.25        | 0.25        |
|           | histogram equalization              | –                | –             | –           | –     | 0.25        | 0.25        |
| Spatial   | rotation / scaling                  | 0.2 / 0.2        | 0.2 / 0.2     | 0.2 / 0.2   | –     | 0.2 / 0.2   | 0.2 / 0.2   |
|           | affine + elastic deformation        | –                | –             | –           | 1.0   | –           | –           |
|           | flip                                | 0.5 <sup>a</sup> | 0.5           | 0.5         | –     | 0.5         | 0.5         |

Table 4. Per-iteration probability of each augmentation, per method (–: not applied). The synthesis rows are each method’s own contrast step: the SynthSeg variants and SRCSM synthesize on every iteration, as in their original papers, PALETTE-Aug on half of them. The remaining rows are the method’s generic augmentation pipeline: nnU-Net’s defaults for the baseline, Auglab’s for the Auglab-based methods (the SynthSeg variants switch off Auglab’s appearance-randomizing operators, which would otherwise act on top of the synthetic image), and SRCSM’s own spatial recipe. <sup>a</sup>nnU-Net mirroring, each axis independently.

## 8. Datasets and evaluation protocol

All evaluation is performed on held-out test cases only. Each model is trained on one modality of its task and evaluated on every test contrast of that task, the training contrast included; training and validation subjects never appear in any reported number.

**Glioma (BraTS2024-GLI [11]).** Post-treatment glioma sub-region segmentation (necrotic core, surrounding non-enhancing FLAIR hyperintensity (SNFH), enhancing tumor, resection cavity), trained on each of T1n, T1c, T2w and FLAIR and evaluated on all four.

**MS (Open-MS [35]).** MS-lesion segmentation with manual consensus annotations, evaluated across FLAIR, T1w and T2w.

**Breast (I-SPY2; external: Duke, I-SPY1, ACRIN-6698 [14, 26, 38, 43, 47]).** Breast-lesion segmentation on malignant-only cohorts. I-SPY2 trains, on its first-post-contrast T1 (T1-CE) and its T2w; its reference is the trial’s functional tumor volume mask. Three external cohorts are test-only. Duke and I-SPY1 come through the MAMA-MIA release, whose expert primary-lesion protocol both share, and contribute T1-CE and pre-contrast columns. ACRIN-6698, the diffusion sub-study of I-SPY2 with no patient in common with our I-SPY2 subjects, contributes the baseline

diffusion-weighted acquisition with the trial’s whole-tumor DWI region, a contrast neither training modality resembles. Because I-SPY2 is mostly acquired unilaterally, bilateral external volumes are cropped to the lesion side and, for the full-thorax acquisitions, also anterior-posteriorly to the breast-plus-chest-wall extent of I-SPY2’s own crops; without the second crop every method’s HD95 is dominated by the thorax. I-SPY2’s DCE series are acquired unilaterally at some sites and bilaterally at others; both field-of-view variants are kept per patient, and the significance tests treat them as one patient.

**Abdomen (CHAOS [31]; external: AMOS [28], SLIVER07 [23]).** Abdominal organ segmentation (liver, spleen, both kidneys), evaluated on the CHAOS MR contrasts (four held-out patients) and CT (20 held-out patients), and on two external cohorts: CT (100) and T2w MRI (20) from AMOS and liver-only CT (20) from SLIVER07. In the pooled CT group, liver comes from CHAOS, AMOS and SLIVER07 and spleen and kidneys from AMOS only (CHAOS CT is annotated for the liver only). The full-torso external volumes are cropped to the axial field of view of CHAOS before prediction (field-of-view matching, below).

**Brain (ON-Harmony [58]).** 31-region healthy-brain segmentation. Reference labels are *silver standard*: FreeSurfer parcellations [13] of each session’s T1w, grouped into 31 regions and rigidly registered with ANTs into the native space

| Task                      | Target                  | Trained on (one modality at a time)             | Other test contrasts | External test cohorts                                                 | Test cases     | Ladder |
|---------------------------|-------------------------|-------------------------------------------------|----------------------|-----------------------------------------------------------------------|----------------|--------|
| <i>Appearance-defined</i> |                         |                                                 |                      |                                                                       |                |        |
| Glioma                    | glioma, 4 sub-regions   | BraTS2024-GLI [11]: T1n, T1c, T2w, FLAIR        | –                    | –                                                                     | 70             | ✓      |
| MS                        | MS lesions              | Open-MS [35]: FLAIR, T1w                        | T2w                  | –                                                                     | 8              | ✓      |
| Breast                    | tumor                   | I-SPY2 [38]: T1-CE, T2w                         | –                    | Duke [14, 47], I-SPY1 [26]: T1-CE, pre-contrast; ACRIN-6698 [43]: DWI | 84 + 829       | ✓      |
| <i>Interface-bounded</i>  |                         |                                                 |                      |                                                                       |                |        |
| Abdomen                   | liver, kidneys, spleen  | CHAOS [31]: T1in, T2spir                        | T1out, CT            | AMOS [28]: CT, T2w; SLIVER07 [23]: CT                                 | 24 + 140       | ✓      |
| Brain                     | 31 brain regions        | ON-Harmony [58]: T1w, T2w, DWI                  | BOLD, EPI, GRE       | –                                                                     | 8 <sup>†</sup> | ✓      |
| Mandible                  | mandible                | ToothFairy2 [6]: CBCT                           | –                    | HaN-Seg [44]: CT, MR-T1; PDDCA [46]: CT                               | 71 + 82        | ✓      |
| Spine                     | vertebrae, discs, canal | CT/Dixon cohort: CT, Dixon in-phase             | Dixon fat, water     | SPIDER [56]: T1w, T2w                                                 | 29 + 171/219   | –      |
| Pelvis                    | 11 bones and muscles    | TotalSegmentator [10, 59]: CT, MRI <sup>‡</sup> | –                    | –                                                                     | 30/27          | ✓      |

Table 5. Per-task dataset roster (summary in Tab. 1, main paper). Test cases: own held-out split + external cohorts; “ $x/y$ ” gives the counts of two contrasts where they differ (per dataset and contrast: Tab. 6). Ladder: the task has the causal ablation of Sec. 4.3. <sup>†</sup>Sessions of four subjects. <sup>‡</sup>Different patients for CT and MRI.

of the same session’s other contrasts. They are not manual annotations, so absolute Dice here is agreement with an automated reference; every method is scored against the same reference. The test set is eight sessions of four subjects, each scanned at two sites.

**Mandible (ToothFairy2; external: HaN-Seg, PDDCA [6, 44, 46]).** Mandible segmentation, trained on dental cone-beam CT, the one non-CT, non-MRI training modality here. ToothFairy2’s annotation coverage is cohort-dependent, so every column of this task is scored on the mandible alone (teeth excluded, matching the external references), which keeps the numbers being averaged commensurable. ToothFairy2 acquires CBCT only, so apart from its own 71 CBCT test cases its held-out contrasts all come from other datasets: HaN-Seg CT (42 cases) and MR-T1 (41) and PDDCA CT (40), the two CT sources pooled as one contrast group. ToothFairy2’s volumes are stored superior-inferior inverted relative to the external cohorts — an inversion invisible to the image header, found by checking anatomy directly — so the external test volumes are flipped into ToothFairy2’s convention before inference, which leaves Dice and HD95 unchanged under the rigid transform. The correction matters for the comparison: a method without flip augmentation (SRCSM) was penalized by the inversion far more than the others, and it is after this correction that SRCSM leads this task.

**Spine (CT/Dixon cohort; external: SPIDER [56]).** Vertebrae, intervertebral discs and spinal canal. Results for this task were computed on a separate pipeline with the same

method suite except SRCSM, trained on CT and Dixon in-phase MRI of a CT/Dixon cohort and tested on that cohort’s held-out CT and Dixon fat, in-phase and water images and on the T1w and T2w of SPIDER. Dice and HD95 are computed on each whole structure per label. No SRCSM arm was run, so Spine is absent from every SRCSM comparison.

**Pelvis (TotalSegmentator [10, 59]).** Eleven pelvic and hip musculoskeletal structures (both hips, sacrum, and the gluteus maximus, medius and minimus and iliopsoas on each side), from the TotalSegmentator CT and MRI releases. The two are unpaired cohorts, and the MRI release mixes sequences by design, so its MRI column is a sequence-heterogeneous test rather than one named contrast. CT is subsampled to 200 cases to keep training time comparable with the other tasks.

**Field-of-view matching.** A model tested on a field of view it never saw in training can fail for reasons that have nothing to do with contrast or modality. Where an external cohort images more of the body than the training cohort, we therefore crop its volumes before prediction, so that the shift left between training and test is, as far as possible, the contrast or modality alone. Abdomen (AMOS, SLIVER07): the axial slab CHAOS covers, placed with the median CHAOS margins around the kidneys (AMOS) or the liver (SLIVER07). Mandible (HaN-Seg, PDDCA): a  $144 \times 128 \times 80$  mm box centered on the mandible, close to the 90th percentile of ToothFairy2’s training extent, enlarged (at most to  $154 \times 154 \times 89$  mm) only when needed to contain the whole mandible. Breast (Duke, ACRIN-

6698): the lesion-side half of bilateral volumes, then, for full-thorax acquisitions, an anterior-posterior window from 27 mm in front of the skin to 147 mm behind it, the measured geometry of I-SPY2’s unilateral acquisitions, extended if needed to contain the lesion; I-SPY1 is natively unilateral and is not cropped. The network sees, and is scored on, only the cropped volume, and the crops are identical for every method.

| Task     | Dataset (role)               | Train            | Test cases               |
|----------|------------------------------|------------------|--------------------------|
| Glioma   | BraTS2024-GLI (train)        | 630              | 70                       |
| MS       | Open-MS (train)              | 22               | 8                        |
| Breast   | I-SPY2 (train)               | 476              | 84 patients <sup>a</sup> |
|          | Duke (external)              | –                | 291                      |
|          | I-SPY1 (external)            | –                | 167 <sup>b</sup>         |
|          | ACRIN-6698 (external)        | –                | 371                      |
| Abdomen  | CHAOS (train)                | 16               | 4 MR, 20 CT              |
|          | AMOS (external)              | –                | 100 CT, 20 T2w           |
|          | SLIVER07 (external)          | –                | 20 CT                    |
|          | ON-Harmony (train)           | 84 <sup>c</sup>  | 8 sessions               |
| Mandible | ToothFairy2 (train)          | 408              | 71                       |
|          | HaN-Seg (external)           | –                | 42 CT, 41 MR-T1          |
|          | PDDCA (external)             | –                | 40 CT                    |
| Spine    | CT/Dixon cohort (train)      | n/a <sup>d</sup> | 27–29 per contrast       |
|          | SPIDER (external)            | –                | 171 T1w, 219 T2w         |
| Pelvis   | TotalSegmentator CT (train)  | 170              | 30                       |
|          | TotalSegmentator MRI (train) | 156              | 27                       |

Table 6. Every dataset, by task. Train: subjects in the three train/validation folds (identical across folds and methods); external cohorts are test-only. Test cases per contrast where they differ. ON-Harmony counts are scan sessions (the eight test sessions are four subjects, each scanned at two sites). <sup>a</sup>One held-out patient can yield two T2w images (both field-of-view variants). <sup>b</sup>After excluding three bilateral cancers and one implant by a rule fixed before prediction. <sup>c</sup>76 for the T2w-trained models. <sup>d</sup>Trained on a separate pipeline. Pelvis MRI keeps 183 of the 298 released cases.

## 9. Statistical testing

**Test construction.** All tests are paired, and the unit of independence is the *patient*, never the fold, the image or the contrast. A patient’s score is its mean over labels and evaluated folds; where the same patient is scored in several held-out contrasts, or by the models of several training modalities, those scores move together as one unit under the null. Treating them as separate samples would be pseudo-replication and would inflate the effective sample size — a real risk here, since several tasks test a handful of patients under many contrasts (Brain: eight sessions of four subjects, six contrasts). For Brain the unit is the session, since its four subjects leave a test that could never reach significance; two sessions of one subject are not fully independent, so its *p*-values are, if anything, optimistic, and we state Brain-only results with this caveat. Per-setting comparisons (Sec. 10) use the same test within each setting. The ablation ladders use the paired Wilcoxon signed-rank

test [60], since there each task is judged on its own: each patient contributes one difference, its mean over the held-out contrasts and training modalities it appears in, so patients, not contrasts, are weighted equally. The fill swap has one correction family: it is tested once per task and Holm-corrected across the seven tasks (Fig. 3 and Tab. 10); its per-setting rows report the same test uncorrected, as a breakdown, and carry no significance marks. All other tests are Holm-corrected [25] for multiple comparisons.

**Why the aggregate tests are sign-flip tests on a macro-average.** The comparisons between methods (Tabs. 3 and 7 to 9) are not paired Wilcoxon signed-rank tests, which is the more usual choice in this literature: SynthSeg+, from the lineage we compare against, marks a method as better than its competitors using a one-sided Bonferroni-corrected Wilcoxon signed-rank test on paired Dice [4]. We therefore give our reason rather than leave it implicit. The *estimand* we test, by contrast, is not unusual at all — it is the same equal-weight-per-stratum principle the Medical Segmentation Decathlon adopted, for the same reason, as set out below.

*The estimand.* What we claim is  $\text{macro}\Delta$ : the mean, over test contrasts, of that contrast’s mean paired per-case difference (ours minus competitor) — equal weight per contrast, which is exactly what every task column of Tab. 3 and every setting column of Tab. 7 reports, so the test and the table measure the same thing. Equal weighting is not cosmetic here. The claim under test is cross-contrast generalization, and a contrast’s held-out case count reflects how large a public cohort happens to be, not how much that contrast should say about generalizing. Abdomen is the extreme case: its CT column pools 140 cases (20 from CHAOS, 100 from AMOS, 20 from SLIVER07), while CHAOS’s three MR contrasts share the same 4 patients and AMOS T2w has 20. Weighting by case count would let CT cast most of the votes in a test whose conclusion is about transfer across contrasts.

This is the same problem, and the same answer, as the Medical Segmentation Decathlon’s. Facing sample sizes that “varied heavily across all tasks and target ROIs,” it ranked algorithms within each target ROI, averaged those ranks per task, and averaged the per-task ranks for the final score — choosing that over a flat average across ROIs precisely because the flat scheme means “tasks with multiple target ROIs ... would be over-weighted” [1]. Our  $\text{macro}\Delta$  applies the same principle one level across, with the held-out contrast in place of the target ROI. Where we differ is only in how that weighted quantity is reached: the Decathlon arrives at it through counts of pairwise Wilcoxon outcomes at an uncorrected per-comparison  $\alpha$ , whereas we test the weighted difference itself and Holm-correct across competitors.

*Why not a pooled Wilcoxon.* The signed-rank test is defined on a single undifferentiated sample of paired differences. Applying it to the concatenation of all contrasts therefore tests a case-weighted null — the imbalance above, reintroduced — and there is no standard reweighting of the signed-rank statistic that recovers the equal-per-contrast estimand. The disagreement this causes runs in both directions: a pooled test can declare a tenth of a Dice point significant because one large cohort supplies most of the cases, or miss a consistent effect confined to the smaller contrasts.

*The test.* Under the null that the two methods are exchangeable on a given patient, that patient’s paired differences are equally likely to carry either sign. Sign-flipping is therefore a valid randomization null for *any* statistic computed from those differences [12], macro $\Delta$  included, which is what makes it applicable where the signed-rank test is not. We report the one-sided  $p$  in the direction of the stated hypothesis.

*Evaluating the tail.* The statistic is a weighted sum of per-patient contributions, so its sign-flip null distribution can be computed rather than sampled. We enumerate all  $2^M$  sign patterns of the  $M=20$  patients with the largest contributions exactly, and treat the remaining, lighter patients as a Gaussian term with their exact null variance; with twenty patients or fewer the test is therefore exact. A plain normal approximation to the whole null, which an earlier version of this analysis used, is badly conservative in exactly the designs we have, where a few patients spanning several contrasts carry most of the null variance: on the Abdomen table it returned  $p = 0.013$  for a comparison that a 200,000-draw Monte-Carlo test places below  $5 \times 10^{-6}$ . The enumerated tail is deterministic and reaches far tails without the  $1/B$  floor of resampling. In the cross-task tests of Tab. 3 (1,456–1,730 patients), the Gaussian remainder carries 7–24% of the null variance for Dice (13–40% for HD95); the exactly enumerated twenty patients carry the rest. Exact enumeration also sets a hard floor for small cohorts: eight patients admit  $2^8 = 256$  sign patterns, so no one-sided  $p$  can fall below  $1/256$  however consistent the effect.

**Contrast-group pooling.** For Abdomen, the same tested modality is sometimes evidenced by more than one source: Abdomen’s CT evidence comes from CHAOS’s own held-out CT split plus AMOS and SLIVER07. Averaging every raw column with equal weight, the natural first implementation, silently lets a modality with more contributing sources outvote one with fewer in both the task-level average and the significance test, since it then supplies more of the equally-weighted columns being averaged. We instead pool disjoint sources of the *same* modality by concatenating their cases into one group before any averaging happens, then average modality groups with equal weight. For Abdomen this additionally means averaging equal-weight

across liver/spleen/kidney *within* the pooled CT evidence, since AMOS’s own CT column already blends all three organs into one number before it reaches this step. The same pooling applies to Mandible (HaN-Seg and PDDCA both supply CT) and Breast (I-SPY2, Duke and I-SPY1 supply T1-CE, and Duke and I-SPY1 also supply pre-contrast); the remaining tasks have no contrast evidenced by more than one source.

## 10. Per-setting results

Table 3 in the main paper reports results at the task level, pooling each task’s test contrasts (the training contrast included), all of its training modalities and its external cohorts. This section breaks the same evaluation out into the eighteen settings. Every number is a cross-fold, cross-label average over all test contrasts of the setting, the training contrast included, under the protocol of Sec. 4.1.

Tables 7 and 8 break each task column of Tab. 3 into its training modalities. Where a margin is not significant in a single setting, the honest reading is that one setting cannot separate the methods, not that they are equivalent: several settings have few test units (eight patients on MS, eight sessions on Brain), which is why the main claims rest on the task- and cross-task-level tests.

**Boundary distance at task level.** Table 9 gives the HD95 counterpart of Tab. 3. Across tasks PALETTE-Aug is significantly better than every competitor, Auglab included. Per task it has the lowest HD95 on Glioma, Abdomen, Brain, MS, Spine and Pelvis (on Abdomen the gap to SRCSM is not significant); it ties Auglab on Breast without a significant gap; and on Mandible it is significantly behind Auglab and the baseline, the one task where the baseline beats PALETTE-Aug. PALETTE alone, shown for reference, has the lowest HD95 on Breast. When HD95 disagrees with Dice for a method (a noticeably better HD95 than its Dice would predict), the usual explanation is low recall: a method that predicts only a small, confident, well-localized fragment of the structure can have a deceptively good boundary distance while missing most of the target, so a strong HD95 paired with a weak Dice is a warning sign rather than evidence of competitiveness.

| Setting              | Baseline | SynthSeg<br>-noEM | SynthSeg<br>-EM | SRCSM       | Auglab | PALETTE-<br>Aug |
|----------------------|----------|-------------------|-----------------|-------------|--------|-----------------|
| Glioma T1n           | 25.3*    | 8.7*              | 43.1*           | 42.7*       | 46.3   | <b>46.4</b>     |
| Glioma T1c           | 28.3*    | 9.5*              | 42.9*           | 40.2*       | 46.6   | <b>47.1</b>     |
| Glioma T2w           | 18.1*    | 8.4*              | 44.1*           | 30.1*       | 47.1   | <b>47.3</b>     |
| Glioma FLAIR         | 21.8*    | 8.7*              | 42.3            | 34.7*       | 39.4*  | <b>43.0</b>     |
| MS FLAIR             | 25.2*    | 4.6*              | 32.8*           | 27.9*       | 32.8*  | <b>39.5</b>     |
| MS T1w               | 16.7*    | 4.0*              | 35.1*           | 27.2*       | 35.8   | <b>37.4</b>     |
| Breast T1-CE         | 34.1*    | 16.6*             | 40.4*           | 16.9*       | 41.1*  | <b>41.9</b>     |
| Breast T2w           | 35.9*    | 16.5*             | 38.9*           | 12.5*       | 39.0*  | <b>40.5</b>     |
| Abdomen T1in         | 29.5*    | 63.3*             | 84.4*           | 84.9        | 84.7*  | <b>86.0</b>     |
| Abdomen T2spir       | 29.6*    | 58.5*             | 83.5*           | <b>86.3</b> | 84.5*  | 85.8            |
| Brain T1w            | 26.4*    | 66.1*             | 67.0*           | 66.2        | 67.1*  | <b>67.9</b>     |
| Brain T2w            | 35.9*    | 66.6              | <b>67.7</b>     | 66.6        | 67.2*  | 67.6            |
| Brain DWI            | 41.2*    | 59.7*             | 63.8            | <b>66.6</b> | 63.6   | 64.7            |
| Mandible CBCT        | 59.2*    | 53.6*             | 76.3*           | <b>79.8</b> | 78.3   | 77.7            |
| Spine CT             | 27.0*    | 46.7*             | 79.1*           | —           | 77.5*  | <b>80.2</b>     |
| Spine Dixon in-phase | 69.9*    | 48.4*             | 82.2*           | —           | 84.2   | <b>84.3</b>     |
| Pelvis CT            | 35.6*    | 17.1*             | 66.5*           | 67.3        | 68.5   | <b>69.2</b>     |
| Pelvis MRI           | 24.0*    | 12.8*             | 58.6*           | 58.5*       | 60.0   | <b>62.0</b>     |

Table 7. Per-setting Dice (%), best per row in bold. Each row is one setting (one training modality of one task), scored on every test contrast of its task, the training contrast and external cohorts included, exactly as its task column of Tab. 3, so the rows of a task average to that column. \*: PALETTE-Aug significantly better than that method within the setting (same patient-level sign-flip test, Holm-corrected across competitors). Spine has no SRCSM run.

| Setting              | Baseline    | SynthSeg<br>-noEM | SynthSeg<br>-EM | SRCSM       | Auglab      | PALETTE-<br>Aug |
|----------------------|-------------|-------------------|-----------------|-------------|-------------|-----------------|
| Glioma T1n           | 41.0*       | 78.0*             | 18.4*           | 17.0*       | <b>15.1</b> | 15.4            |
| Glioma T1c           | 40.3*       | 78.3*             | 19.5*           | 20.3*       | <b>15.1</b> | 15.4            |
| Glioma T2w           | 54.9*       | 79.7*             | 16.5*           | 29.2*       | <b>14.8</b> | 15.0            |
| Glioma FLAIR         | 43.1*       | 76.0*             | 17.3            | 20.3*       | 19.2*       | <b>16.6</b>     |
| MS FLAIR             | 39.9*       | 41.5*             | 23.6            | <b>23.0</b> | 26.6        | 24.3            |
| MS T1w               | 30.5*       | 42.3*             | 25.9            | 27.8*       | 26.3*       | <b>24.2</b>     |
| Breast T1-CE         | 61.8*       | 103.2*            | 61.7*           | 92.7*       | <b>57.3</b> | 59.0            |
| Breast T2w           | <b>55.2</b> | 99.1*             | 60.1*           | 92.1*       | 56.9        | 56.3            |
| Abdomen T1in         | 119.0*      | 78.0*             | 25.4            | 29.4*       | 30.4*       | <b>22.3</b>     |
| Abdomen T2spir       | 125.1*      | 83.2*             | 22.0            | <b>21.7</b> | 25.1        | 23.2            |
| Brain T1w            | 44.3*       | 11.2*             | 7.2             | <b>6.6</b>  | 7.4         | 6.9             |
| Brain T2w            | 27.8*       | 9.5*              | <b>4.0</b>      | 4.6*        | 4.3*        | 4.1             |
| Brain DWI            | 18.8*       | 5.7*              | 4.7             | 4.7*        | 4.7         | <b>4.6</b>      |
| Mandible CBCT        | <b>21.7</b> | 32.9*             | 24.9*           | 21.9        | 22.1        | 23.5            |
| Spine CT             | 78.3*       | 110.4*            | 4.5*            | —           | 5.8*        | <b>3.9</b>      |
| Spine Dixon in-phase | 21.6*       | 113.2*            | 3.5*            | —           | 3.6         | <b>3.1</b>      |
| Pelvis CT            | 146.2*      | 252.0*            | <b>20.0</b>     | 22.3        | 23.5        | 21.6            |
| Pelvis MRI           | 152.3*      | 271.6*            | 48.0            | 56.9*       | 47.8*       | <b>36.4</b>     |

Table 8. Per-setting HD95 (mm), best per row in bold. Each row is one setting (one training modality of one task), scored on every test contrast of its task, the training contrast and external cohorts included, exactly as its task column of Tab. 3, so the rows of a task average to that column. \*: PALETTE-Aug significantly better than that method within the setting (same patient-level sign-flip test, Holm-corrected across competitors). Spine has no SRCSM run.

|                                   | Glioma      | MS          | Breast      | Abdomen     | Brain      | Mandible    | Spine      | Pelvis      | $p$ vs ours          |
|-----------------------------------|-------------|-------------|-------------|-------------|------------|-------------|------------|-------------|----------------------|
| Baseline                          | 44.8        | 35.2        | 58.5        | 122.1       | 30.3       | <b>21.7</b> | 49.9       | 149.3       | $2 \times 10^{-47}$  |
| SynthSeg-noEM                     | 78.0        | 41.9        | 101.1       | 80.6        | 8.8        | 32.9        | 111.8      | 261.8       | $2 \times 10^{-58}$  |
| SynthSeg-EM                       | 17.9        | 24.7        | 60.9        | 23.7        | 5.3        | 24.9        | 4.0        | 34.0        | $2.7 \times 10^{-6}$ |
| SRCSM                             | 21.7        | 25.4        | 92.4        | 25.6        | 5.3        | 21.9        | —          | 39.6        | $1 \times 10^{-43}$  |
| Auglab                            | 16.1        | 26.5        | <b>57.1</b> | 27.8        | 5.4        | 22.1        | 4.7        | 35.6        | $2 \times 10^{-4}$   |
| PALETTE alone (ours)              | 16.6        | 23.3        | 56.3        | 24.9        | 4.9        | 22.3        | —          | 35.2        | $0.23^{\S}$          |
| <b>PALETTE-Aug (ours)</b>         | <b>15.6</b> | <b>24.3</b> | 57.7        | <b>22.7</b> | <b>5.2</b> | 23.5        | <b>3.5</b> | <b>29.0</b> | —                    |
| competitors beaten ( $p < 0.05$ ) | 5/5         | 3/5         | 3/5         | 4/5         | 2/5        | 2/5         | 4/4        | 5/5         |                      |

Table 9. Task-level HD95 (mm, ↓). Same pooling and test as Tab. 3, and likewise no cross-task average; the bottom row counts the competitors PALETTE-Aug beats within each task (on Mandible, Auglab and the baseline are significantly ahead of it). Spine has no SRCSM run and no ladder. PALETTE alone: the transform without Auglab, left out of the bolding;  $^{\S}$ PALETTE-Aug against PALETTE alone over the seven tasks with both. SRCSM includes its test-time source matching.

## 11. Ablation-ladder details

Tables 10 and 11 give the fill swap of every setting, Fig. 5 its HD95 counterpart per task, Tab. 12 every rung, and Tab. 13 every training  $\rightarrow$  test contrast pair.

**Full ladder: numbers and significance.** Figure 3 (main paper) shows the rung sequence visually; Tab. 12 gives the underlying out-of-domain Dice at every rung of all sixteen ladders, with significance for each step against the rung immediately before it (Holm-corrected paired Wilcoxon on patients, within each ladder’s family of steps). Most steps that add an ingredient help, as expected. The significant *decreases* (marked  $\downarrow$ ) are the Voronoi step on Glioma T1n and Abdomen T1in, which the real-fill step then more than recovers; the fill swap on Breast T1-CE and Brain T1w and T2w; and the +Auglab step on MS T1w, both Abdomen settings, Brain DWI and Mandible, the task where the complete method places third (Sec. 4.2), plus Mandible’s step to synthetic-contrast validation. On Mandible, PALETTE alone (the real-fill rung, 68.2) is the ladder’s maximum and adding the Auglab recipe takes it back down; adding PALETTE to Auglab also lowers Mandible Dice (Tab. 3), so the two components interact negatively there and we cannot attribute the shortfall to either; we have not diagnosed which Auglab operator is responsible. Abdomen T1in, MS T1w and Brain DWI also peak at the real-fill rung, while on all four Glioma settings and on Breast and Pelvis the +Auglab step adds Dice.

Voronoi’s own marginal contribution is task-dependent rather than uniformly positive. A separate paired significance check on all-contrast Dice (the estimand of Tab. 17, not the OOD slice above) shows that Voronoi significantly *helps* MS FLAIR and (by a small margin) Brain T1w, is flat on Abdomen T1in (its OOD step above is a small significant decrease), and significantly *hurts* Glioma T1n. The full method is still best on Glioma T1n (Tab. 7; not significantly ahead of Auglab): whatever Voronoi costs there is more than recovered by the subsequent real-intensity fill step, so the net effect of the complete pipeline is positive even though one internal component is not.

Table 13 breaks the fill swap down by training  $\rightarrow$  test contrast pair: the numbers behind the thin curves of Fig. 3.

**Per-label breakdown: is the fill-swap effect uniform within a task?** Glioma T2w’s pooled fill-swap effect is small and not significant (Tab. 10), with one significantly negative held-out contrast inside it (T2w $\rightarrow$ T1n, Tab. 13). This raises a further question: within that one (training, eval) pair, does the fill-swap step affect every anatomical sub-region equally, or is the worsening itself concentrated in specific labels? We break Glioma T2w $\rightarrow$ T1n down by its four labels: necrotic core (NCR), surrounding non-

| Task / setting                 | noise | $\Delta$ | rel. (%) | $p$     |
|--------------------------------|-------|----------|----------|---------|
| <i>Interface-bounded</i>       |       |          |          |         |
| <b>Abdomen</b> (all settings)  | 83.9  | +2.58*   | +3.1     | 1.7e-14 |
| Abdomen T1in                   | 85.7  | +1.43    | +1.7     | 3.1e-13 |
| Abdomen T2spir                 | 82.0  | +3.72    | +4.5     | 1.2e-10 |
| <b>Brain</b> (all settings)    | 64.0  | -0.72*   | -1.1     | 0.023   |
| Brain T1w                      | 63.1  | -1.48    | -2.3     | 0.0078  |
| Brain T2w                      | 64.2  | -1.40    | -2.2     | 0.0078  |
| Brain DWI                      | 64.6  | +0.54    | +0.8     | 0.0078  |
| <b>Mandible</b> (all settings) | 68.1  | +0.07*   | +0.1     | 0.023   |
| Mandible CBCT <sup>‡</sup>     | 68.1  | +0.07    | +0.1     | 0.0091  |
| <b>Pelvis</b> (all settings)   | 49.2  | +3.53*   | +7.2     | 0.00097 |
| Pelvis CT                      | 35.3  | +4.76    | +13.5    | 0.02    |
| Pelvis MRI                     | 63.1  | +2.59    | +4.1     | 0.00095 |
| <i>Appearance-defined</i>      |       |          |          |         |
| <b>Glioma</b> (all settings)   | 35.8  | +4.65*   | +13.0    | 6.4e-10 |
| Glioma T1n                     | 39.8  | +5.98    | +15.0    | 7.8e-11 |
| Glioma T1c                     | 30.5  | +12.73   | +41.8    | 3.6e-13 |
| Glioma T2w                     | 41.0  | +1.13    | +2.8     | 0.096   |
| Glioma FLAIR                   | 31.9  | -1.25    | -3.9     | 0.14    |
| <b>MS</b> (all settings)       | 26.8  | +6.82*   | +25.4    | 0.023   |
| MS FLAIR                       | 19.6  | +7.69    | +39.3    | 0.0078  |
| MS T1w                         | 34.1  | +5.94    | +17.4    | 0.0078  |
| <b>Breast</b> (all settings)   | 35.3  | +0.49*   | +1.4     | 0.00088 |
| Breast T1-CE                   | 30.8  | -1.16    | -3.8     | 6.1e-5  |
| Breast T2w                     | 39.8  | +2.13    | +5.3     | 3.9e-15 |

Table 10. Replacing SynthSeg-style noise fill with the real-intensity remap, *with the partition and region target means unchanged*. noise: out-of-domain Dice with noise fill;  $\Delta$ : Dice change at the swap; rel.:  $\Delta$  relative to the noise-fill Dice. Significance is tested once per task (bold rows; the test annotated in Fig. 3): two-sided paired Wilcoxon on patients, each patient contributing its mean difference over the held-out contrasts and training modalities it appears in, Holm-corrected across the seven tasks (\*:  $p < 0.05$ ). The indented setting rows break each task into its training modalities; their  $p$  is the same test within the setting, uncorrected and without significance marks, since the claim is made at the task level. With eight test patients (MS) or eight sessions of four subjects (Brain) a two-sided  $p$  cannot fall below  $1/128$ . <sup>‡</sup>Held-out contrasts are cross-dataset (HaN-Seg, PDDCA), since ToothFairy2 is CBCT only. Per pair: Tab. 13; HD95: Tab. 11.

enhancing FLAIR hyperintensity (SNFH), enhancing tumor (ET) and resection cavity (RC);  $n = 70$  test cases for SNFH and ET, present in essentially every case (NCR  $n = 51$ , RC  $n = 67$ , since not every tumor has a necrotic core or a resection cavity).

The worsening is concentrated in SNFH (Tabs. 14 and 15). SNFH has the largest gap between directions (+23.2 Dice) and the largest significant sign reversal: +14.8 when trained on T1n, -11.9 when trained on T2w. RC reverses too, at a smaller size. NCR and ET show no significant effect in either direction. Figure 6 shows this on one case: with real fill, the T2w-trained model’s SNFH pre-




Figure 5. Figure 3 (main paper) on HD95 instead of Dice: same tasks, grouping and visual scheme. Annotation: the task’s HD95 change at the fill swap in mm (negative is better); \*: significant. Glioma, Abdomen, Breast and Mandible improve significantly, Brain worsens significantly, and MS and Pelvis do not move significantly. Per-setting numbers: Tab. 11.

diction almost vanishes (label Dice 0.90  $\rightarrow$  0.06), while its ET prediction is almost unchanged (0.61  $\rightarrow$  0.68).

BraTSGLI02108101, fold 0 -- T2w-trained model evaluated on T1n




Figure 6. One Glioma test case, fold 0, T2w-trained model evaluated on T1n. Zoomed crops before (noise-fill) and after (real-fill) the fill-swap step, for the SNFH-dominant (top) and ET-dominant (bottom) lesion components of the same case. Color = ground truth wherever the prediction agrees; gray = false positive; outline = ground truth.

One candidate mechanism, flagged for future work rather than tested here: a texture pattern learned from one contrast’s rendering of a diffuse, texture-defined label like edema may not simply fail to transfer to a different contrast’s rendering of the same tissue; it may actively mislead, if the model has come to rely on a contrast-specific cue that the other contrast does not reproduce. This would be an instance of shortcut learning [16]: a network relying on a domain-specific cue that does not reproduce across acquisition protocols, the mechanism Ouyang et al. [42] identify behind failed single-source domain generalization in cross-modality and cross-sequence medical segmentation, rather than on the intended, causal feature, which a simplicity-biased network need not prefer even when it would generalize better [49]. We have not tested this mechanistically here; it is the most concrete hypothesis consistent with the label-specific, direction-dependent pattern above, not a demonstrated cause.

| Task / setting                 | noise | $\Delta$ (mm) | rel. (%) | $p$     |
|--------------------------------|-------|---------------|----------|---------|
| <i>Interface-bounded</i>       |       |               |          |         |
| <b>Abdomen</b> (all settings)  | 31.9  | -8.83*        | -27.7    | 0.045   |
| Abdomen T1in                   | 23.3  | -1.36         | -5.9     | 0.4     |
| Abdomen T2spir                 | 40.5  | -16.30        | -40.3    | 0.035   |
| <b>Brain</b> (all settings)    | 5.0   | +0.61*        | +12.3    | 0.031   |
| Brain T1w                      | 5.5   | +1.52         | +27.6    | 0.0078  |
| Brain T2w                      | 4.7   | +0.39         | +8.3     | 0.0078  |
| Brain DWI                      | 4.7   | -0.17         | -3.5     | 0.11    |
| <b>Mandible</b> (all settings) | 34.3  | -1.17*        | -3.4     | 0.021   |
| Mandible CBCT <sup>†</sup>     | 34.3  | -1.17         | -3.4     | 0.0042  |
| <b>Pelvis</b> (all settings)   | 59.6  | -2.79         | -4.7     | 0.43    |
| Pelvis CT                      | 67.1  | -3.61         | -5.4     | 0.031   |
| Pelvis MRI                     | 52.2  | -0.48         | -0.9     | 0.56    |
| <i>Appearance-defined</i>      |       |               |          |         |
| <b>Glioma</b> (all settings)   | 22.7  | -3.30*        | -14.5    | 1.5e-8  |
| Glioma T1n                     | 20.5  | -3.89         | -19.0    | 6.5e-9  |
| Glioma T1c                     | 27.6  | -9.77         | -35.4    | 6.5e-12 |
| Glioma T2w                     | 18.8  | -1.51         | -8.0     | 0.0047  |
| Glioma FLAIR                   | 23.9  | +1.97         | +8.2     | 0.011   |
| <b>MS</b> (all settings)       | 27.7  | -2.46         | -8.9     | 0.39    |
| MS FLAIR                       | 31.5  | -3.02         | -9.6     | 0.023   |
| MS T1w                         | 23.8  | -1.91         | -8.0     | 0.2     |
| <b>Breast</b> (all settings)   | 55.5  | -3.00*        | -5.4     | 5.4e-7  |
| Breast T1-CE                   | 70.1  | -3.43         | -4.9     | 0.00051 |
| Breast T2w                     | 40.9  | -2.57         | -6.3     | 2.3e-7  |

Table 11. Same fill-swap step as Tab. 10, on HD95 (negative is better; noise: HD95 with noise fill), with the same tests: task rows Holm-corrected across the seven tasks (\*:  $p < 0.05$ ), setting rows uncorrected. HD95 is less consistent than Dice within a task (on Glioma it improves from three training contrasts and worsens from FLAIR), so we state the claim on Dice and report HD95 here rather than as confirmation of it.

| Rung          | Glioma T1n | Glioma T1c | Glioma T2w | Glioma FLAIR | MS FLAIR | MS T1w | Breast T1-CE | Breast T2w |
|---------------|------------|------------|------------|--------------|----------|--------|--------------|------------|
| base          | 16.5       | 15.9       | 6.0        | 11.0         | 2.4      | 1.1    | 23.7         | 33.5       |
| +k-means      | 38.4***    | 29.6***    | 38.6***    | 26.1***      | 11.2*    | 9.8*   | 28.8***      | 37.0***    |
| +label remap  | 41.8***    | 30.3       | 40.2**     | 26.4         | 17.5*    | 22.4*  | 28.5         | 38.8***    |
| +Voronoi      | 39.8***↓   | 30.5       | 41.0       | 31.9***      | 19.6     | 34.1*  | 30.8***      | 39.8**     |
| real fill     | 45.8***    | 43.2***    | 42.1       | 30.7         | 27.3*    | 40.0*  | 29.6***↓     | 41.9***    |
| +Auglab       | 46.9       | 44.2       | 46.8***    | 40.9***      | 30.6     | 36.4*↓ | 34.7***      | 41.9       |
| +Auglab, s.v. | 47.6*      | 46.5***    | 46.9       | 41.2         | 31.3     | 34.7   | —            | —          |

| Rung          | Abdomen T1in | Abdomen T2spir | Brain T1w | Brain T2w | Brain DWI | Mandible CBCT | Pelvis CT | Pelvis MRI |
|---------------|--------------|----------------|-----------|-----------|-----------|---------------|-----------|------------|
| base          | 17.0         | 19.7           | 13.6      | 26.4      | 34.6      | 39.4          | 0.7       | 10.0       |
| +k-means      | 82.2***      | 77.4***        | 62.2*     | 63.3*     | 60.5*     | 64.6***       | 22.6***   | 55.9***    |
| +label remap  | 86.1***      | 81.4***        | 62.9      | 63.2      | 63.6*     | 65.6***       | 30.0***   | 63.3***    |
| +Voronoi      | 85.7***↓     | 82.0***        | 63.1      | 64.2*     | 64.6      | 68.1***       | 35.3**    | 63.1       |
| real fill     | 87.1***      | 85.7***        | 61.6*↓    | 62.8*↓    | 65.2*     | 68.2**        | 40.0*     | 65.7**     |
| +Auglab       | 86.4***↓     | 85.4***↓       | 63.3      | 63.7      | 62.7*↓    | 67.6***↓      | 46.3*     | 67.2       |
| +Auglab, s.v. | 86.1         | 85.9           | 63.6*     | 64.1      | 61.5      | 67.5***↓      | —         | —          |

Table 12. Out-of-domain Dice (%) at every rung of every ladder (Fig. 3), with significance for each step against the rung above it (“s.v.”, synthetic-contrast validation, against real-contrast validation): \* $p < 0.05$ , \*\* $p < 0.01$ , \*\*\* $p < 0.001$  (Holm-corrected paired Wilcoxon on patients, within each ladder; Tab. 10 tests the fill swap once per task instead); ↓ marks a significant *decrease*. Rungs: baseline, +k-means, +label remap, +Voronoi (noise fill), real fill (PALETTE alone), +Auglab (PALETTE-Aug); Fig. 3 abbreviates them base, +km, +lbl, +vor, +real, +Auglab. MS and Brain have eight test units each, so their smallest attainable  $p$  after correction is just under 0.05. Breast and Pelvis have no synthetic-validation arm. Mandible’s held-out contrasts all come from external cohorts.

| Task                      | Train → test         | Noise | Real | $\Delta$ | Rel. (%) | $p$                 |
|---------------------------|----------------------|-------|------|----------|----------|---------------------|
| <i>Appearance-defined</i> |                      |       |      |          |          |                     |
| Glioma                    | T1n → T1c            | 49.8  | 50.7 | +0.9*    | +1.7     | 0.04                |
|                           | T1n → T2w            | 39.2  | 45.5 | +6.3*    | +16.0    | $3 \times 10^{-10}$ |
|                           | T1n → FLAIR          | 30.4  | 41.2 | +10.8*   | +35.5    | $4 \times 10^{-11}$ |
|                           | T1c → T1n            | 40.7  | 43.6 | +2.9*    | +7.2     | $6 \times 10^{-4}$  |
|                           | T1c → T2w            | 25.4  | 43.6 | +18.2*   | +71.9    | $1 \times 10^{-12}$ |
|                           | T1c → FLAIR          | 25.3  | 42.3 | +17.0*   | +67.3    | $1 \times 10^{-12}$ |
|                           | T2w → T1n            | 41.8  | 37.0 | -4.7*    | -11.3    | $2 \times 10^{-4}$  |
|                           | T2w → T1c            | 37.9  | 41.4 | +3.6*    | +9.4     | $1 \times 10^{-3}$  |
|                           | T2w → FLAIR          | 43.2  | 47.8 | +4.6*    | +10.6    | $2 \times 10^{-5}$  |
|                           | FLAIR → T1n          | 26.0  | 22.6 | -3.5*    | -13.3    | $9 \times 10^{-3}$  |
| MS                        | FLAIR → T1c          | 32.5  | 31.3 | -1.2     | -3.7     | 0.76                |
|                           | FLAIR → T2w          | 37.2  | 38.2 | +0.9     | +2.5     | 0.76                |
|                           | FLAIR → T1w          | 15.6  | 21.8 | +6.2*    | +39.5    | 0.02                |
|                           | FLAIR → T2w          | 23.5  | 32.8 | +9.2*    | +39.2    | 0.02                |
|                           | T1w → FLAIR          | 38.3  | 41.9 | +3.6*    | +9.4     | 0.04                |
|                           | T1w → T2w            | 29.9  | 38.2 | +8.3*    | +27.7    | 0.02                |
|                           | T1-CE → T2w          | 21.9  | 23.1 | +1.2     | +5.5     | 0.05                |
|                           | T1-CE → pre-contrast | 35.7  | 31.0 | -4.8*    | -13.4    | $2 \times 10^{-9}$  |
|                           | T1-CE → DWI          | 34.7  | 34.8 | +0.1     | +0.3     | 0.71                |
|                           | T2w → T1-CE          | 51.0  | 52.9 | +1.9*    | +3.7     | $2 \times 10^{-4}$  |
| Breast                    | T2w → pre-contrast   | 40.6  | 40.5 | -0.1     | -0.3     | 0.31                |
|                           | T2w → DWI            | 27.9  | 32.5 | +4.6*    | +16.5    | $3 \times 10^{-16}$ |
| <i>Interface-bounded</i>  |                      |       |      |          |          |                     |
| Abdomen                   | T1in → T1out         | 84.9  | 86.1 | +1.2     | +1.4     | 0.50                |
|                           | T1in → T2spir        | 86.9  | 87.7 | +0.8     | +1.0     | 0.88                |
|                           | T1in → CT            | 83.4  | 85.2 | +1.9*    | +2.2     | $2 \times 10^{-9}$  |
|                           | T1in → T2w           | 87.5  | 89.4 | +1.9*    | +2.1     | $6 \times 10^{-6}$  |
|                           | T2spir → T1in        | 78.0  | 83.6 | +5.5     | +7.1     | 0.25                |
|                           | T2spir → T1out       | 82.4  | 87.7 | +5.2     | +6.4     | 0.25                |
|                           | T2spir → CT          | 78.9  | 81.5 | +2.6*    | +3.3     | $4 \times 10^{-8}$  |
|                           | T2spir → T2w         | 88.7  | 90.2 | +1.5     | +1.7     | 0.05                |
|                           | T1w → T2w            | 87.9  | 87.6 | -0.3     | -0.3     | 0.07                |
|                           | T1w → BOLD           | 49.0  | 46.5 | -2.5*    | -5.1     | 0.04                |
| Brain                     | T1w → DWI            | 71.8  | 70.5 | -1.3     | -1.8     | 0.06                |
|                           | T1w → EPI            | 48.8  | 46.3 | -2.5     | -5.1     | 0.12                |
|                           | T1w → GRE            | 57.8  | 56.9 | -0.8     | -1.4     | 0.07                |
|                           | T2w → T1w            | 90.6  | 89.7 | -0.9*    | -1.0     | 0.04                |
|                           | T2w → BOLD           | 50.2  | 48.7 | -1.4     | -2.9     | 0.06                |
|                           | T2w → DWI            | 72.6  | 71.9 | -0.7     | -1.0     | 0.08                |
|                           | T2w → EPI            | 50.4  | 48.7 | -1.8     | -3.5     | 0.12                |
|                           | T2w → GRE            | 57.1  | 54.9 | -2.1     | -3.7     | 0.06                |
|                           | DWI → T1w            | 79.0  | 80.6 | +1.6*    | +2.0     | 0.04                |
|                           | DWI → T2w            | 79.6  | 80.3 | +0.7*    | +0.9     | 0.04                |
| Mandible                  | DWI → BOLD           | 56.6  | 56.8 | +0.2     | +0.3     | 1.00                |
|                           | DWI → EPI            | 57.6  | 57.8 | +0.2     | +0.3     | 1.00                |
|                           | DWI → GRE            | 50.4  | 50.4 | +0.1     | +0.1     | 1.00                |
|                           | CBCT → HaN-Seg CT    | 83.2  | 83.7 | +0.5     | +0.6     | 0.17                |
|                           | CBCT → HaN-Seg MR-T1 | 54.3  | 53.3 | -1.0     | -1.8     | 0.23                |
|                           | CBCT → PDDCA CT      | 80.7  | 82.4 | +1.7*    | +2.1     | $5 \times 10^{-6}$  |
|                           | CT → MRI             | 35.3  | 40.0 | +4.8*    | +13.5    | 0.02                |
|                           | MRI → CT             | 63.1  | 65.7 | +2.6*    | +4.1     | $1 \times 10^{-3}$  |

Table 13. The fill swap for each training → test contrast pair: out-of-domain Dice (%) of the noise-fill rung and of the real-fill rung (PALETTE alone), their difference, the difference relative to the noise-fill Dice, and the pair’s paired Wilcoxon test on patients, Holm-corrected across the held-out contrasts of its setting (\*:  $p < 0.05$ , the pairs colored in Fig. 3; Tab. 10 tests each setting as a whole). External cohorts are pooled with the training cohort by contrast, except Mandible, whose held-out contrasts all come from external cohorts and are listed per cohort. 50 pairs.

| Label | T1n→T2w | T2w→T1n | gap   |
|-------|---------|---------|-------|
| NCR   | 18.3    | 21.5    | -3.2  |
| SNFH  | 66.4    | 43.2    | +23.2 |
| ET    | 27.2    | 26.9    | +0.4  |
| RC    | 50.6    | 43.1    | +7.5  |

Table 14. Absolute OOD Dice (%) of the real-fill rung (PALETTE alone), per label, both transfer directions.

| Label | T1n→T2w $\Delta$                     | T2w→T1n $\Delta$                    |
|-------|--------------------------------------|-------------------------------------|
| NCR   | +0.4 ( $p = 0.25$ )                  | +0.8 ( $p = 1.0$ )                  |
| SNFH  | +14.8* ( $p = 5.4 \times 10^{-12}$ ) | -11.9* ( $p = 1.6 \times 10^{-7}$ ) |
| ET    | +2.0 ( $p = 0.054$ )                 | -0.2 ( $p = 1.0$ )                  |
| RC    | +3.5* ( $p = 6 \times 10^{-5}$ )     | -6.5* ( $p = 0.005$ )               |

Table 15. Fill-swap step (+Voronoi noise fill → real fill), per label, both directions.  $p$ : Holm-corrected paired Wilcoxon, family = the four labels within each direction; \*: significant.

## 12. Synthetic contrast at validation time

Throughout the main paper PALETTE-Aug is applied with synthetic contrast during training only, and the best checkpoint is chosen on real-contrast validation images (*val000*), as for every competing method and every ladder rung; this keeps the comparison like-for-like. PALETTE-Aug can additionally be applied to the validation set, so that model selection is performed under randomized contrast as well (*val100*).

Validation-time randomization improves Glioma and Abdomen on both metrics (Tab. 16), most visibly on boundary distance for Abdomen ( $22.7 \rightarrow 20.6$  mm). Across tasks the HD95 gain is significant but the Dice change is not, and MS and Brain Dice move slightly the other way, so we do not claim it as a separate contribution; we report it because it costs nothing and because it indicates that the benefit of contrast randomization extends to model selection and is not confined to the gradient updates.

| Task    | Dice $\uparrow$ |             | HD95 (mm) $\downarrow$ |             |
|---------|-----------------|-------------|------------------------|-------------|
|         | val000          | val100      | val000                 | val100      |
| Glioma  | 45.9            | <b>46.6</b> | 15.6                   | <b>15.1</b> |
| Abdomen | 85.9            | <b>86.0</b> | 22.7                   | <b>20.6</b> |
| Brain   | <b>66.7</b>     | 66.5        | 5.2                    | <b>4.7</b>  |
| MS      | <b>38.5</b>     | 37.9        | 24.3                   | <b>24.2</b> |

Table 16. Applying PALETTE-Aug at validation time as well as at training time, on the four tasks with a task-level table for both configurations (Mandible’s val100 arm is scored only in its ladder, Tab. 12; the Breast, Spine and Pelvis arms were not trained). Across these tasks the two do not differ on Dice (same cross-task test as Tab. 3, two-sided  $p = 0.93$ ), and val100 has a lower HD95 (two-sided  $p = 0.031$ ,  $-0.8$  mm). We report the like-for-like val000 configuration throughout the main paper.

## 13. Component ablation on all test contrasts

To attribute the gain of the full method to its individual components, rather than to synthetic-contrast augmentation in general, we build an additive ablation ladder on top of the nnU-Net baseline (*Base*, Sec. 4.1): first the unsupervised  $k$ -means intensity partition alone (§3.1, step 1); then the anatomical label-remap step (step 4); then the Voronoi spatial sub-parcellation (step 2); and finally the full signed affine remap of Eq. (1) (step 3), which together constitute PALETTE alone. Every rung is evaluated without Auglab, so that the ladder isolates PALETTE’s own internal components from the separate question of whether it composes with a prior augmentation pipeline, which we address next. The synthetic-data fraction is fixed at 50% of training iterations throughout, as one of several augmentation constants (alongside, e.g., the  $k$ -means class count and the affine magnitude range of §3.1) that are set once and not tuned per dataset; validation uses real-contrast data only (no synthetic contrast at validation time).

**Composability with Auglab.** PALETTE-Aug is deployed, throughout this paper, as one additional operator inside the Auglab pipeline rather than as a replacement for it (Sec. 4.1); Auglab is a full prior augmentation method in its own right, not a component of PALETTE-Aug, so stacking it on top of the ladder above is a composability/deployment check, not a further component-ablation step. Table 17 keeps the two separate: the top five rows isolate PALETTE’s internal components, evaluated with no Auglab throughout, while the bottom row reports PALETTE-Aug plus Auglab — the configuration used everywhere else in this paper — set apart by a rule rather than presented as the next rung of the same ladder.

| Component                  | Abdomen<br>T1in | MS<br>FLAIR | Glioma<br>T1n | Brain<br>T1w |
|----------------------------|-----------------|-------------|---------------|--------------|
| Baseline (nnU-Net default) | 36.1            | 25.2        | 26.0          | 26.4         |
| + $k$ -means partition     | 84.8            | 30.6        | 41.7          | 67.4         |
| + label remap              | 87.7            | 34.6        | 44.3          | 68.0         |
| + Voronoi sub-parcellation | 87.2            | 36.1        | 42.8          | 68.1         |
| real fill (PALETTE alone)  | 87.9            | 40.4        | 47.6          | 66.9         |
| + Auglab (PALETTE-Aug)     | 87.4            | 39.5        | 47.5          | 67.9         |

Table 17. Component ablation (top five rows, additive, no Auglab) and composability check (bottom row). From the  $k$ -means row on, each rung adds one PALETTE-Aug component on top of the row above; the first step also replaces nnU-Net’s default intensity augmentation with spatial augmentation only. Dice (%) over all test contrasts of the setting, the training contrast included, with equal weight per contrast (the all-contrast series of the same ladders as Tab. 12, which reports the held-out contrasts only). Abdomen T1in: CHAOS’s own test contrasts only, without the external cohorts. The baseline and PALETTE-Aug rows are the runs of Tab. 7; their values can differ slightly from it (Glioma), because the ladder averages over the labels present in each case before averaging cases, whereas Tab. 7 averages each label over cases first.

## 14. Texture fidelity per method

The per-method NGF analysis is in the main paper (Fig. 4); Tab. 18 gives its per-task values. Figure 7 follows NGF and Dice along the ladder rungs on MS FLAIR (an earlier, FLAIR-only NGF run on 30 scans): the three noise-fill rungs stay at NGF 0.45–0.46 while their Dice rises from 11.2 to 19.6, through contrast randomization, not texture; only the real-fill swap raises both. (Table 18 averages FLAIR and T1w over 20 scans each, hence 0.42 for MS noise fill there.)

### Auglab keeps the most texture among the competitors.

Auglab keeps less source texture than PALETTE alone on six of seven tasks but far more than the label-style methods (Tab. 18), which fits its standing as the best competitor on five tasks (Tab. 3). The large margins appear against the methods that discard texture.

| Method        | Glioma       | MS           | Breast       | Abdomen      | Brain        | Mandible     | Pelvis       | Mean         |
|---------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| PALETTE alone | <b>0.836</b> | <b>0.842</b> | <b>0.809</b> | <b>0.805</b> | 0.721        | <b>0.904</b> | <b>0.863</b> | <b>0.826</b> |
| PALETTE-Aug   | 0.722        | 0.713        | 0.693        | 0.701        | 0.681        | 0.718        | 0.705        | 0.705        |
| Auglab        | 0.757        | 0.758        | 0.728        | 0.735        | <b>0.727</b> | 0.730        | 0.733        | 0.738        |
| SRCSM         | 0.663        | 0.658        | 0.645        | 0.636        | 0.626        | 0.600        | 0.662        | 0.641        |
| SynthSeg-EM   | 0.540        | 0.454        | 0.494        | 0.535        | 0.535        | 0.465        | 0.469        | 0.499        |
| SynthSeg-noEM | 0.337        | 0.337        | 0.332        | 0.374        | 0.504        | 0.360        | 0.339        | 0.369        |
| noise fill    | 0.465        | 0.423        | 0.448        | 0.463        | 0.463        | 0.373        | 0.416        | 0.436        |

Table 18. Texture fidelity per task: NGF similarity of the augmented volume to its source scan (1 = all texture kept,  $\approx 1/3$  = isotropic noise) on the seven tasks with a ladder (Spine has none), mean over contrasts of the mean over scans (16–20 per contrast, 5 draws each). PALETTE alone is the ladder’s real-fill rung. PALETTE-Aug stacks PALETTE (in half of the draws) on the full Auglab chain, so its draws also lose what the Auglab augmentations remove. PALETTE alone keeps more texture than every other method on every task, except Auglab on Brain ( $\Delta = -0.005$ ,  $p = 0.68$ ), where the difference is not significant (paired one-sided Wilcoxon over scans, Holm-corrected over the 42 comparisons; largest significant  $p = 4.1 \times 10^{-3}$ ).




Figure 7. Noise versus real fill on MS FLAIR. Left: out-of-domain Dice at each rung. Right: NGF similarity of the augmented volume to its source (mean  $\pm$  s.e. over 30 scans). The noise-fill rungs keep only the partition edges while Dice rises; the real-fill swap (shaded) lifts both.