

PALETTE-Aug: A Simple Texture-Preserving Augmentation for Medical Imaging Domain Generalization

Anonymous CVPR submission

Paper ID ****

Abstract

Introduction. Deep learning for medical image segmentation generalizes poorly across imaging modalities, and annotating data for every new domain is prohibitively costly. Synthesis-based augmentation methods such as SynthSeg address this by generating training images from label maps alone, but they require dense label maps and discard the source image’s texture. We introduce PALETTE-Aug, a texture-preserving alternative, and ask when texture matters for generalization. **Methods.** We benchmark PALETTE-Aug against five methods on eight segmentation tasks: five whose targets are bounded by an anatomical interface (healthy brain, multi-organ, healthy spine, mandible, pelvis) and three defined by tissue appearance (brain tumor, multiple sclerosis lesions, breast lesions). Each method trains on one modality and is evaluated in-domain and out-of-domain. An ablation isolates the effect of texture preservation. **Results.** PALETTE-Aug wins the benchmark overall: it has the highest Dice on six of the eight tasks and significantly beats every competitor on five of them. In a seven-task ablation, replacing noise fill with real texture raises per-task out-of-domain Dice by a relative 15% on brain tumor and 16% on multiple sclerosis lesions, both appearance-defined, but by at most 7% on interface-bounded tasks. Breast lesions, also appearance-defined, show no Dice gain. **Conclusion.** Texture preservation helps most on targets defined by tissue appearance rather than an anatomical interface, though not on all of them. PALETTE-Aug generalizes best overall across eight segmentation tasks.

1. Introduction

Medical image segmentation models are trained on one acquisition setting and are frequently applied to another: a different scanner, sequence, or imaging modality than the one seen during training. MRI contrast is a protocol choice, not a physical constant, so the same anatomy has different

intensities across sequences. A network trained only on T1-weighted brain images can fail outright on T2-weighted or FLAIR images [3]. Acquiring and annotating training data for every contrast a model might encounter is not practical: manual segmentation is expensive, and clinical sites rarely share raw data across institutions [33, 51]. The setting that matters for deployment is therefore single-source domain generalization: train on one labeled dataset in one contrast, then generalize to unseen contrasts and modalities without target-domain data [36, 45, 55, 57].

The dominant single-source strategy is label-generative domain randomization, popularized by SynthSeg [3]. It re-samples intensities from a per-label Gaussian mixture over a dense anatomical label map, a voxel-wise parcellation of the whole image rather than the sparse annotation a segmentation dataset normally provides. The resulting images have arbitrary, unrealistic contrast, forcing the network to rely on shape and layout [3]. This is effective, but carries two structural costs. First, it is *label-driven*: the dense parcellation it needs is richer than the task’s own annotation, so it reintroduces the expensive labeling it aims to avoid. Second, and less discussed, it *does not preserve texture*: replacing every labeled region with piecewise Gaussian noise discards the real gradients, partial-volume boundaries, and unlabeled fine structure the acquired scan contains. Layout survives; texture does not.

We ask when keeping texture, rather than discarding it, helps contrast robustness. Our central finding is that the answer depends on the *task*: preserving real intensity structure helps most where the target is defined by tissue appearance, and much less where an anatomical interface separates it from its surroundings. Prior work shows that texture carries diagnostic signal in appearance-defined tasks: texture features separate MS lesions from normal-appearing white matter [56], and diffuse glioma infiltrates beyond every margin visible on conventional MRI [17, 47].

PALETTE-Aug, our single-source augmentation, keeps this texture and needs no dense label map. PALETTE-Aug groups the foreground intensities of a scan into classes with 1-D k -means, splits each class into spatial sub-regions

with a Voronoi partition, and applies a signed random affine remap to the *real* intensities of each sub-region. Edges, gradients and partial-volume detail therefore survive inside each sub-region, because existing intensities are remapped rather than replaced by noise. Only the contrast between regions is randomized, including sign inversions that no acquisition protocol produces. An optional last step uses labels, and only the sparse task labels that any segmentation dataset already has.

We test the mechanism by intervention. Keeping the partition and the region target means fixed, we replace SynthSeg-style noise fill with the real-intensity remap in sixteen training settings and measure the Dice change relative to the noise-fill Dice. Pooled per task, the gain is largest on the appearance-defined targets, brain tumor (+15%) and MS lesions (+16%), against at most 7% on interface-bounded tasks. The converse does not hold: several appearance-defined settings, breast lesions among them, show no Dice gain, so we state the result in one direction only (§4.3). A segmentation-free gradient measurement confirms that the swap changes texture: noise fill keeps only the partition edges, and real fill recovers most of the source texture.

We train one nnU-Net [22] configuration, varying only the epoch budget per dataset, and evaluate cross-contrast, cross-modality and cross-dataset transfer on eight tasks: healthy brain, MS lesion, brain tumor, abdominal organ, mandible, breast lesion, spine and pelvis. We compare against nnU-Net with its default augmentation, two variants of SynthSeg (the original and our re-implementation of its EM label-subdivision extension), SRCSM [48], and Auglab [34], a published cross-modality augmentation pipeline. PALETTE-Aug wins the benchmark overall: it has the best Dice on six of the eight tasks, significantly beats every competitor on five of them, and is ahead of each competitor across tasks on both Dice and boundary distance.

In summary, we make the following contributions:

- **When texture preservation matters, tested by intervention.** Replacing noise fill with the real-intensity remap, with the partition fixed, gives the largest gains relative to the starting Dice on appearance-defined targets, though not on every one. A segmentation-free gradient measurement confirms that the swap changes texture.
- **PALETTE-Aug**, a single-source contrast augmentation that needs no dense label map, only the task’s own labels. In every comparison it runs as one added operator inside the Auglab pipeline; the ablation also tests it alone.
- **Breadth evidence.** Across cross-contrast, cross-modality and cross-dataset transfer on eight tasks and sixteen datasets, with one configuration reused throughout, PALETTE-Aug wins the benchmark overall: it is best on six tasks and significantly beats every competitor on five; where it does not win, we test whether the leader is truly

ahead.

2. Related Work

2.1. Domain generalization for medical segmentation

Domain generalization (DG) targets unseen distributions without access to target data, unlike unsupervised domain adaptation, which assumes unlabeled target images and solves an easier problem [44]. Most medical DG methods are *multi-source*, interpolating across several labeled domains by meta-learning, feature alignment, frequency-space exchange, multi-frequency architectures, foundation-model scale, or diffusion-based stylization [2, 30, 33, 35, 42, 51] – all of which still rely on source diversity. We instead work in the *single-source* regime (one labeled dataset, no target data, no second source), which is routine in clinical deployment [36, 57]. This rules out cross-domain interpolation and makes architecture-coupled methods orthogonal to our contribution: our mechanism acts at the input level and leaves the segmentation network a fixed nnU-Net [22].

2.2. Synthesis-based contrast-agnostic training

The dominant single-source strategy is label-generative domain randomization, introduced by SynthSeg [3]: intensities are sampled from a per-label Gaussian mixture over an anatomical segmentation, yielding images of arbitrary contrast that force contrast-invariance. This family has the two costs noted in Sec. 1: it needs a dense anatomical label map per subject, and it replaces real texture with piecewise noise. SIAM [49] trades label *quantity* for *quality* with curated templates, but still needs a full dense parcellation. The lineage has been specialized to white-matter and MS lesions and to atlas-based expectation-maximization [6, 7, 28]. Among SynthSeg’s own variants we take the original (*noEM*), which synthesizes from the label map alone and never looks at the acquired scan, and the *EM* extension, which first splits every label, background included, into intensity-coherent sub-regions by expectation-maximization clustering of the real image, so that sparse task labels yield a richer map to synthesize from. SynthSeg introduced the EM extension for cardiac MRI [3, §5.4]; we re-implement it.

2.3. Texture as a segmentation signal

Independent work establishes that local texture, not only intensity level, carries information a segmentation model can use: second-order gray-level statistics discriminate image content beyond simple intensity [18], and quantitative texture features from tumor regions carry prognostic signal across CT, MRI and PET [27]. Closer to our setting, Kang et al. disentangle structure from texture in cross-domain im-

age translation with a texture co-occurrence loss, targeting the same failure mode we do — naive translation distorts fine texture and degrades downstream segmentation [24]; a texture-fidelity loss has likewise been added to GAN-based lesion augmentation [15]. For diffuse white-matter lesions, it has been argued directly that the lack of a sharp edge defeats boundary-based methods and favors texture-based segmentation [26]. None of this prior work, to our knowledge, connects texture preservation in an *augmentation* to which segmentation targets it helps; that connection, tested causally, is our contribution.

2.4. Single-source intensity and appearance augmentation

Where the first line of single-source methods synthesizes from a label map, a second manufactures appearance diversity directly from the real image, with no generative label model. General-purpose libraries stack intensity, spatial and artifact transforms on an nnU-Net trainer; the recent Auglab pipeline [22, 34] is one such GPU-accelerated framework and is the strongest competitor across our tasks, and BigAug stacks a comparable set of coarse intensity, quality and spatial perturbations under the same single-source DG protocol we use [57]. Our method inserts the PALETTE-Aug transform as one additional operator inside the Auglab pipeline, so the comparison against Auglab directly isolates PALETTE-Aug’s contribution. Closer to our operator are random-filter methods, which vary appearance with an untrained nonlinear filter: GIN applies randomly re-initialized convolutions for label-free appearance variants [36]; related variants establish the random-filter principle for natural images, perturb intensities per class using the task mask, add learned diffusion generators, or perturb the Fourier amplitude spectrum instead of the image itself [8, 31, 39, 45, 55, 58]. Our operator instead partitions the foreground into intensity classes and Voronoi sub-regions and applies a signed *affine* remap of the real intensities within each, invertible and gradient-preserving by construction. The closest peer is SRCSM [48]. Like PALETTE-Aug it perturbs region-wise in 3-D with one label-consuming step: it applies a fresh untrained convolution per anatomical class, smooths the resulting boundary artifacts, and histogram-matches test images to the source. They differ in the transform: SRCSM applies a random nonlinear convolution, PALETTE-Aug an affine remap that keeps the relative intensity structure inside each region.

3. Method

We describe PALETTE-Aug (§3.1), how it is applied during training (§3.2), and the causal ablation used to isolate texture preservation as the source of its effect (§3.3).

3.1. The PALETTE-Aug transform

Let x be a training volume min-max normalized to $[0, 1]$, and F its foreground voxels ($x_v > 0.01$). PALETTE-Aug applies four steps at every training iteration (Fig. 1).

1. Intensity partition. We run 1-D k -means on the foreground intensities $\{x_v : v \in F\}$ with k sampled per iteration from $\{2, \dots, 6\}$, producing k intensity classes $C_1, \dots, C_k$ (Fig. 1, panel (b)). This partition groups voxels by intensity alone, with no label map, so it applies unchanged to any modality or contrast. Occasionally it is skipped and step 3 applies one global remap to the whole foreground; all step probabilities are in Sec. 7.

2. Spatial sub-parcellation. A Voronoi partition seeded with points sampled within each class C_i splits it into sub-regions $R_{i,1}, \dots, R_{i,m_i}$ that are spatially localized and similar in intensity (Fig. 1c). Two voxels of equal intensity in different parts of the anatomy can thus fall in different sub-regions and receive independent contrast changes. Each class is usually, but not always, sub-split this way (Sec. 7).

3. Signed affine remap. For each sub-region R , let $\bar{x}_R$ denote its mean intensity, $\mu_R \sim U(0, 1)$ a resampled target mean, and draw $\alpha_R \sim \pm U(0.5, 2.0)$ (a magnitude in $[0.5, 2.0]$ with an independently sampled sign). Every voxel $v \in R$ is remapped as

$$y_v = \mu_R + \alpha_R (x_v - \bar{x}_R). \quad (1)$$

Because Eq. (1) is a per-region affine function of the *real* intensity x_v , it preserves the relative structure of the signal within each region: edges, gradients, and partial-volume transitions between adjacent voxels are scaled, not replaced. Only the contrast between regions is randomized; $\alpha_R < 0$ also inverts local contrast, which no acquisition protocol produces. Background voxels are zeroed at this step (and remain zero through the rest of the pipeline). Half the time, the result is then Gaussian-blurred (Sec. 7).

4. Anatomical-label affine remap (optional). The unsupervised partition of steps 1–3 need not follow anatomy: two adjacent tissues of similar intensity can, by chance, receive the same target contrast, erasing their boundary. Step 4 corrects this: for each anatomical label ℓ , we draw a fresh (μ_ℓ, α_ℓ) and apply Eq. (1) *again* to ℓ ’s voxels of the step-3 output (Fig. 1d), so the label remap composes with the unsupervised one and keeps its internal structure. Since labels are independent of the unsupervised partition, the label remap restores a full-strength contrast step at every true label border, whether or not step 3 kept it; the anatomy itself is unchanged. It uses only the task’s own labels and is

optional: without it, PALETTE-Aug reduces to steps 1–3, e.g. on unlabeled data. An optional second blur and a foreground z-score normalization then give the training input (Fig. 1e).

3.2. Training protocol

PALETTE-Aug is applied on-the-fly as one contrast-augmentation operator inside an existing GPU-accelerated nnU-Net augmentation pipeline [22, 34], alongside that pipeline’s standard spatial and artifact augmentations; it adds no learned parameters and runs on the GPU inside that pipeline. One configuration (k -range, sub-region count, blur schedule, α range; Sec. 7) is set once and reused unchanged for every dataset and modality; no dataset gets its own hyperparameter sweep.

3.3. Causal ablation: isolating texture preservation

PALETTE-Aug differs from SynthSeg in two ways at once: the partition (unsupervised k -means and Voronoi vs. an anatomical label map) and the fill (remapped real intensities vs. Gaussian noise). A direct comparison cannot separate them.

We construct a controlled variant, *noise-fill*, with the same k -means/Voronoi partition and the same random target mean μ_R for each sub-region as Eq. (1). Instead of remapping the real intensities, it fills each sub-region with SynthSeg-style Gaussian noise, $\mu_R + \sigma_R \epsilon$ with $\epsilon \sim \mathcal{N}(0, 1)$ i.i.d. per voxel and σ_R drawn uniformly from $[0.05, 0.25]$ per sub-region; the label step is noise-filled the same way. The partition and the target means are PALETTE-Aug’s own, so a difference between the two isolates whether the real intensity structure inside each region is kept or replaced by noise. We report this comparison separately for *interface-bounded* targets (a normal structure bounded by another tissue class) and *appearance-defined* targets (a pathology identified by abnormal signal within a tissue), assigned by anatomy; the texture hypothesis predicts a larger gain on the latter.

4. Experiments

4.1. Setup

Protocol. Each model trains on one real source modality and is tested on every modality of the same task, the protocol Auglab uses for spine [34], here over eight tasks and with external cohorts wherever one exists. No target-domain data is seen in training. All methods share one nnU-Net [22] configuration and schedule and differ only in input augmentation; the epoch budget is fixed per dataset (Sec. 7). Our full method adds the PALETTE-Aug transform as one operator to the Auglab pipeline [34], so it and Auglab differ by exactly that transform. Each method is trained on three train/validation folds, and all three models are scored

on the same held-out test set. Significance tests are paired, use the patient (for Brain, the session) as the unit, and are Holm-corrected [20] (Sec. 9).

Datasets. Table 1 lists the eight tasks, the modalities each is trained on and the contrasts it is tested on, across MRI, CT and dental CBCT. Every (task, training modality) pair is one *setting*, eighteen in all; each setting trains its own models and is tested on every contrast of its task, the training contrast included. *Out-of-domain* (OOD) results exclude the training contrast, also in external cohorts. The training cohorts are, in table order, BraTS2024-GLI [10], Open-MS [29], I-SPY2 [32], CHAOS [25], ON-Harmony [52], ToothFairy2 [5], a CT/Dixon spine cohort and TotalSegmentator [9, 53]. The external test cohorts are Duke [13, 41], I-SPY1 [21] and ACRIN-6698 [37] (Breast), AMOS [23] and SLIVER07 [19] (Abdomen), HaN-Seg [38] and PDDCA [40] (Mandible), and SPIDER [50] (Spine); they are extra held-out test columns of their task, not separate tasks (sources per contrast: Tab. 5). ON-Harmony’s reference labels are silver-standard FreeSurfer [12] parcellations of the T1w, registered to the other contrasts. External volumes that image more of the body than the training cohort are cropped to its field of view before prediction, identically for every method (Sec. 8).

Baselines. Table 2 compares the five competitors: nnU-Net with its default augmentation (*Baseline*), SynthSeg [3] without and with its EM label-subdivision extension [3, §5.4] (re-implemented), both given each task’s training labels, which are a dense parcellation only on Brain; SRCSM [48]; and Auglab [34], a competitor in its own right. Each competitor is run as its authors propose: the SynthSeg variants and SRCSM synthesize on every training iteration with their own augmentation recipe, and Auglab uses its default configuration (Tab. 4). Our method adds one operator to Auglab, so comparing the two also tests that operator.

4.2. Main comparison

PALETTE-Aug is the best method overall. PALETTE-Aug has the best Dice on six of the eight tasks. Across tasks it is significantly better than every competitor ($p < 10^{-10}$ each, Tab. 3). Within single tasks it significantly outperforms every competitor on five tasks (Glioma, Abdomen, MS, Breast, Spine), three of five on Mandible and Pelvis, and two of five on Brain. The cross-task test generalizes to new patients on these eight tasks, not to new tasks. Counting each task as one unit instead, PALETTE-Aug beats Auglab on seven of eight tasks (one-sided sign test $p = 0.035$) and SRCSM on five of the seven tasks where SRCSM was run ($p = 0.23$).




Figure 1. The PALETTE-Aug pipeline on one real CHAOS T2SPIR slice (top) and one real ON-Harmony T1w slice (bottom). Every panel is a deterministic function of the same real image (no external model). Panels (b)–(d) paint each region, sub-region or label flat with its target mean μ (Eq. (1)) to show the partition alone; (e), the actual training input, applies the signed affine remap of Eq. (1), which keeps the real texture inside every region.

Task	Target	Trained on (one model per modality)	Tested on (every model; [†] from an external cohort)
Appearance-defined targets: abnormal signal inside a tissue			
Glioma	glioma, 4 sub-regions	T1n T1c T2w FLAIR	T1n, T1c, T2w, FLAIR
MS	MS lesions	FLAIR T1w	FLAIR, T1w, T2w
Breast	tumor	T1-CE T2w	T1-CE ^(†) , T2w, pre-contrast [†] , DWI [†]
Interface-bounded targets: a structure bounded by another tissue			
Abdomen	liver, kidneys, spleen	T1in T2spir	T1in, T1out, T2spir, T2w [†] , CT ^(†)
Brain	31 brain regions	T1w T2w DWI	T1w, T2w, DWI, BOLD, EPI, GRE
Mandible	mandible	CBCT	CBCT, CT [†] , MR-T1 [†]
Spine	vertebrae, discs, canal	CT Dixon in-phase	CT, Dixon in-phase, fat, water, T1w [†] , T2w [†]
Pelvis	11 bones and muscles	CT MRI	CT, MRI

Table 1. The eight tasks and eighteen settings. Each training modality trains its own models, which are tested on every contrast listed for the task, the training contrast included (^(†): from the training cohort and an external cohort). All tasks but Spine have an ablation ladder. Datasets and case counts per contrast: Tab. 5.

Given sparse labels, SynthSeg-noEM paints one Gaussian background around the labeled structures and falls below the baseline on six of the eight tasks, every one with sparse labels; SynthSeg-EM, which subdivides the labels by clustering the image, closes most of that gap.

No competitor is reliably second. The best competitor alternates between Auglab (Glioma, MS, Breast, Spine) and SRCSM (Brain, Mandible, Pelvis), the two tie on Abdomen, and SRCSM fails on Breast (16.4 Dice). Picking the right one would require knowing the answer in advance. PALETTE-Aug is first or second on every task but one.

Where PALETTE-Aug does not win. The one-sided test above cannot show that a competitor is ahead, so on Brain and Mandible we also ran the reverse test on the same patient-level design (Holm-corrected across competitors within the task). On Brain SRCSM leads by 0.1 Dice and neither method is significantly ahead. On Mandible SRCSM (+1.3 Dice, $p = 6 \times 10^{-4}$) and Auglab (+0.6, $p = 2.5 \times 10^{-3}$) are significantly ahead, and PALETTE-Aug is third. In the Mandible ablation (§4.3), real fill helps significantly; Dice drops only once Auglab is added (Tab. 12).

Method	Intensities	Varied per	No dense labels	Real texture
Baseline	real	—	✓	✓
SynthSeg-noEM	Gaussian	label	×	×
SynthSeg-EM	Gaussian	EM sub-label	✓	×
SRCSM	random convolution	label	✓	~
Auglab	real, generic aug.	image	✓	✓
PALETTE-Aug	real, affine remap	sub-region	✓	✓

Table 2. What each method trains on. Sub-region: unsupervised k -means and Voronoi partition (Sec. 3). No dense labels: works from the task’s own labels (SynthSeg-noEM runs on them but fills everything unlabeled with one Gaussian). ~: the real image passes through a random filter. Augmentation probabilities: Tab. 4.

Boundary distance. The 95th-percentile Hausdorff distance (HD95) is reported in the supplement (Tab. 9). PALETTE-Aug has the lowest HD95 on six of the eight tasks (Glioma, Abdomen, Brain, MS, Spine, Pelvis) and is significantly better than every competitor across tasks. It ties Auglab on Breast and is significantly behind SRCSM, Auglab and the baseline on Mandible.

4.3. Causal ablation: when does texture matter?

Does keeping the real texture help beyond PALETTE-Aug’s partition, and on which targets? We test this with a ladder in which each rung adds one ingredient: baseline $\rightarrow$ $+k$ -means $\rightarrow$ $+label$ remap $\rightarrow$ $+Voronoi$ $\rightarrow$ real-intensity fill $\rightarrow$ $+Auglab$. The three partition rungs fill every region with SynthSeg-style Gaussian noise. All rungs share one pipeline (Sec. 3) and differ only in which steps are on; before $+Auglab$, the only other augmentation is spatial (with flips). This step from $+Voronoi$ to real fill, the *fill swap*, changes only the fill: partition, sub-region count, target-mean distribution and both blur passes are unchanged. Any change at this step therefore comes from keeping the real within-region intensity structure instead of noise.

In our data every large relative gain from real fill is on an appearance-defined target, but not every appearance-defined target shows one (Fig. 2 and Tab. 10). We report each gain relative to its noise-fill Dice, because the same few points are a small share of a high starting Dice and a large share of a low one (Sec. 5 notes the alternative); the three largest gains are the same relative or absolute. They are Glioma T1c (+43.6%, +13.3 Dice) and T1n (+18.1%, +7.2), both significant after correction over the sixteen rows, and MS FLAIR (+39.3%, +7.7), which improves for all eight test patients but is not significant after that correction. No interface-bounded row gains more than 8.7% (Pelvis CT, not significant after correction); the largest absolute interface gain, Abdomen T2spir, is 5.1 points from a noise-fill Dice of 82.0. Pooled per task (Fig. 2), the step is worth +14.7% on Glioma and +15.8% on MS,




Figure 2. The fill swap (shaded band), one panel per task with a ladder, grouped by target type. Thin lines: one model, i.e. one training modality tested on one held-out contrast (external cohorts included), colored at the swap by that pair’s significance; values per pair: Tab. 13. Bold line: the task average. Annotation: the task’s Dice change at the swap relative to its noise-fill Dice; *: significant (paired Wilcoxon on patients pooling all training modalities, Holm-corrected across the seven tasks; Brain’s eight sessions come from four subjects). Per-setting numbers: Tab. 10; HD95: Fig. 4.

against +7.3% on Pelvis, +3.9% on Abdomen, +2.0% on Mandible and −1.0% on Brain. Not every appearance-defined setting gains: MS T1w, Glioma T2w and FLAIR, and both Breast rows (−2.8% and +1.3%, four cohorts pooled by contrast) show no Dice gain.

ON-Harmony labels 31 brain regions in the FreeSurfer label set [12]: ventricles, cortex, white matter, cerebel-

	Glioma	Abdomen	Brain	MS	Mandible	Breast	Spine	Pelvis	p vs ours
Baseline	23.4	29.5	35.0	21.0	59.2	35.0	48.5	29.8	$< 10^{-10}$
SynthSeg-noEM	8.8	60.9	64.8	4.3	53.6	16.5	47.5	15.0	$< 10^{-10}$
SynthSeg-EM	43.1	84.0	66.3	33.9	76.3	39.6	80.6	62.5	$< 10^{-10}$
SRCSM	41.5	84.6	66.6	33.0	79.0	16.4	—	65.3	$< 10^{-10}$
Auglab	44.8	84.6	66.1	34.3	78.3	40.1	80.9	64.3	$< 10^{-10}$
PALETTE-Aug (ours)	45.9	85.9	66.5	38.5	77.7	41.2	82.3	65.6	—
PALETTE alone (ours)	44.2	86.4	67.1	38.5	78.3	40.3	—	65.7	—
competitors beaten ($p < 0.05$)	5/5	5/5	2/5	5/5	3/5	5/5	4/4	3/5	

Table 3. Task-level Dice (%). Each column pools the task’s test cases of every contrast (the training contrast included), all source modalities and any external cohorts, weighting contrast groups equally so a modality with several external sources does not outvote one with fewer (Sec. 9). p : Holm-corrected one-sided sign-flip test on the macro-averaged per-patient difference, stratified by task with tasks weighted equally. We print no cross-task mean Dice: averaging whole-liver and breast-lesion Dice gives a number no clinician would recognize. The bottom row counts the competitors PALETTE-Aug beats within each task alone (same test). Spine has no SRCSM run (Tab. 5). PALETTE alone: the PALETTE transform without the Auglab pipeline (the ladder’s real-fill rung, Sec. 4.3), shown for reference and left out of the bolding; Spine has no ladder. HD95 is in Tab. 9.

lum, brainstem and eight deep grey nuclei. Nearly all are bounded by another tissue class (CSF against tissue, grey against white matter), so we group it with Abdomen as interface-bounded. At the task level it is the only task where the swap significantly hurts, with sessions as units. It is also the only task labeling nearly the whole brain; with one such task, we cannot tell whether that is why.

Does the fill swap change texture? The ladder shows that the fill swap changes Dice, not that it changes texture. We check this without any segmentation, with the normalized-gradient-field (NGF) similarity [16] between each augmented volume and its source scan. NGF measures how well image gradients align: it is 1 when all edges and fine structure are kept and about 1/3 for isotropic noise. Figure 3 reports it for every method on the seven ablation tasks. The noise fill keeps only the partition edges (mean NGF 0.44); the real fill, on the same partition, raises it to 0.83, the highest of all methods. PALETTE alone keeps more texture than every competitor on six tasks and ties with Auglab on Brain (Tab. 18). SynthSeg keeps little (0.50 with EM, 0.37 without, near the noise floor). PALETTE-Aug keeps less than PALETTE alone (0.70): it applies PALETTE in half of its draws, on top of the Auglab chain, which acts on every draw.

Supporting analyses. The supplement adds a component ablation (Sec. 11), per-setting results, every ladder rung and the HD95 ablation (Sec. 9).

5. Discussion and Limitations

Why anatomy predicts the largest effects. Liver, spleen and kidney are encapsulated organs with a clear physical boundary, and edge-following methods have long seg-




Figure 3. Texture kept by each augmentation: NGF similarity of the augmented volume to its source scan on the seven ablation tasks (violin: all scans, each averaged over 5 draws; dots: task means; bar: mean of the task means). The noise fill keeps only the partition edges; the real fill keeps the most.

mented them well [46]. An augmentation that scrambles their internal texture but keeps the boundary loses little. Diffuse glioma has no such boundary: tumor cells infiltrate beyond both the enhancing margin and the FLAIR abnormality [17, 47], so the target must be inferred from appearance. MS lesions, though typically well demarcated, are identified by abnormal signal within white matter, and texture features separate them from normal-appearing white matter [56]. This account interprets the intervention of Sec. 4.3; it is not independent evidence.

Not every appearance-defined target gains. MS T1w, Glioma T2w and FLAIR, and both Breast rows show no Dice gain. We have two hypotheses. First, the effect depends sharply on the training contrast (Glioma: +13.3 from T1c, +7.2 from T1n, about zero from T2w and FLAIR), so whether its texture transfers may matter as much as the target; with T2w training, real fill helps some labels and hurts others (Sec. 9). Second, Breast pools four cohorts from different trials, which may dilute a cohort-specific effect. On interface-bounded tasks, where texture adds little relative to the starting Dice, most of the gain arrives with the texture-free partition rungs (Tab. 12), that is, from the wide range of region-to-region contrasts PALETTE-Aug samples.

Limitations. We assign each target to a group by its anatomy, not by a measured boundary statistic. Target type is confounded with difficulty: all interface-bounded rows but one start above 63% noise-fill Dice and every appearance-defined row below 42%, so dividing by the starting Dice favors the latter; as a share of the remaining error, Abdomen T2spir would rank first. The exception, Pelvis CT, starts at 35% and gains +8.7%, below the large appearance-defined gains. The between-group comparison is descriptive: a permutation test on the seven per-task gains gives $p = 0.2$. MS has eight test patients and Brain eight sessions of four subjects, so their single-setting effects can be consistent without reaching significance, and Brain’s session-level tests are optimistic. Spine was run on a separate pipeline and has no SRCSM arm. The cross-task test covers new patients, not new tasks; with the task as the unit, the advantage over SRCSM is not significant. We applied each competitor as its paper proposes and did not develop it further; in particular, we did not combine SRCSM or SynthSeg with the full Auglab pipeline, which would be a new method of our own design rather than theirs. PALETTE-Aug targets the single-source setting and is complementary to multi-source domain generalization, which assumes several labeled source domains.

6. Conclusion

We asked when preserving real image texture matters for contrast-robust segmentation. Relative to the noise-fill Dice, preserving texture helps most where the target is defined by tissue appearance rather than an anatomical interface: +15% on brain tumor and +16% on MS lesions per task, against at most 7% on interface-bounded tasks. Not every appearance-defined target gains. PALETTE-Aug, which preserves texture and needs no dense label map, wins the eight-task benchmark overall and significantly beats every competitor on five tasks. The task-level account makes a testable prediction for new tasks: a large relative gain from texture preservation should not be expected where the target has an anatomical interface.

References

- [1] Michela Antonelli, Annika Reinke, Spyridon Bakas, Keyvan Farahani, Annette Kopp-Schneider, Bennett A. Landman, Geert Litjens, Bjoern Menze, Olaf Ronneberger, Ronald M. Summers, et al. The Medical Segmentation Decathlon. *Nature Communications*, 13(1):4128, 2022. 4
- [2] Jie Bao, Zhixin Zhou, Wen Jung Li, and Rui Luo. Structure-aware stylized image synthesis for robust medical image segmentation. *arXiv preprint arXiv:2412.04296*, 2024. 2
- [3] Benjamin Billot, Douglas N. Greve, Oula Puonti, Axel Thielscher, Koen Van Leemput, Bruce Fischl, Adrian V. Dalca, and Juan Eugenio Iglesias. SynthSeg: Segmentation of brain MRI scans of any contrast and resolution without retraining. *Medical Image Analysis*, 86:102789, 2023. 1, 2, 4
- [4] Benjamin Billot, Colin Magdamo, You Cheng, Steven E. Arnold, Sudeshna Das, and Juan Eugenio Iglesias. Robust machine learning segmentation for large-scale analysis of heterogeneous clinical brain MRI datasets. *Proceedings of the National Academy of Sciences*, 120(9):e2216399120, 2023. 4
- [5] Federico Bolelli, Luca Lumetti, et al. Multi-structure segmentation in CBCT volumes: The ToothFairy2 challenge. *Medical Image Analysis*, 112:104095, 2026. 4, 2, 3
- [6] Vicent Caselles-Ballester, Eloy Martínez-Heras, Giuseppe Pontillo, Zoe Mendelsohn, Elena M. Marrón, Juan Luis García Fernández, Laia Subirats, Jon Stutters, Jeremy Chataway, Frederik Barkhof, Sara Llufrui, and Ferran Prados. TimeLesSeg: Unified contrast-agnostic cross-sectional and longitudinal MS lesion segmentation via a stochastic generative model. *arXiv preprint arXiv:2605.07955*, 2026. 2
- [7] Stefano Cerri, Oula Puonti, Dominik S. Meier, Jens Wuerfel, Mark Muehlau, Hartwig R. Siebner, and Koen Van Leemput. A contrast-adaptive method for simultaneous whole-brain and lesion segmentation in multiple sclerosis. *NeuroImage*, 225:117471, 2021. 2
- [8] Boqi Chen, Yuanzhi Zhu, Yunke Ao, Sebastiano Caprara, Reto Sutter, Gunnar Rätsch, Ender Konukoglu, and Anna Susmelj. Generalizable single-source cross-modality medical image segmentation via invariant causal mechanisms. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 3592–3602, 2025. 3
- [9] Tugba Akinci D’Antonoli, Lucas K. Berger, Ashraya K. Indrakanti, et al. TotalSegmentator MRI: Robust sequence-independent segmentation of multiple anatomic structures in MRI. *Radiology*, 314(2):e241613, 2025. 4, 2, 3
- [10] Maria Correia de Verdier, Rachit Saluja, Louis Gagnon, Dominic LaBella, Ujjwal Baid, Nourel Hoda Tahon, Martha Foltyn-Dumitru, Jikai Zhang, Ahmed W. Moawad, et al. The 2024 brain tumor segmentation (BraTS) challenge: Glioma segmentation on post-treatment MRI. *arXiv preprint arXiv:2405.18368*, 2024. 4, 2
- [11] Michael D. Ernst. Permutation methods: A basis for exact inference. *Statistical Science*, 19(4):676–685, 2004. 5
- [12] Bruce Fischl. FreeSurfer. *NeuroImage*, 62(2):774–781, 2012. 4, 6, 1

- [13] Lidia Garrucho, Kaisar Kushibar, Claire-Anne Reidel, Smriti Joshi, Richard Osuala, et al. A large-scale multicenter breast cancer DCE-MRI benchmark dataset with expert segmentations. *Scientific Data*, 12(1):453, 2025. 4, 2, 3
- [14] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020. 6
- [15] Qiu Guan, Yizhou Chen, Zihan Wei, Ali Asghar Heidari, Haigen Hu, Xu-Hua Yang, Jianwei Zheng, Qianwei Zhou, Huiling Chen, and Feng Chen. Medical image augmentation for lesion detection using a texture-constrained multichannel progressive GAN. *Computers in Biology and Medicine*, 145: 105444, 2022. 3
- [16] Eldad Haber and Jan Modersitzki. Intensity gradient based registration and fusion of multi-modal images. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 726–733, 2006. 7
- [17] Alexander F. Haddad, Jacob S. Young, Ramin A. Morshed, and Mitchel S. Berger. FLAIRectomy: Resecting beyond the contrast margin for glioblastoma. *Brain Sciences*, 12(5):544, 2022. 1, 7
- [18] Robert M. Haralick, K. Shanmugam, and Its'hak Dinstein. Textural features for image classification. *IEEE Transactions on Systems, Man, and Cybernetics*, SMC-3(6):610–621, 1973. 2
- [19] Tobias Heimann, Bram van Ginneken, Martin A. Styner, Yulia Arzhaeva, Volker Aurich, Christian Bauer, Andreas Beck, Christoph Becker, Reinhard Beichel, György Bekes, et al. Comparison and evaluation of methods for liver segmentation from CT datasets. *IEEE Transactions on Medical Imaging*, 28(8):1251–1265, 2009. 4, 2, 3
- [20] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979. 4
- [21] Nola M. Hylton et al. Locally advanced breast cancer: MR imaging for prediction of response to neoadjuvant chemotherapy—results from ACRIN 6657/I-SPY TRIAL. *Radiology*, 263(3):663–672, 2012. 4, 2, 3
- [22] Fabian Isensee, Paul F. Jaeger, Simon A. A. Kohl, Jens Petersen, and Klaus H. Maier-Hein. nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18(2):203–211, 2021. 2, 3, 4, 1
- [23] Yuanfeng Ji, Haotian Bai, Chongjian Ge, Jie Yang, Ye Zhu, Ruimao Zhang, Zhen Li, Lingyan Zhang, Wanling Ma, Xiang Wan, and Ping Luo. AMOS: A large-scale abdominal multi-organ benchmark for versatile medical image segmentation. In *Advances in Neural Information Processing Systems*, pages 36722–36732, 2022. 4, 2, 3
- [24] Myeongkyun Kang, Philip Chikontwe, Dongkyu Won, Miguel Luna, and Sang Hyun Park. Structure-preserving image translation for multi-source medical image domain adaptation. *Pattern Recognition*, 144:109840, 2023. 3
- [25] A. Emre Kavur, N. Sinem Gezer, Mustafa Barış, Sinem Aslan, Pierre-Henri Conze, Vladimir Groza, Duc Duy Pham, Soumick Chatterjee, Philipp Ernst, Savaş Özkan, et al. CHAOS challenge – combined (CT-MR) healthy abdominal organ segmentation. *Medical Image Analysis*, 69:101950, 2021. 4, 2
- [26] Frithjof Kruggel, Joseph Suresh Paul, and Hermann-Josef Gertz. Texture-based segmentation of diffuse lesions of the brain’s white matter. *NeuroImage*, 39(3):987–996, 2008. 3
- [27] Philippe Lambin, Emmanuel Rios-Velazquez, Ralph Leijenaar, Sara Carvalho, Ruben G. P. J. van Stiphout, Patrick Granton, Catharina M. L. Zegers, Robert Gillies, Ronald Boellard, Andre Dekker, and Hugo J. W. L. Aerts. Radiomics: Extracting more information from medical images using advanced feature analysis. *European Journal of Cancer*, 48(4):441–446, 2012. 2
- [28] Pablo Laso, Stefano Cerri, Annabel Sorby-Adams, Jennifer Guo, Farrah Mateen, Philipp Goebel, Jiaming Wu, Peirong Liu, Hongwei Li, Sean I. Young, Benjamin Billot, Oula Puonti, et al. Quantifying white matter hyperintensity and brain volumes in heterogeneous clinical and low-field portable MRI. In *IEEE International Symposium on Biomedical Imaging (ISBI)*, pages 1–5, 2024. 2
- [29] Žiga Lesjak, Alfiia Galimzianova, Aleš Koren, Matjaž Lukač, Koen Van Leemput, and Žiga Špiclin. A novel public MR image dataset of multiple sclerosis patients with lesion segmentations based on multi-rater consensus. *Neuroinformatics*, 16(1):51–63, 2018. 4, 2
- [30] Da Li, Yongxin Yang, Yi-Zhe Song, and Timothy M. Hospedales. Learning to generalize: Meta-learning for domain generalization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018. 2
- [31] Heng Li, Haojin Li, Jianyu Chen, Mingyang Ou, Hai Shu, and Heng Miao. RaffesDG: Random frequency filtering enabled single-source domain generalization for medical image segmentation. *arXiv preprint arXiv:2405.01228*, 2024. 3
- [32] Wen Li, David C. Newitt, Jessica Gibbs, Lisa J. Wilmes, Ella F. Jones, Vignesh A. Arasu, et al. I-SPY 2 breast dynamic contrast enhanced MRI trial (ISPY2). The Cancer Imaging Archive, 2022. Version 1, data set. 4, 2, 3
- [33] Quande Liu, Cheng Chen, Jing Qin, Qi Dou, and Pheng-Ann Heng. FedDG: Federated domain generalization on medical image segmentation via episodic learning in continuous frequency space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1013–1023, 2021. 1, 2
- [34] Nathan Molinier, Hendrik Möller, Thomas Dagonneau, Anna Curto-Vilalta, Robert Graf, Matan Atad, Daniel Rueckert, Jan S. Kirschke, and Julien Cohen-Adad. One sequence to segment them all: Efficient data augmentation for CT and MRI cross-domain 3D spine segmentation. *arXiv preprint arXiv:2605.03098*, 2026. 2, 3, 4
- [35] Ju-Hyeon Nam, Nur Suriza Syazwani, Su Jung Kim, and Sang-Chul Lee. Modality-agnostic domain generalizable medical image segmentation by multi-frequency in multi-scale attention. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11480–11491, 2024. 2
- [36] Cheng Ouyang, Chen Chen, Surui Li, Zeju Li, Chen Qin, Wenjia Bai, and Daniel Rueckert. Causality-inspired single-

- source domain generalization for medical image segmentation. *IEEE Transactions on Medical Imaging*, 42(4):1095–1106, 2023. 1, 2, 3, 6
- [37] Savannah C. Partridge et al. Diffusion-weighted MRI findings predict pathologic response in neoadjuvant treatment of breast cancer: The ACRIN 6698 multicenter trial. *Radiology*, 289(3):618–627, 2018. 4, 2, 3
- [38] Gašper Podobnik, Primož Strojjan, Primož Peterlin, Bulat Ibragimov, and Tomaž Vrtovec. HaN-Seg: The head and neck organ-at-risk CT and MR segmentation dataset. *Medical Physics*, 50(3):1917–1927, 2023. 4, 2, 3
- [39] Qiang Qiao, Wenyu Wang, Meixia Qu, Kun Su, Bin Jiang, and Qiang Guo. Medical image segmentation via single-source domain generalization with random amplitude spectrum synthesis. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 435–445, 2024. 3
- [40] Patrik F. Raudaschl et al. Evaluation of segmentation methods on head and neck CT: Auto-segmentation challenge 2015. *Medical Physics*, 44(5):2020–2036, 2017. 4, 2, 3
- [41] Ashirbani Saha, Michael R. Harowicz, Lars J. Grimm, Connie E. Kim, Sujata V. Ghate, Ruth Walsh, and Maciej A. Mazurowski. A machine learning approach to radiogenomics of breast cancer: a study of 922 subjects and 529 DCE-MRI features. *British Journal of Cancer*, 119(4):508–516, 2018. 4, 2, 3
- [42] Sourya Sengupta, Satrajit Chakrabarty, Keerthi Sravan Ravi, Gopal Avinash, and Ravi Soni. SynthFM: Training modality-agnostic foundation models for medical image segmentation without real medical data. In *IEEE International Symposium on Biomedical Imaging (ISBI)*, pages 1–4, 2025. 2
- [43] Harshay Shah, Kaustav Tamuly, Aditi Raghunathan, Prateek Jain, and Praneeth Netrapalli. The pitfalls of simplicity bias in neural networks. In *Advances in Neural Information Processing Systems*, 2020. 6
- [44] Hyungseob Shin, Hyeongyu Kim, Sewon Kim, Yohan Jun, Taejoon Eo, and Dosik Hwang. SDC-UDA: Volumetric unsupervised domain adaptation framework for slice-direction continuous cross-modality medical image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7412–7421, 2023. 2
- [45] Zixian Su, Kai Yao, Xi Yang, Qiufeng Wang, Jie Sun, and Kaizhu Huang. Rethinking data augmentation for single-source domain generalization in medical image segmentation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 2366–2374, 2023. 1, 3
- [46] Kenji Suzuki, Hieu Trung Huynh, Yipeng Liu, Dominic Calabrese, Karen Zhou, Aytakin Oto, and Masatoshi Hori. Computerized segmentation of liver in hepatic CT and MRI by means of level-set geodesic active contouring. In *Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 2984–2987, 2013. 7
- [47] Jun Tang, Liangpeng Chen, Ziyang Li, Cong Li, and Deling Li. Imaging strategies for delineating glioma margins: Aiming at precise therapy. *Fundamental Research*, 6(4):2644–2653, 2026. 1, 7
- [48] Franz Thaler, Martin Urschler, Mateusz Koziński, Matthias A. F. Gsell, Gernot Plank, and Darko Štern. Semantic-aware random convolution and source matching for domain generalization in medical image segmentation. *IEEE Access*, 14: 64338–64356, 2026. 2, 3, 4
- [49] Romain Valabregue, Ines Khemir, Eric Badinet, François Rousseau, Guillaume Auzias, and Reuben Dorent. SIAM: Head and brain MRI segmentation from few high-quality templates via synthetic training. *arXiv preprint arXiv:2605.02737*, 2026. 2
- [50] Jasper W. van der Graaf, Miranda L. van Hooff, Constantinus F. M. Buckens, Matthieu Rutten, Job L. C. van Susante, Robert Jan Kroeze, Marinus de Kleuver, Bram van Ginneken, and Nikolas Lessmann. Lumbar spine segmentation in MR images: a dataset and a public benchmark. *Scientific Data*, 11:264, 2024. 4, 2, 3
- [51] Shujun Wang, Lequan Yu, Kang Li, Xin Yang, Chi-Wing Fu, and Pheng-Ann Heng. DoFE: Domain-oriented feature embedding for generalizable fundus image segmentation on unseen datasets. *IEEE Transactions on Medical Imaging*, 39(12):4237–4248, 2020. 1, 2
- [52] Shaun Warrington, Andrea Torchi, Olivier Mougin, Jon Campbell, Asante Ntata, Martin Craig, Stephanía Assimopoulos, Fidel Alfaro-Almagro, Karla L. Miller, Mark Jenkinson, Paul S. Morgan, and Stamatios N. Sotiropoulos. A multi-site, multi-modal travelling-heads resource for brain MRI harmonisation. *Scientific Data*, 12:609, 2025. 4, 1, 2
- [53] Jakob Wasserthal et al. TotalSegmentator: Robust segmentation of 104 anatomic structures in CT images. *Radiology: Artificial Intelligence*, 5(5):e230024, 2023. 4, 2, 3
- [54] Frank Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945. 4
- [55] Zhenlin Xu, Deyi Liu, Junlin Yang, Colin Raffel, and Marc Niethammer. Robust and generalizable visual representation learning via random convolutions. In *International Conference on Learning Representations (ICLR)*, 2021. 1, 3
- [56] Jing Zhang, Longzheng Tong, Lei Wang, and Ning Li. Texture analysis of multiple sclerosis: a comparative study. *Magnetic Resonance Imaging*, 26(8):1160–1166, 2008. 1, 7
- [57] Ling Zhang, Xiaosong Wang, Dong Yang, Thomas Sanford, Stephanie Harmon, Baris Turkbey, Bradford J. Wood, Holger Roth, Andriy Myronenko, Daguang Xu, and Ziyue Xu. Generalizing deep learning for medical image segmentation to unseen domains via deep stacked transformation. *IEEE Transactions on Medical Imaging*, 39(7):2531–2540, 2020. 1, 2, 3
- [58] Haoyu Zhao, Wenhui Dong, Rui Yu, Zhou Zhao, Bo Du, and Yongchao Xu. MoreStyle: Relax low-frequency constraint of Fourier-based image reconstruction in generalizable medical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 434–444, 2024. 3

PALETTE-Aug: A Simple Texture-Preserving Augmentation for Medical Imaging Domain Generalization

Supplementary Material

7. Training configuration

All methods share one nnU-Net [22] configuration; only the input augmentation differs. Nothing below is tuned per dataset.

Training schedule. 3D full-resolution nnU-Net, SGD with Nesterov momentum 0.99, initial learning rate 10^{-2} with the nnU-Net `polyLR` decay, 250 mini-batches per epoch, combined Dice and cross-entropy loss, deep supervision on. Patch size, spacing and batch size are whatever the nnU-Net planner selects for that dataset, identically for every method. The epoch budget is fixed per dataset and shared by all methods on it: 2500 (Glioma), 2000 (Brain, MS), 1000 (Mandible, Breast) and 200 (Abdomen, Pelvis); Spine was trained on a separate pipeline. Each method is trained as three models on three train/validation folds, all scored on the same held-out test set; folds are the same across methods.

PALETTE-Aug hyper-parameters. Set once and never tuned per dataset: k -means class count drawn per sample from $\{2, \dots, 6\}$, skipped in favor of a single global remap with probability 0.10; Voronoi sub-parcellation of each intensity class with probability 0.60 (else left as one region); signed affine coefficient $\alpha \in \pm U(0.5, 2.0)$ applied per sub-region as in Eq. (1) (always applied, to whatever partition results from the k -means/Voronoi steps); per-label affine decoupling applied independently per label with probability 0.5 (skipped for any label with under 4 foreground voxels in that sample); Gaussian blur σ drawn from $\{0, 0, 0, 0.3, 0.5, 0.8\}$, applied twice (after the region-level and after the label-level remap), so applied half the time at each point. The transform is applied on the GPU, on-the-fly, to a fixed 50% of training iterations (*train050*); the remaining half of iterations use the un-transformed real image.

Validation-time configuration. *val000* (the configuration reported in the main paper) applies no synthetic contrast to the validation split, matching every competing method and keeping model selection like-for-like. *val100* applies PALETTE-Aug to the entire validation split; it is reported in Sec. 10.

Baseline configurations. Both SynthSeg variants run through Auglab's GPU implementation of the SynthSeg

generator, with uniform intensity priors and with the generator's own spatial and flip steps disabled (the pipeline's spatial augmentation is shared by all methods). They receive each task's training label map; SynthSeg-EM additionally subdivides it on the fly by expectation-maximization clustering of the image (2 sub-regions per foreground label, 3–10 for the background, 20 iterations), where the original work precomputes these maps [3, §5.4]. Each sub-region then receives a randomly sampled intensity, as in SynthSeg. Both SynthSeg variants and SRCSM synthesize on every training iteration because their original papers train that way, so we keep that setting when running them; PALETTE-Aug synthesizes on 50% of iterations (Tab. 4 lists every augmentation's probability). SRCSM is a training-time augmentation only: no test-time intensity matching is applied, and every method is predicted with the same nnU-Net predictor, without test-time augmentation. SRCSM and Auglab otherwise use their published default configurations. No competing method was re-tuned in our favor, and all use the same folds, schedule and epoch budget as above.

Compute. Training was performed on NVIDIA L40S and H100 GPUs. A single fold costs from under an hour (Abdomen, 200 epochs) to roughly three days (Glioma, 2500 epochs); the full study is dominated by the six methods $\times$ eighteen settings $\times$ three folds of the main comparison, plus the sixteen ablation ladders of Sec. 4.3.

8. Datasets and evaluation protocol

All evaluation is performed on held-out test cases only. For each task the model is trained on a single source contrast and evaluated on every contrast of that task's test split, including the training contrast itself; training and validation subjects never appear in any reported number.

ON-Harmony [52] — 31-region healthy-brain segmentation. Reference labels are *silver standard*: FreeSurfer parcellations [12] of each session's T1w, grouped into 31 regions and rigidly registered with ANTs into the native space of the same session's other contrasts. They are not manual annotations, so absolute Dice here is agreement with an automated reference; every method is scored against the same reference. The test set is eight sessions of four subjects, each scanned at two sites.

	Transform	Baseline	SynthSeg-noEM	SynthSeg-EM	SRCSM	Auglab	PALETTE-Aug
Synthesis	SynthSeg generator	–	1.0	1.0	–	–	–
	semantic random convolution	–	–	–	1.0	–	–
	PALETTE transform	–	–	–	–	–	0.5
Intensity	Gaussian noise	0.1	0.1	0.1	–	0.1	0.1
	Gaussian blur	0.2	0.2	0.2	–	0.2	0.2
	brightness	0.15	0.15	0.15	–	0.15	0.15
	contrast	0.15	0.15	0.15	–	0.15	0.15
	gamma / inverted gamma	0.3 / 0.1	0.3 / 0.1	0.3 / 0.1	–	0.3 / 0.1	0.3 / 0.1
	simulated low resolution	0.25	0.25	0.25	–	0.25	0.25
	bias field	–	0.15	0.15	–	0.15	0.15
	random intensity function	–	0.1	0.1	–	0.1	0.1
	Scharr / unsharp mask	–	–	–	–	0.25 / 0.25	0.25 / 0.25
	random convolution	–	–	–	–	0.2	0.2
	label-wise intensity redistribution	–	–	–	–	0.4	0.4
	intensity inversion	–	–	–	–	0.25	0.25
	histogram equalization	–	–	–	–	0.25	0.25
Spatial	rotation / scaling	0.2 / 0.2	0.2 / 0.2	0.2 / 0.2	–	0.2 / 0.2	0.2 / 0.2
	affine + elastic deformation	–	–	–	1.0	–	–
	flip	0.5 ^a	0.5	0.5	–	0.5	0.5

Table 4. Per-iteration probability of each augmentation, per method (–: not applied). The synthesis rows are each method’s own contrast step: the SynthSeg variants and SRCSM synthesize on every iteration, as in their original papers, PALETTE-Aug on half of them. The remaining rows are the method’s generic augmentation pipeline: nnU-Net’s defaults for the baseline, Auglab’s for the Auglab-based methods (the SynthSeg variants switch off Auglab’s appearance-randomizing operators, which would otherwise act on top of the synthetic image), and SRCSM’s own spatial recipe. ^annU-Net mirroring, each axis independently.

Task	Target	Trained on (one modality at a time)	Other test contrasts	External test cohorts	Test cases	Ladder
<i>Appearance-defined</i>						
Glioma	glioma, 4 sub-regions	BraTS2024-GLI [10]: T1n, T1c, T2w, FLAIR	–	–	70	✓
MS	MS lesions	Open-MS [29]: FLAIR, T1w	T2w	–	8	✓
Breast	tumor	I-SPY2 [32]: T1-CE, T2w	–	Duke [13, 41], I-SPY1 [21]: T1-CE, pre-contrast; ACRIN-6698 [37]: DWI	84 + 829	✓
<i>Interface-bounded</i>						
Abdomen	liver, kidneys, spleen	CHAOS [25]: T1in, T2spir	T1out, CT	AMOS [23]: CT, T2w; SLIVER07 [19]: CT	24 + 140	✓
Brain	31 brain regions	ON-Harmony [52]: T1w, T2w, DWI	BOLD, EPI, GRE	–	8 ^a	✓
Mandible	mandible	ToothFairy2 [5]: CBCT	–	HaN-Seg [38]: CT, MR-T1; PDDCA [40]: CT	71 + 82	✓
Spine	vertebrae, discs, canal	CT/Dixon cohort: CT, Dixon in-phase	Dixon fat, water	SPIDER [50]: T1w, T2w	29 + 171/219	–
Pelvis	11 bones and muscles	TotalSegmentator [9, 53]: CT, MRI ^b	–	–	30/27	✓

Table 5. Per-task dataset roster (summary in Tab. 1, main paper). Test cases: own held-out split + external cohorts; *a/b* gives two contrasts’ counts where they differ (Tab. 6). Ladder: the task has the causal ablation of Sec. 4.3. ^aSessions of four subjects. ^bDifferent patients for CT and MRI.

892	Open-MS [29] — MS-lesion segmentation with manual	each of T1n, T1c, T2w and T2-FLAIR and evaluated on all	898
893	consensus annotations, evaluated across FLAIR, T1w and	four.	899
894	T2w.		
895	BraTS2024-GLI [10] — post-treatment glioma sub-	CHAOS [25] — abdominal organ segmentation (liver,	900
896	region segmentation (necrotic core, surrounding FLAIR hy-	spleen, both kidneys), evaluated on the CHAOS MR con-	901
897	perintensity, enhancing tumor, resection cavity), trained on	trasts (four held-out MR patients) and CT (20 held-out pa-	902
		tients), and on two external cohorts: CT and T2w MRI from	903

904 AMOS [23] and liver-only CT from SLIVER07 [19]. In
905 the pooled CT group, liver comes from CHAOS, AMOS
906 and SLIVER07 and spleen and kidneys from AMOS only
907 (CHAOS CT is annotated for the liver only). CHAOS im-
908 ages a restricted axial field of view, so the full-torso external
909 volumes are cropped to the CHAOS-equivalent axial slab
910 (median CHAOS margins around the kidneys for AMOS
911 and the liver for SLIVER07) before prediction; the network
912 sees, and is scored on, only the cropped volume.

913 **ToothFairy2, HaN-Seg and PDDCA** [5, 38, 40] —
914 mandible segmentation, trained on dental cone-beam CT,
915 the one non-CT, non-MRI training modality here. Tooth-
916 Fairy2’s annotation coverage is cohort-dependent, so every
917 column of this task is scored on the mandible alone (teeth
918 excluded, matching the external references), which keeps
919 the numbers being averaged commensurable. ToothFairy2
920 acquires CBCT only, so apart from its own 71 CBCT test
921 cases its held-out contrasts all come from other datasets:
922 HaN-Seg CT and MR-T1 (42 cases each) and PDDCA CT
923 (40 cases), the two CT sources pooled as one contrast group.
924 ToothFairy2’s volumes are stored superior-inferior inverted
925 relative to the external cohorts — an inversion invisible to
926 the image header, found by checking anatomy directly —
927 so the external test volumes are flipped into ToothFairy2’s
928 convention before inference, which leaves Dice and HD95
929 unchanged under the rigid transform. The correction mat-
930 ters for the comparison: a method without flip augmentation
931 (SRCSM) was penalised by the inversion far more than the
932 others, and it is after this correction that SRCSM leads this
933 task.

934 **Breast: I-SPY2, Duke-Breast-MRI, I-SPY1, ACRIN-**
935 **6698** [13, 21, 32, 37, 41] — breast-lesion segmentation
936 on malignant-only cohorts. I-SPY2 (CC BY 4.0) trains, on
937 its first-post-contrast T1 (T1-CE) and its T2w; its reference
938 is the trial’s functional tumor volume mask. Three external
939 cohorts are test-only. Duke-Breast-MRI and I-SPY1 come
940 through the MAMA-MIA release, whose expert primary-
941 lesion protocol both share, and contribute T1-CE and pre-
942 contrast columns; Duke-Breast-MRI is CC BY-NC 4.0, a
943 non-commercial term no other cohort here carries. ACRIN-
944 6698, the diffusion sub-study of I-SPY2 with no patient in
945 common with our I-SPY2 subjects, contributes the baseline
946 diffusion-weighted acquisition with the trial’s whole-tumor
947 DWI region, a contrast neither training modality resembles.
948 Because I-SPY2 is mostly acquired unilaterally, bilateral
949 external volumes are cropped to the lesion side and, for
950 the full-thorax acquisitions, also anterior-posteriorly to the
951 breast-plus-chest-wall extent of I-SPY2’s own crops; with-
952 out the second crop every method’s HD95 is dominated by
953 the thorax. I-SPY2’s DCE series are acquired unilaterally
954 at some sites and bilaterally at others; both field-of-view

variants are kept per patient, and the significance tests treat
them as one patient.

Field-of-view matching. A model tested on a field of
view it never saw in training can fail for reasons that have
nothing to do with contrast or modality. Where an external
cohort images more of the body than the training cohort,
we therefore crop its volumes before prediction, so that the
shift left between training and test is, as far as possible,
the contrast or modality alone. CHAOS (AMOS, SLIVER07):
the axial slab CHAOS covers, placed with the median CHAOS
margins around the kidneys (AMOS) or the liver (SLIVER07).
Mandible (HaN-Seg, PDDCA): a $144 \times 128 \times 80$ mm box
centred on the mandible, close to the 90th percentile of
ToothFairy2’s training extent, enlarged (at most to $154 \times 154 \times 89$ mm)
only when needed to contain the whole mandible. Breast
(Duke-Breast-MRI, ACRIN-6698): the lesion-side half of
bilateral volumes, then, for full-thorax acquisitions, an
anterior-posterior window from 27 mm in front of the skin
to 147 mm behind it, the measured geometry of I-SPY2’s
unilateral acquisitions, extended if needed to contain the
lesion; I-SPY1 is natively unilateral and is not cropped.
The network sees, and is scored on, only the cropped volume.
Each crop is placed using the reference annotation (organ,
mandible or lesion position). This reveals the target’s
position, not its extent, and is identical for every method,
so it does not favour any method; but absolute Dice and
HD95 are better than a pipeline that had to localize the
target first would obtain.

Spine [50] — vertebrae, intervertebral discs and spinal
canal. Results for this task were computed on a separate
pipeline with the same method suite except SRCSM, trained
on CT and Dixon in-phase MRI of a CT/Dixon cohort and
tested on that cohort’s held-out CT and Dixon fat, in-phase
and water images and on the T1w and T2w of SPIDER.
Dice and HD95 are computed on each whole structure per
label (the pipeline’s `global_bin` metrics). No SRCSM
arm was run, so Spine is absent from every SRCSM com-
parison.

Pelvis [9, 53] — eleven pelvic and hip musculoskeletal
structures (both hips, sacrum, and the gluteus maximus,
medius and minimus and iliopsoas on each side), from the
TotalSegmentator CT (CC BY 4.0) and MRI (CC BY-NC-SA
2.0) releases. The two are unpaired cohorts, and the MRI
release mixes sequences by design, so its MRI column is
a sequence-heterogeneous test rather than one named
contrast. CT is subsampled to 200 cases to keep training
time comparable with the other tasks.

Task	Train	Test
CHAOS	16	24
ON-Harmony	84 ^a	8
Open-MS	22	8
BraTS2024-GLI	630	70
ToothFairy2	408	71
I-SPY2	476	84
Pelvis (CT / MRI)	170 / 156	30 / 27

Table 6. Per-task subject counts for the training cohorts, pooled across the 3 folds used throughout (± 1 between folds). ON-Harmony counts are scan sessions; the eight test sessions are four subjects, each scanned at two sites. External cohorts are test-only and counted separately: AMOS and SLIVER07 (CHAOS), HaN-Seg and PDDCA (Mandible), Duke-Breast-MRI (291), I-SPY1 (167, after excluding three bilateral cancers and one implant by a rule fixed before prediction) and ACRIN-6698 (371) for Breast, and SPIDER (171 T1w, 219 T2w) for Spine, whose own CT/Dixon test split has 27–29 cases per contrast. I-SPY2 counts are patients; one held-out patient can yield more than one T2w image (both field-of-view variants). Pelvis MRI keeps 183 of the 298 released cases. ^a76 for the T2w-trained models.

9. Per-setting results

Tab. 3 in the main paper reports results at the task level, pooling each task’s held-out test contrasts, all of its source modalities and its external cohorts. This section reports the same evaluation broken out into the eighteen individual source-modality settings. Every number is a cross-fold, cross-class average over all held-out contrasts of that setting, under the protocol of Sec. 4.1.

Statistical test construction. All tests are paired, and the unit of independence is the *patient*, never the fold, the image or the contrast. A patient’s score is its mean over labels and evaluated folds; where the same patient is scored in several held-out contrasts, or by the models of several training modalities, those scores move together as one unit under the null. Treating them as separate samples would be pseudo-replication and would inflate the effective sample size — a real risk here, since several tasks test a handful of patients under many contrasts (Brain: eight sessions of four subjects, six contrasts). For Brain the unit is the session, since its four subjects leave a test that could never reach significance; two sessions of one subject are not fully independent, so its *p*-values are, if anything, optimistic, and we do not rest any claim on Brain alone. Per-setting comparisons (below) use the same test within each setting; the ablation ladders use the paired Wilcoxon signed-rank test [54], since there each setting is judged on its own. All tests are Holm-corrected [20] for multiple comparisons.

Why the aggregate tests are sign-flip tests on a macro-average. The aggregate tests (Tab. 3 and the per-dataset headline blocks) are not paired Wilcoxon signed-rank tests, which is the more usual choice in this literature: SynthSeg+, from the lineage we compare against, marks a method as better than its competitors using a one-sided Bonferroni-corrected Wilcoxon signed-rank test on paired Dice [4]. We therefore give our reason rather than leave it implicit. The *estimand* we test, by contrast, is not unusual at all — it is the same equal-weight-per-stratum principle the Medical Segmentation Decathlon adopted, for the same reason, as set out below.

The estimand. What we claim is $\text{macro}\Delta$: the mean, over held-out contrasts, of that contrast’s mean paired per-case difference (ours minus competitor) — equal weight per contrast, which is exactly what the *all* column of every table reports, so the test and the table measure the same thing. Equal weighting is not cosmetic here. The claim under test is cross-contrast generalization, and a contrast’s held-out case count reflects how large a public cohort happens to be, not how much that contrast should say about generalizing. Abdomen is the extreme case: its held-out CT evidence includes an external cohort of 100 patients, while each of its three MR contrasts has 4 (cases counted once per training modality: 200 against 8). Weighting by case count would let one CT cohort cast more than half the votes in a test whose conclusion is about transfer across MR contrasts.

This is the same problem, and the same answer, as the Medical Segmentation Decathlon’s. Facing sample sizes that “varied heavily across all tasks and target ROIs,” it ranked algorithms within each target ROI, averaged those ranks per task, and averaged the per-task ranks for the final score — choosing that over a flat average across ROIs precisely because the flat scheme means “tasks with multiple target ROIs ... would be over-weighted” [1]. Our $\text{macro}\Delta$ applies the same principle one level across, with the held-out contrast in place of the target ROI. Where we differ is only in how that weighted quantity is reached: the Decathlon arrives at it through counts of pairwise Wilcoxon outcomes at an uncorrected per-comparison α , whereas we test the weighted difference itself and Holm-correct across competitors.

Why not a pooled Wilcoxon. The signed-rank test is defined on a single undifferentiated sample of paired differences. Applying it to the concatenation of all contrasts therefore tests a case-weighted null — the imbalance above, reintroduced — and there is no standard reweighting of the signed-rank statistic that recovers the equal-per-contrast estimand. The disagreement this causes runs in both directions: a pooled test can declare a tenth of a Dice point significant because one large cohort supplies most of the cases, or miss a consistent effect confined to the smaller contrasts.

The test. Under the null that the two methods are ex-

changeable on a given patient, that patient's paired differences are equally likely to carry either sign. Sign-flipping is therefore a valid randomization null for *any* statistic computed from those differences [11], macro Δ included, which is what makes it applicable where the signed-rank test is not. We report the one-sided p in the direction of the stated hypothesis, with the two-sided p retained as a secondary column, and give a 95% confidence interval from a bootstrap that resamples patients within groups of patients sharing the same set of contrasts, so each source keeps its observed size.

Evaluating the tail. The statistic is a weighted sum of per-patient contributions, so its sign-flip null distribution can be computed rather than sampled. We enumerate all 2^M sign patterns of the $M=20$ patients with the largest contributions exactly, and treat the remaining, lighter patients as a Gaussian term with their exact null variance; with twenty patients or fewer the test is therefore exact. A plain normal approximation to the whole null, which an earlier version of this analysis used, is badly conservative in exactly the designs we have, where a few patients spanning several contrasts carry most of the null variance: on the Abdomen table it returned $p = 0.013$ for a comparison that a 200,000-draw Monte-Carlo test places below 5×10^{-6} . The enumerated tail is deterministic and reaches far tails without the $1/B$ floor of resampling. Its corollary is a hard floor for small cohorts: eight patients admit $2^8 = 256$ sign patterns, so no one-sided p can fall below $1/256$ however consistent the effect.

What is reported alongside. Every generated significance report also carries a per-contrast table of mean Δ with the raw, uncorrected Wilcoxon p for each contrast individually, and the ablation ladders use the paired Wilcoxon throughout. The standard test is reported; it is simply not the one the pooled headline numbers rest on.

Contrast-group pooling. For Abdomen, the same tested modality is sometimes evidenced by more than one source: CHAOS's CT evidence comes from its own held-out CT split plus AMOS and SLIVER07. Averaging every raw column with equal weight – the natural first implementation – silently lets a modality with more contributing sources outvote one with fewer in both the displayed *all* average and the significance test, since it then supplies more of the equally-weighted columns being averaged. We instead pool disjoint sources of the *same* modality by concatenating their cases into one group before any averaging happens, then average modality groups with equal weight – for Abdomen this additionally means averaging equal-weight across liver/spleen/kidney *within* the pooled CT evidence, since AMOS's own CT column already blends all three organs into one number before it reaches this step. The same pooling applies to Mandible (HaN-Seg and PDDCA both supply CT) and Breast (I-SPY2, Duke-Breast-MRI and I-

Setting	Base	noEM	EM	SRCSM	Auglab	Ours
Glioma T1n	25.3*	8.7*	43.1*	40.3*	46.3	46.4
Glioma T1c	28.3*	9.5*	42.9*	44.5*	46.6	47.1
Glioma T2w	18.1*	8.4*	44.1*	43.8*	47.1	47.3
Glioma FLAIR	21.8*	8.7*	42.3	37.5*	39.4*	43.0
Abdomen T1in	29.5*	63.3*	84.4*	84.5	84.7*	86.0
Abdomen T2spir	29.6*	58.5*	83.5*	84.8*	84.5*	85.8
Brain T1w	27.5*	66.0*	67.3*	66.4	67.6*	68.6
Brain T2w	34.8*	66.7	67.5	66.7	67.2	67.7
Brain DWI	42.6*	61.6	64.2	66.8	63.4	63.1
MS FLAIR	25.2*	4.6*	32.8*	34.6*	32.8*	39.5
MS T1w	16.7*	4.0*	35.1*	31.4*	35.8	37.4
Mandible CBCT	59.2*	53.6*	76.3*	79.0	78.3	77.7
Breast T1-CE	34.1*	16.6*	40.4*	18.0*	41.1*	41.9
Breast T2w	35.9*	16.5*	38.9*	14.8*	39.0*	40.5
Spine CT	27.0*	46.7*	79.1*	—	77.5*	80.2
Spine Dixon in	69.9*	48.4*	82.2*	—	84.2	84.3
Pelvis CT	35.6*	17.1*	66.5*	71.4	68.5	69.2
Pelvis MRI	24.0*	12.8*	58.6*	59.2*	60.0	62.0

Table 7. Per-setting Dice (%), best per row in bold. Each row is one training modality of one task, scored on exactly the held-out contrasts and external cohorts of its task column in Tab. 3, so the rows of a task average to that column. *: PALETTE-Aug significantly better than that method within the setting (same patient-level sign-flip test, Holm-corrected across competitors). Spine has no SRCSM run.

SPY1 supply T1-CE, and Duke-Breast-MRI and I-SPY1 also supply pre-contrast); the remaining tasks have no contrast evidenced by more than one source.

Tables 7 and 8 break each task column of Tab. 3 into its training modalities. Where a margin is not significant in a single setting, the honest reading is that one setting cannot separate the methods, not that they are equivalent: several settings have few test units (eight patients on MS, eight sessions on Brain), which is why the main claims rest on the task- and cross-task-level tests.

Boundary distance at task level. Tab. 9 gives the HD95 counterpart of Tab. 3. Across tasks PALETTE-Aug is significantly better than every competitor, Auglab included. Per task it has the lowest HD95 on Glioma, Abdomen, Brain, MS, Spine and Pelvis (on Abdomen the gap to SRCSM is not significant); it ties Auglab on Breast without a significant gap; and on Mandible it is significantly behind SRCSM, Auglab and the baseline, the one task where the baseline beats PALETTE-Aug. When HD95 disagrees with Dice for a method – a noticeably better HD95 than its Dice would predict – the usual explanation is low recall: a method that predicts only a small, confident, well-localized fragment of the structure can have a deceptively good boundary distance while missing most of the target, so a strong HD95 paired with a weak Dice is a warning sign

Setting	Base	noEM	EM	SRCSM	Auglab	Ours
Glioma T1n	41.0*	78.0*	18.4*	20.8*	15.1	15.4
Glioma T1c	40.3*	78.3*	19.5*	14.9	15.1	15.4
Glioma T2w	54.9*	79.7*	16.5*	15.7*	14.8	15.0
Glioma FLAIR	43.1*	76.0*	17.3	18.5*	19.2*	16.6
Abdomen T1in	119.0*	78.0*	25.4	22.7	30.4*	22.3
Abdomen T2spir	125.1*	83.2*	22.0	26.0	25.1	23.2
Brain T1w	33.4*	10.4*	5.4	6.7	5.2*	4.3
Brain T2w	23.7*	4.6*	4.0	4.6*	4.2*	4.0
Brain DWI	14.7*	4.9	4.6	4.5	4.6	4.7
MS FLAIR	39.9*	41.5*	23.6	23.8	26.6	24.3
MS T1w	30.5*	42.3*	25.9	25.5	26.3*	24.2
Mandible CBCT	21.7	32.9*	24.9*	19.3	22.1	23.5
Breast T1-CE	61.8*	103.2*	61.7*	92.4*	57.3	59.0
Breast T2w	55.2	99.1*	60.1*	91.2*	56.9	56.3
Spine CT	78.3*	110.4*	4.5*	—	5.8*	3.9
Spine Dixon in	21.6*	113.2*	3.5*	—	3.6	3.1
Pelvis CT	146.2*	252.0*	20.0	16.0	23.5	21.6
Pelvis MRI	152.3*	271.6*	48.0	51.9*	47.8*	36.4

Table 8. Per-setting HD95 (mm, ↓), best per row in bold. Each row is one training modality of one task, scored on exactly the held-out contrasts and external cohorts of its task column in Tab. 3, so the rows of a task average to that column. *: PALETTE-Aug significantly better than that method within the setting (same patient-level sign-flip test, Holm-corrected across competitors). Spine has no SRCSM run.

rather than evidence of competitiveness.

Full ladder: numbers and significance. Fig. 2 (main paper) shows the rung sequence visually; Tab. 12 gives the underlying out-of-domain Dice at every rung of all sixteen ladders, with significance for each step against the rung immediately before it (Holm-corrected paired Wilcoxon on patients, within each ladder’s family of steps). Most steps that add an ingredient help, as expected. Eight cells are significant *decreases* (marked ↓). Two are the fill-swap step on Brain T1w and T2w. Two are the Voronoi sub-parcellation step, on Glioma T1n and Abdomen T1in, which the real-fill step then more than recovers. Three are the +Auglab rung — on Brain DWI, Abdomen T1in and Mandible, the task where the complete method places third (Sec. 4.2) — and one is Mandible’s val000→val100 step. There PALETTE-Aug alone (the +real rung, 69.5) is the ladder’s maximum and adding the Auglab recipe takes it back down; adding PALETTE-Aug to Auglab also lowers Mandible Dice (Tab. 3), so the two components interact negatively there and we cannot attribute the shortfall to either. Both Abdomen ladders peak at the same rung. We have not diagnosed which Auglab operator is responsible.

We likewise do not hide that Voronoi’s own marginal contribution is task-dependent rather than uniformly positive — a separate paired significance check on all-contrast Dice (the estimand of Tab. 17, not the OOD slice above)

shows: Voronoi significantly *helps* MS FLAIR and (by a small margin) Brain T1w, is flat on Abdomen T1in (its OOD step above is a small significant decrease), and significantly *hurts* Glioma T1n under the real-fill deployed configuration as well, not only here under noise fill. The full method is still best on Glioma T1n (Tab. 7; not significantly ahead of Auglab): whatever Voronoi costs there is more than recovered by the subsequent real-intensity fill step, so the net effect of the complete pipeline is positive even though one internal component is not.

Table 13 breaks the fill swap down by training → test contrast pair: the numbers behind the thin curves of Fig. 2.

Per-label breakdown: is the fill-swap effect uniform within a task? Glioma T2w’s pooled fill-swap effect is small and not significant (Tab. 10), with one significantly negative held-out contrast inside it (T2w→T1n, −4.1 Dice). This raises a further question: within that one (training, eval) pair, does the fill-swap step affect every anatomical sub-region equally, or is the worsening itself concentrated in specific labels? We break Glioma T2w→T1n down by its four labels — necrotic core (NCR), oedema (SNFH), enhancing tumor (ET), resection cavity (RC); $n = 70$ test cases for SNFH and ET, present in essentially every case (NCR $n = 51$, RC $n = 67$, since not every tumor has a necrotic core or a resection cavity).

The worsening is concentrated in SNFH (oedema) (Tabs. 14 and 15). SNFH has the largest gap between directions (+19.9 Dice) and the largest significant sign reversal: +14.7 when trained on T1n, −8.8 when trained on T2w. RC reverses too, at a smaller size. NCR shows no significant effect in either direction, and ET improves only in the T1n→T2w direction. Fig. 5 shows this on one case: with real fill, the T2w-trained model’s SNFH prediction shrinks to a thin band near the ventricle, while its ET prediction is almost unchanged.

One candidate mechanism, flagged for future work rather than tested here: a texture pattern learned from one contrast’s rendering of a diffuse, texture-defined label like oedema may not simply fail to transfer to a different contrast’s rendering of the same tissue — it may actively mislead, if the model has come to rely on a contrast-specific cue that the other contrast does not reproduce. This would be an instance of shortcut learning [14] — specifically, a network relying on a domain-specific cue that does not reproduce across acquisition protocols, the mechanism Ouyang et al. [36] identify behind failed single-source domain generalization in cross-modality and cross-sequence medical segmentation — rather than the intended, causal feature, which a simplicity-biased network need not prefer even when it would generalize better [43]. We have not tested this mechanistically here; it is the most concrete hypothesis consistent with the label-specific, direction-dependent pattern above,

	Glioma	Abdomen	Brain	MS	Mandible	Breast	Spine	Pelvis	p vs ours
Baseline	44.8	122.1	24.0	35.2	21.7	58.5	49.9	149.3	4.7e-48
SynthSeg-noEM	78.0	80.6	6.7	41.9	32.9	101.1	111.8	261.8	2.9e-58
SynthSeg-EM	17.9	23.7	4.7	24.7	24.9	60.9	4.0	34.0	1.8e-6
SRCSM	17.5	24.4	5.3	24.6	19.3	91.8	—	33.9	2.0e-34
Auglab	16.1	27.8	4.7	26.5	22.1	57.1	4.7	35.6	1.0e-4
PALETTE-Aug (ours)	15.6	22.7	4.3	24.3	23.5	57.7	3.5	29.0	—
PALETTE alone (ours)	15.8	24.9	4.6	25.6	22.3	56.3	—	35.2	—

Table 9. Task-level HD95 (mm, ↓). Same pooling and test as Tab. 3, and likewise no cross-task average. Spine has no SRCSM run. PALETTE alone: the transform without Auglab, shown for reference and left out of the bolding.




Figure 4. Figure 2 (main paper) on HD95 instead of Dice: same tasks, grouping and visual scheme. Annotation: the task’s HD95 change at the fill swap in mm (negative is better); *: significant. Glioma, Abdomen and Breast improve significantly and Brain worsens; MS, Mandible and Pelvis do not move significantly. Per-row numbers: Tab. 11.

not a demonstrated cause.

10. Synthetic contrast at validation time

Throughout the main paper PALETTE-Aug is applied with synthetic contrast during training only (*train050/val000*), because every competing method is likewise evaluated with a real-contrast validation set; this keeps the comparison exactly like-for-like. PALETTE-Aug can additionally be applied to the validation set, so that model selection is performed under randomized contrast as well (*train050/val100*).

Validation-time randomization improves most tasks on both metrics (Tab. 16), most visibly on boundary distance for Abdomen (22.7 → 20.6 mm). The effect is not significant across tasks, and MS Dice moves slightly the other

way, so we do not claim it as a separate contribution; we report it because it costs nothing and because it indicates that the benefit of contrast randomization extends to model selection and is not confined to the gradient updates.

11. Component ablation

To attribute the gain of the full method to its individual components, rather than to synthetic-contrast augmentation in general, we build an additive ablation ladder on top of the nnU-Net baseline (*Base*, Sec. 4.1): first the unsupervised k -means intensity partition alone (§3.1, step 1); then the anatomical label-remap step (step 4); then the Voronoi spatial sub-parcellation (step 2); and finally the full signed affine remap of Eq. (1) (step 3), which together constitute PALETTE-Aug on its own. Every rung is evaluated without

Task	noise	Δ	rel. (%)	p
<i>Interface-bounded</i>				
Abdomen T1in	85.7	+1.43	+1.7	4.6e-12
Abdomen T2spir	82.0	+5.11	+6.2	1.2e-20
Brain T1w	63.2	-1.34	-2.1	0.078*
Brain T2w	64.2	-1.55	-2.4	0.078*
Brain DWI	64.6	+0.78	+1.2	0.11*
Mandible CBCT [‡]	68.1	+1.36	+2.0	5.3e-6
Pelvis CT	35.3	+3.08	+8.7	0.33
Pelvis MRI	63.1	+4.12	+6.5	2.0e-6
<i>Appearance-defined</i>				
MS FLAIR	19.6	+7.69	+39.3	0.078*
MS T1w	34.1	+0.79	+2.3	1.0*
Glioma T1n	39.8	+7.22	+18.1	3.3e-11
Glioma T1c	30.5	+13.27	+43.6	5.0e-12
Glioma T2w	41.0	+0.69	+1.7	0.61
Glioma FLAIR	31.9	-0.19	-0.6	1.0
Breast T1-CE	30.8	-0.85	-2.8	0.33
Breast T2w	39.8	+0.50	+1.3	0.11

Table 10. Replacing SynthSeg-style noise fill with the real-intensity remap, with the *partition and region target means unchanged*, for every training modality of every task with a ladder. noise: out-of-domain Dice with noise fill; Δ : Dice change at the swap; rel.: Δ relative to the noise-fill Dice (bold: above 15%). p : two-sided paired Wilcoxon on patients, pooled over the row’s held-out contrasts, Holm-corrected across all sixteen rows. *With eight test patients (MS) or eight sessions of four subjects, two sites each (Brain), no row reaches $p < 0.05$ after correction even when all eight move together: MS FLAIR (8/8 improve), Brain T1w and T2w (8/8 worsen). [‡]Held-out contrasts are cross-dataset (HaN-Seg, PDDCA), since ToothFairy2 is CBCT only. Per-task summary: Fig. 2 (main paper); HD95: Tab. 11.

Auglab, so that the ladder isolates PALETTE-Aug’s own internal components from the separate question of whether it composes with a prior augmentation pipeline, which we address next. The synthetic-data fraction is fixed at 50% of training iterations throughout, as one of several augmentation constants (alongside, e.g., the k -means class count and the affine magnitude range of §3.1) that are set once and not tuned per dataset; validation uses real-contrast data only (no synthetic contrast at validation time).

Composability with Auglab. PALETTE-Aug is deployed, throughout this paper, as one additional operator inside the Auglab pipeline rather than as a replacement for it (Sec. 4.1); Auglab is a full prior augmentation method in its own right, not a component of PALETTE-Aug, so stacking it on top of the ladder above is a composability/deployment check, not a further mechanism-ablation step. Tab. 17 keeps the two separate: the top five rows isolate PALETTE-Aug’s internal mechanism, evaluated with no Auglab throughout, while the bottom row reports PALETTE-Aug plus Auglab

Task	Δ HD95 (mm)	p
<i>Interface-bounded</i>		
Abdomen T1in	-1.36	1.0
Abdomen T2spir	-18.47	3.8e-12
Brain T1w	+1.38	0.078
Brain T2w	+0.38	0.078
Brain DWI	-0.05	1.0
Mandible CBCT	-0.20	0.56
Pelvis CT	+4.19	1.0
Pelvis MRI	-0.49	1.0
<i>Appearance-defined</i>		
MS FLAIR	-3.02	0.19
MS T1w	+3.94	1.0
Glioma T1n	-4.59	7.4e-10
Glioma T1c	-9.91	4.9e-11
Glioma T2w	-1.86	0.010
Glioma FLAIR	-1.71	0.29
Breast T1-CE	-3.21	1.0e-4
Breast T2w	-1.82	0.0029

Table 11. Same fill-swap step as Tab. 10 (main paper), on HD95 (negative is better), with p Holm-corrected across the sixteen rows. Five of the six significant improvements are on appearance-defined targets, including both Breast rows, whose Dice does not move; the sixth, and the largest in millimetres, is Abdomen T2spir. HD95 is less consistent than Dice within a task (its sign flips between the two modalities of MS and of Pelvis), so we state the claim on Dice and report HD95 here rather than as confirmation of it.

Rung	MS FLAIR	MS T1w	Glioma T1n	Glioma T1c	Glioma T2w	Glioma FLAIR	Abdomen T1in	Abdomen T2spir
base	2.4	1.1	16.5	15.9	6.0	11.0	17.0	19.7
+km	11.2*	9.8*	38.4***	29.6***	38.6***	26.1***	82.2***	77.4***
+lbl	17.5*	22.4*	41.8***	30.3	40.2**	26.4	86.1***	81.4***
+vor	19.6	34.1*	39.8***↓	30.5	41.0	31.9***	85.7***↓	82.0***
+real	27.3*	34.9	47.0***	43.7***	41.6	31.7	87.1***	87.1***
+AL(v000)	30.6	36.4	46.9	44.2	46.8***	40.9***	86.4***↓	85.4
+AL(v100)	31.3	34.7	47.6*	46.5***	46.9	41.2	86.1	85.9

Rung	Brain T1w	Brain T2w	Brain DWI	Mandible CBCT	Breast T1-CE	Breast T2w	Pelvis CT	Pelvis MRI
base	14.8	24.5	36.1	39.4	23.7	33.5	0.7	10.0
+km	62.5*	63.5*	60.4*	64.6***	28.8***	37.0***	22.6***	55.9***
+lbl	62.9	63.6	64.5*	65.6***	28.5	38.8***	30.0***	63.3***
+vor	63.2	64.2*	64.6	68.1***	30.8***	39.8**	35.3*	63.1
+real	61.9*↓	62.6*↓	65.4*	69.5***	29.9	40.3*	38.3	67.2***
+AL(v000)	64.2	63.9	60.9*↓	67.6***↓	34.7***	41.9***	46.3*	67.2
+AL(v100)	64.2	64.4*	60.9	67.5***↓	—	—	—	—

Table 12. Out-of-domain Dice (%) at every rung of every ladder (Fig. 2), with significance for each step against the rung above it (val100 against val000): * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ (Holm-corrected paired Wilcoxon on patients, within each ladder); ↓ marks a significant *decrease*. Rungs: base(line), + k -means, +label-remap, +Voronoi (noise fill), +real-intensity fill (PALETTE-Aug alone), +Auglab (val000/val100). MS and Brain have eight test units each, so their smallest attainable p after correction is just under 0.05. Breast and Pelvis have no val100 arm. Mandible’s held-out axis is cross-dataset.

— the configuration used everywhere else in this paper —

BraTSGLI02108101, fold 0 -- T2w-trained model evaluated on T1n




Figure 5. BraTSGLI02108101, fold 0, T2w-trained model evaluated on T1n. Zoomed crops before (noise-fill) and after (real-fill) the fill-swap step, for the SNFH-dominant (top) and ET-dominant (bottom) lesion components of the same case. Color = ground truth wherever the prediction agrees; grey = false positive; outline = ground truth.

set apart by a rule rather than presented as the next rung of the same ladder.

12. Texture fidelity per method

The per-method NGF analysis is in the main paper (Fig. 3); Tab. 18 gives its per-task values. Figure 6 follows NGF and Dice along the ladder rungs on MS FLAIR: the three noise-fill rungs stay at NGF 0.45–0.46 while their Dice rises from 11.2 to 19.6, through contrast randomization, not texture; only the real-fill swap raises both.

Why Auglab is the closest competitor. Auglab is the strongest competitor on the cross-task test, and NGF suggests why: it keeps most of the source texture, less than PALETTE alone on six of seven tasks but far more than the label-style methods (Tab. 18). The large margins appear against the methods that discard texture.




Figure 6. Noise versus real fill on MS FLAIR. Left: out-of-domain Dice at each rung. Right: NGF similarity of the augmented volume to its source (mean $\pm$ s.e. over 30 scans). The noise-fill rungs keep only the partition edges while Dice rises; the real-fill swap (shaded) lifts both.

Task	Train → test	Noise	Real	Δ	Rel.	p
<i>Interface-bounded</i>						
Abdomen	T1in → T1out	84.9	86.1	+1.2	+1.4%	0.50
	T1in → T2spir	86.9	87.7	+0.8	+1.0%	0.88
	T1in → CT	83.4	85.2	+1.9	+2.2%	2×10^{-9}
	T1in → T2w	87.5	89.4	+1.9	+2.1%	6×10^{-6}
	T2spir → T1in	78.0	85.9	+7.9	+10.1%	0.25
	T2spir → T1out	82.4	88.4	+6.0	+7.3%	0.25
Brain	T2spir → CT	78.9	83.6	+4.6	+5.9%	9×10^{-19}
	T2spir → T2w	88.7	90.6	+1.9	+2.1%	0.03
	T1w → T2w	87.9	87.7	-0.2	-0.2%	0.09
	T1w → BOLD	49.4	46.9	-2.5	-5.1%	0.04
	T1w → DWI	72.1	71.1	-1.0	-1.4%	0.09
	T1w → EPI	48.6	46.8	-1.7	-3.6%	0.25
	T1w → GRE	58.0	56.8	-1.2	-2.2%	0.09
	T2w → T1w	90.3	89.6	-0.7	-0.7%	0.04
	T2w → BOLD	50.2	49.1	-1.1	-2.2%	0.07
	T2w → DWI	72.8	72.0	-0.8	-1.1%	0.08
	T2w → EPI	50.6	48.0	-2.6	-5.1%	0.12
	T2w → GRE	56.9	54.4	-2.6	-4.5%	0.06
	DWI → T1w	79.4	81.2	+1.8	+2.3%	0.27
	DWI → T2w	79.7	80.0	+0.4	+0.5%	0.94
Mandible	DWI → BOLD	57.0	57.7	+0.7	+1.3%	0.44
	DWI → EPI	56.9	57.4	+0.5	+0.9%	0.94
	DWI → GRE	50.1	50.5	+0.5	+0.9%	0.94
	CBCT → HaN-Seg CT	83.2	84.1	+0.9	+1.0%	5×10^{-3}
Pelvis	CBCT → HaN-Seg MR-T1	54.3	55.7	+1.5	+2.7%	0.03
	CBCT → PDDCA CT	80.7	82.4	+1.6	+2.0%	8×10^{-6}
	CT → MRI	35.3	38.3	+3.1	+8.7%	0.07
	MRI → CT	63.1	67.2	+4.1	+6.5%	2×10^{-7}
<i>Appearance-defined</i>						
Glioma	T1n → T1c	49.8	51.8	+2.0	+4.0%	1×10^{-3}
	T1n → T2w	39.2	46.4	+7.1	+18.2%	2×10^{-11}
	T1n → FLAIR	30.4	42.9	+12.5	+41.2%	2×10^{-12}
	T1c → T1n	40.7	44.1	+3.4	+8.4%	4×10^{-3}
	T1c → T2w	25.4	44.4	+19.0	+74.8%	1×10^{-12}
	T1c → FLAIR	25.3	42.7	+17.4	+68.9%	1×10^{-12}
	T2w → T1n	41.8	36.3	-5.5	-13.1%	1×10^{-4}
	T2w → T1c	37.9	41.1	+3.3	+8.7%	3×10^{-4}
	T2w → FLAIR	43.2	47.5	+4.3	+9.9%	1×10^{-5}
	FLAIR → T1n	26.0	23.8	-2.3	-8.7%	0.12
	FLAIR → T1c	32.5	32.1	-0.4	-1.1%	0.94
	FLAIR → T2w	37.2	39.3	+2.1	+5.6%	0.01
	FLAIR → T1w	15.6	21.8	+6.2	+39.5%	0.02
	FLAIR → T2w	23.5	32.8	+9.2	+39.2%	0.02
MS	T1w → FLAIR	38.3	34.5	-3.8	-9.9%	0.25
	T1w → T2w	29.9	35.2	+5.3	+17.9%	0.02
	T1-CE → T2w	21.9	22.9	+0.9	+4.2%	0.01
	T1-CE → pre-contrast	35.7	31.6	-4.1	-11.6%	1×10^{-6}
Breast	T1-CE → DWI	34.7	35.4	+0.7	+1.9%	0.01
	T2w → T1-CE	51.0	51.6	+0.6	+1.3%	0.60
	T2w → pre-contrast	40.6	40.7	+0.0	+0.1%	0.33
	T2w → DWI	27.9	28.7	+0.8	+3.0%	0.17

Table 13. The fill swap for each training → test contrast pair: out-of-domain Dice (%) of the noise-fill rung and of the real-fill rung (PALETTE alone), their difference, the difference relative to the noise-fill Dice, and the pair’s paired test (uncorrected; bold: $p < 0.05$, the pairs coloured in Fig. 2). External cohorts are pooled with the training cohort by contrast; Mandible’s held-out contrasts all come from external cohorts. 50 pairs.

Label	T1n→T2w	T2w→T1n	gap
NCR	21.5	20.0	+1.4
SNFH	66.3	46.4	+19.9
ET	27.5	29.0	-1.5
RC	52.3	42.8	+9.5

Table 14. Absolute OOD Dice (%) of the real-fill rung (PALETTE-Aug alone), per label, both transfer directions.

Label	T1n→T2w Δ	T2w→T1n Δ
NCR	+1.9 (p=0.10)	+1.0 (p=0.77)
SNFH	+14.7 (p=4.3e-11)	-8.8 (p=1e-4)
ET	+2.3 (p=0.017)	+1.6 (p=0.24)
RC	+2.9 (p=0.0003)	-6.9 (p=0.0004)

Table 15. Fill-swap step (+Voronoi noise fill → real fill), per label, both directions. p : Holm-corrected paired Wilcoxon, family = the 4 labels within each direction. **Bold** = significant.

Task	Dice $\uparrow$		HD95 (mm) $\downarrow$	
	val000	val100	val000	val100
Glioma	45.9	46.6	15.6	15.1
Abdomen	85.9	86.0	22.7	20.6
Brain	66.5	66.6	4.3	4.3
MS	38.5	37.9	24.3	24.2

Table 16. Applying PALETTE-Aug at validation time as well as at training time, on the four tasks with a task-level table for both configurations (Mandible’s val100 arm is scored only in its ladder, Tab. 12; the Breast, Spine and Pelvis arms were not trained). Across these tasks the two show no significant difference in either direction (same cross-task test as Tab. 3; two-sided Dice $p = 0.74$, HD95 $p = 0.096$), though val100 is significantly ahead within Glioma and Brain individually. We report the like-for-like val000 configuration throughout the main paper.

Component	Abdomen	MS	Glioma	Brain
	T1in	FLAIR	T1n	T1w
Base (nnU-Net default)	36.1	25.2	26.0	27.5
+ k -means partition	84.8	30.6	41.7	67.6
+ label remap	87.7	34.6	44.3	68.0
+ Voronoi sub-parcellation	87.2	36.1	42.8	68.2
+ affine remap (PALETTE-Aug alone)	87.9	40.4	48.5	63.6
+ Auglab (Ours , deployed)	87.4	39.5	47.5	68.6

Table 17. Mechanism ablation (top five rows, additive, no Auglab) and deployment composability check (bottom row). Each rung adds exactly one PALETTE-Aug component on top of the row above. Dice (%) over all test contrasts of the setting, the training contrast included, with equal weight per contrast (the all-contrast series of the same ladders as Tab. 12, which reports the held-out contrasts only). Abdomen T1in: CHAOS’s own test contrasts only, without the external cohorts.

Method	Glioma	MS	Breast	Abdomen	Brain	Mandible	Pelvis	Mean
PALETTE alone	0.836	0.842	0.809	0.805	0.721	0.904	0.863	0.826
PALETTE-Aug	0.722	0.713	0.693	0.701	0.681	0.718	0.705	0.705
Auglab	0.757	0.758	0.728	0.735	0.727	0.730	0.733	0.738
SRCSM	0.663	0.658	0.645	0.636	0.626	0.600	0.662	0.641
SynthSeg-EM	0.540	0.454	0.494	0.535	0.535	0.465	0.469	0.499
SynthSeg-noEM	0.337	0.337	0.332	0.374	0.504	0.360	0.339	0.369
noise fill (rung 4)	0.465	0.423	0.448	0.463	0.463	0.373	0.416	0.436

Table 18. Texture fidelity per task: NGF similarity of the augmented volume to its source scan (1 = all texture kept, $\approx 1/3$ = isotropic noise), mean over contrasts of the mean over scans (16–20 per contrast, 5 draws each). PALETTE alone is the ladder’s real-fill rung. PALETTE-Aug stacks PALETTE (in half of the draws) on the full Auglab chain, so its draws also lose what the Auglab augmentations remove. PALETTE alone keeps more texture than every other method on every task, except Auglab on Brain ($\Delta = -0.005$, $p = 0.68$), where the difference is not significant (paired one-sided Wilcoxon over scans, Holm-corrected over the 42 comparisons; largest significant $p = 4.1 \times 10^{-3}$).